

We are losing track

A case for pervasive causal
metadata in distributed
systems

Rodrigo Fonseca
Brown University

Who

I'm an assistant professor at Brown University

interested in Networking, Operating Systems, Distributed Systems

www.cs.brown.edu/~rfonseca

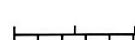
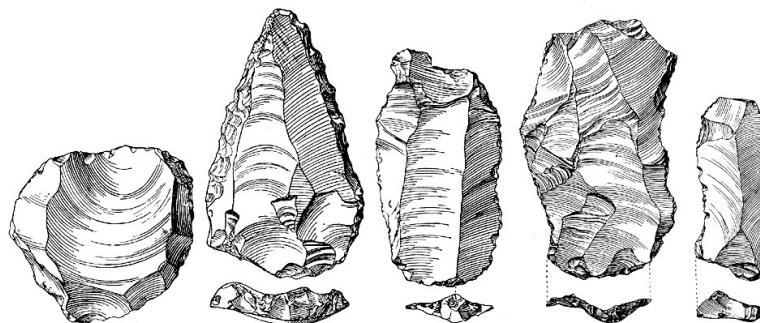


Much of this work with George Porter, Jonathan Mace, Raja Sambasivan, Ryan Roelke, Jonathan Leavitt, Sandy Riza, and many others.

In the beginning...

... life was simple

- Activity happening in one thread ~ meaningful
- Hardware support for understanding execution
 - Stack hugely helpful (e.g. profiling, debugging)
- Single-machine systems
 - OS had global view
 - Timestamps in logs made sense
- **gprof, gdb, dtrace, strace, top, ...**



Source: Anthropology: Nelson, Gilbert, Wong, Miller, Price (2012)



But then things got complicated

- **Within a node**
 - Threadpools, queues (e.g., SEDA), multi-core
 - Single-threaded event loops, callbacks, continuations
- **Across multiple nodes**
 - SOA, Ajax, Microservices, Dunghill
 - Complex software stacks
- **Stack traces, thread ids, thread local storage, logs all telling a small part of the story**



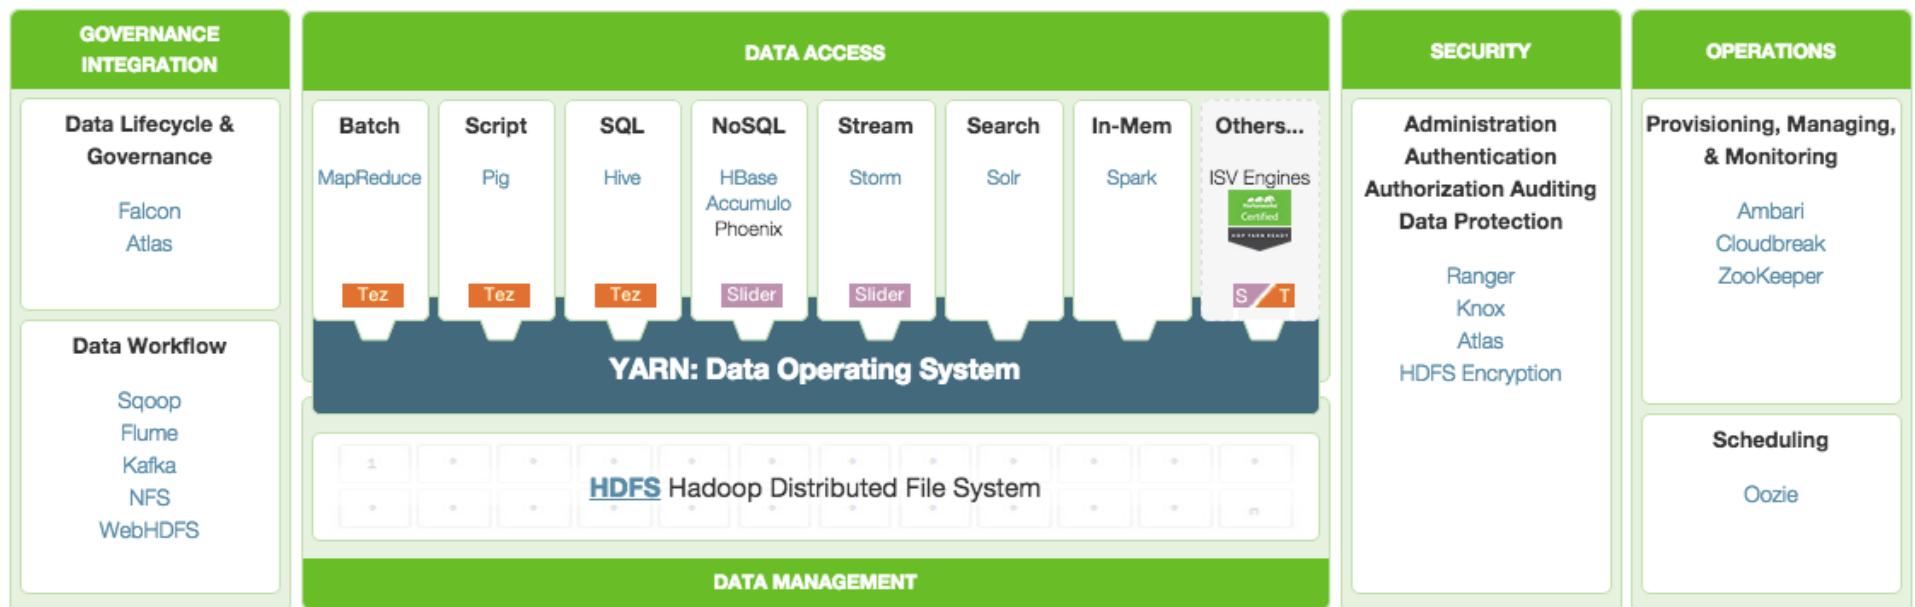
Dynamic dependencies



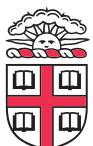
Netflix “Death Star” Microservices Dependencies

@bruce m wong

Hadoop Stack



Source: Hortonworks



Callback Hell

```
1 function hell (win) {
2     // for listener purpose
3     return function () {
4         loadLink(win, REMOTE_SRC+'/assets/css/style.css', function () {
5             loadScript(win, REMOTE_SRC+'/lib/async.js', function () {
6                 loadScript(win, REMOTE_SRC+'/lib/easyXDM.js', function () {
7                     loadScript(win, REMOTE_SRC+'/lib/json2.js', function () {
8                         loadScript(win, REMOTE_SRC+'/lib/underscore.min.js', function () {
9                             loadScript(win, REMOTE_SRC+'/lib/backbone.min.js', function () {
10                                loadScript(win, REMOTE_SRC+'/dev/base_dev.js', function () {
11                                    loadScript(win, REMOTE_SRC+'/assets/js/deps.js', function () {
12                                        loadScript(win, REMOTE_SRC+'/src/'+win.loader_path+'/loader.js', function () {
13                                            async.eachSeries(SCRIPTS, function (src, callback) {
14                                                loadScript(win, BASE_URL+src, callback);
15                                            });
16                                        });
17                                    });
18                                });
19                            });
20                        });
21                    });
22                });
23            });
24        });
25    );
26 }
```

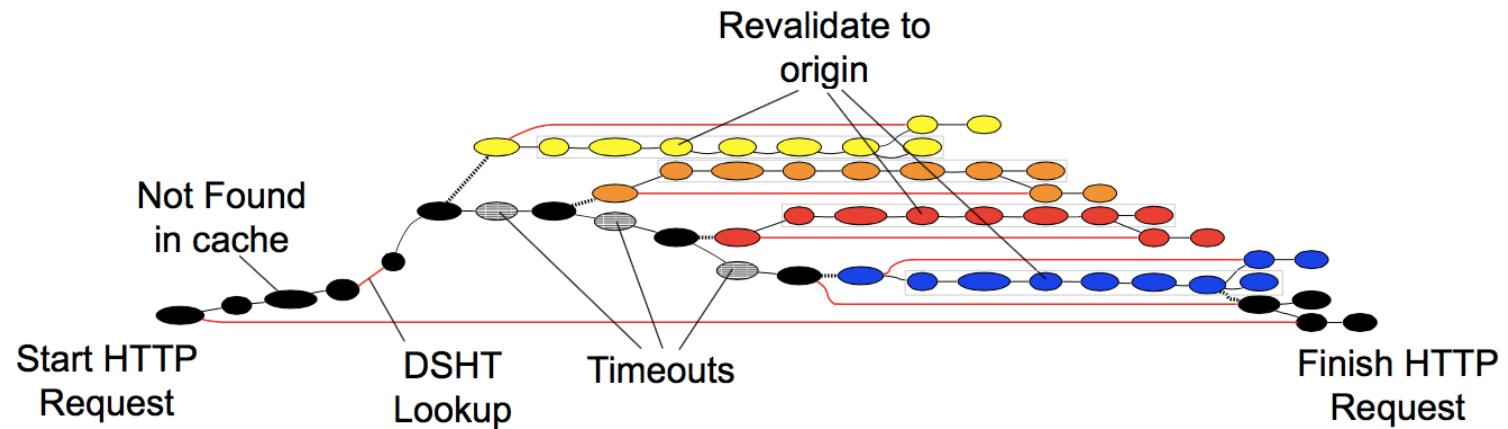


End-to-End Tracing

- **Capture the flow of execution back**
 - Through non-trivial concurrency/deferral structures
 - Across components
 - Across machines



End-to-End Tracing

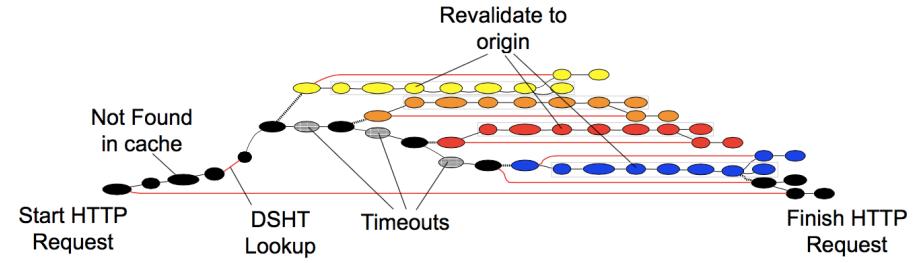
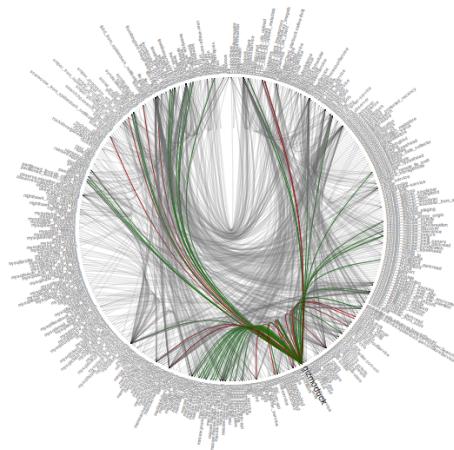


Source: X-Trace, 2008

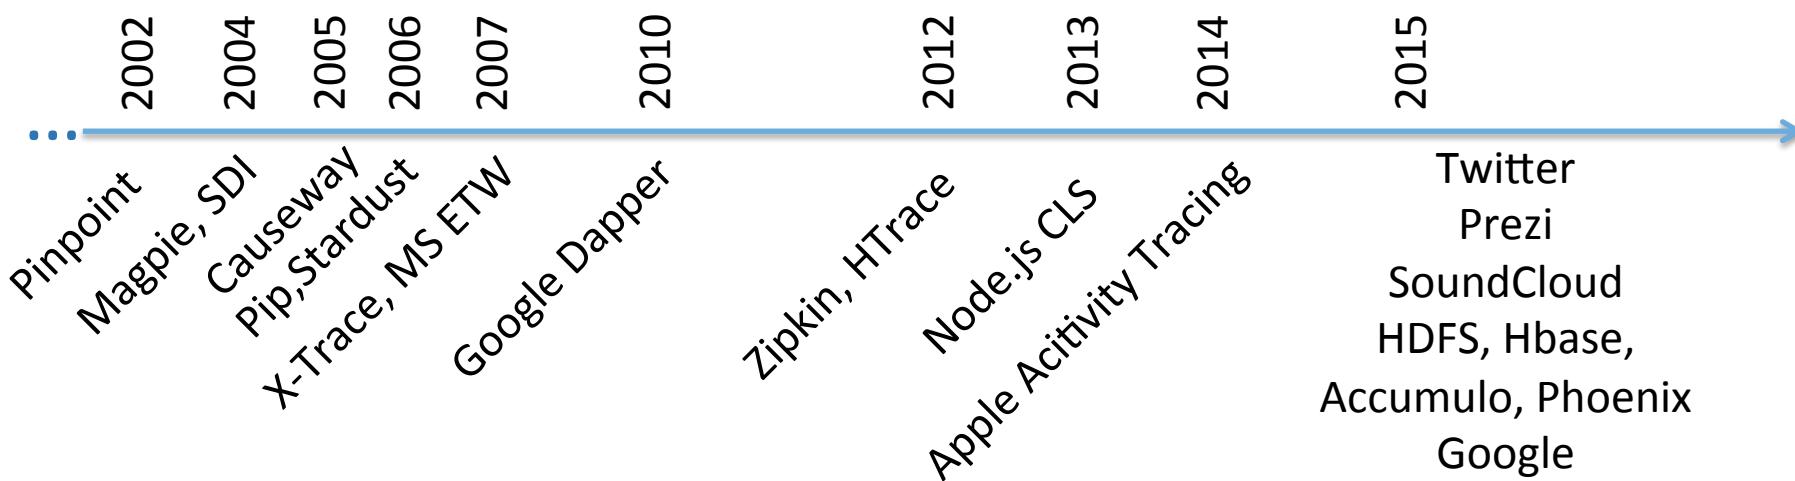
End-to-End Tracing



Source: AppNeta



End-to-End Tracing



AppNeta
AppDynamics
NewRelic



End-to-End Tracing

- **Propagate metadata along with the execution***
 - Usually a request or task id
 - Plus some link to the past (forming DAG, or call chain)
- **Successful**
 - Debugging
 - Performance tuning
 - Profiling
 - Root-cause analysis
 - ...



* Except for Magpie

- Propagate metadata along with the execution



Causal Metadata Propagation

Can be extremely useful and valuable

But...

requires instrumenting your system

(which we repeatedly have found to be doable)



**Of course, you may not want to
do this**



- **You will find IDs that already go part of the way**
- **You will use your existing logs**
 - Which are a pain to gather in one place
 - A bigger pain to join on these IDs
 - Especially because the clocks of your machines are slightly out of sync
- **Then maybe you will sprinkle a few IDs where things break**
- **You will try to infer causality by using incomplete information**



“10th Rule of Distributed System Monitoring*”

“Any sufficiently complicated distributed system contains an ad-hoc, informally-specified, siloed implementation of causal metadata propagation.”



*This is, of course, inspired by Greenspun’s 10th Rule of Programming

Causal Metadata Propagation

- **End-to-End tracing**
 - Similar, but incompatible contents
- **Same *propagation***
 - Flow along thread while working on same activity
 - Store and retrieve when deferred (queues, callbacks)
 - Copy when forking, merge when joining
 - Serialize and send with messages
 - Deserialize and set when receiving messages



Causal Metadata Propagation

- **Not hard, but subtle sometimes**
- **Requires commitment, touches many places in the code**
- **Difficult to completely automate**
 - Sometimes the causality is at a layer above the one being instrumented
- **You will want to do this only once...**

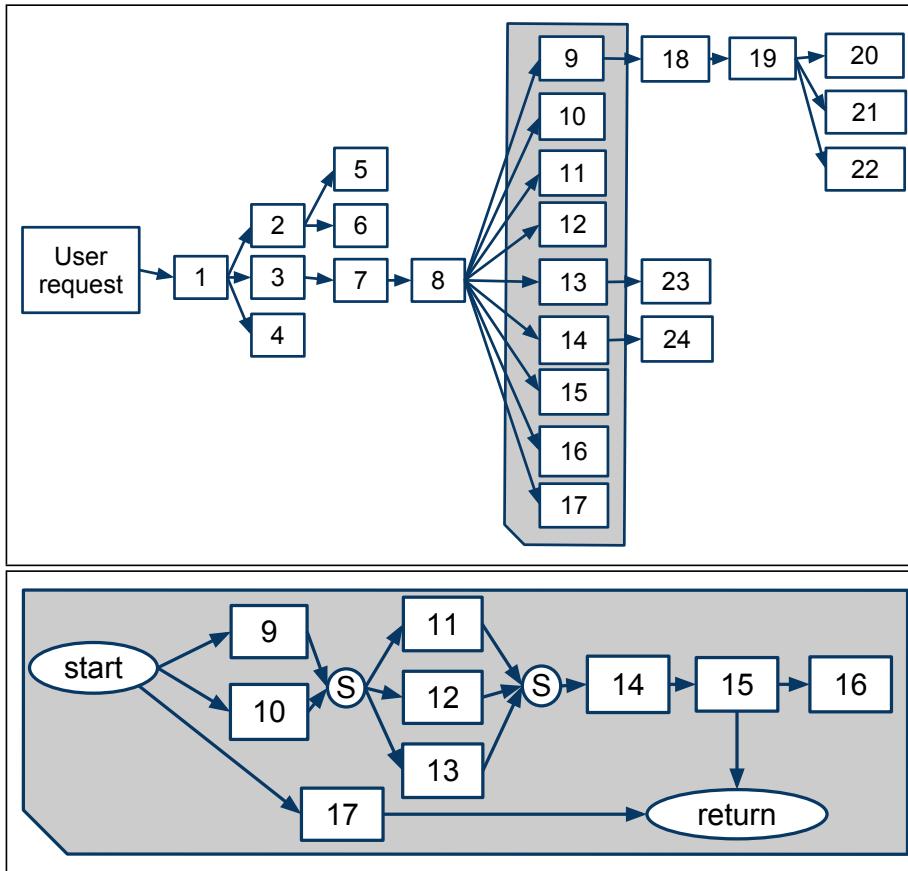


Causal Metadata Propagation

... or you won't have another chance



Modeling the Parallel Execution of Black-Box Services. Mann et al., HotCloud 2011 (Google)



The Dapper Span model doesn't natively distinguish the causal dependencies among siblings



Causal Metadata Propagation

- Propagation currently coupled with the data model
- Multiple different uses for causal metadata

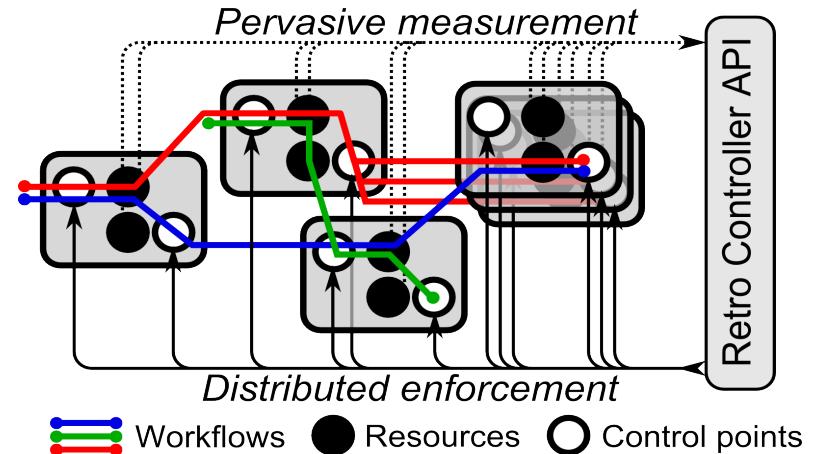


A few more (different) examples

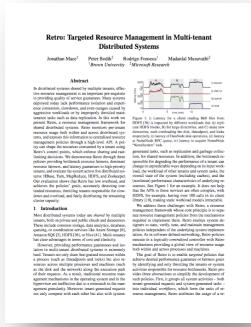
- ...
- **Timecard – Ravindranath et al., SOSP’13**
- **TaintDroid – Enck at al., OSDI’10**
- ...



Retro

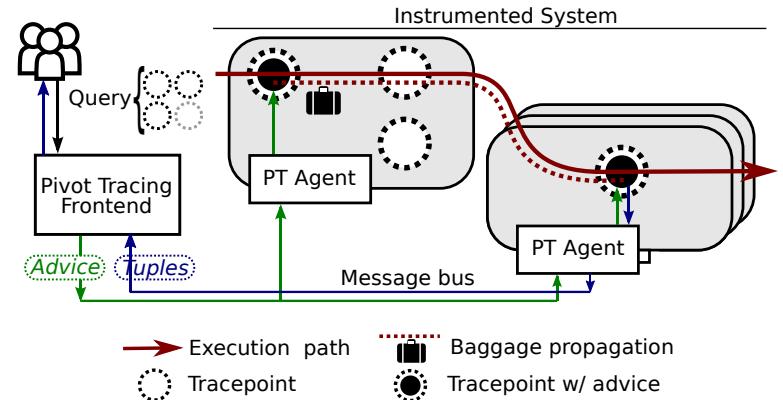


- Propagates TenantID across a system for real-time resource management
- Instrumented most of the Hadoop stack
- Allows several policies – e.g., DRF, LatencySLO
- Treats background / foreground tasks uniformly



Jonathan Mace, Peter Bodik, Madanlal Musuvathi, and Rodrigo Fonseca. Retro: targeted resource management in multi-tenant distributed systems. In NSDI '15

Pivot Tracing



- Dynamic instrumentation + Causal Tracing

```
From incr In DataNodeMetrics.incrBytesRead  
Join c1 In First(ClientProtocols) On cl -> incr  
GroupBy cl.procName  
Select cl.procName SUM(incr.delta)
```

- Queries → Dynamic Instrumentation → *Query-specific metadata* → Results
Implemented generic metadata layer, which we called *baggage*



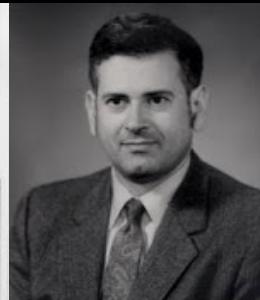
Jonathan Mace, Ryan Roelke, and Rodrigo Fonseca. Pivot Tracing: Dynamic Causal Monitoring for Distributed Systems. SOSP 2015

So, where are we?

- **Multiple interesting uses of causal metadata**
- **Multiple incompatible instrumentations**
 - Coupling propagation with content
- **Systems that increasingly talk to each other**
 - c.f. Death Star



1973

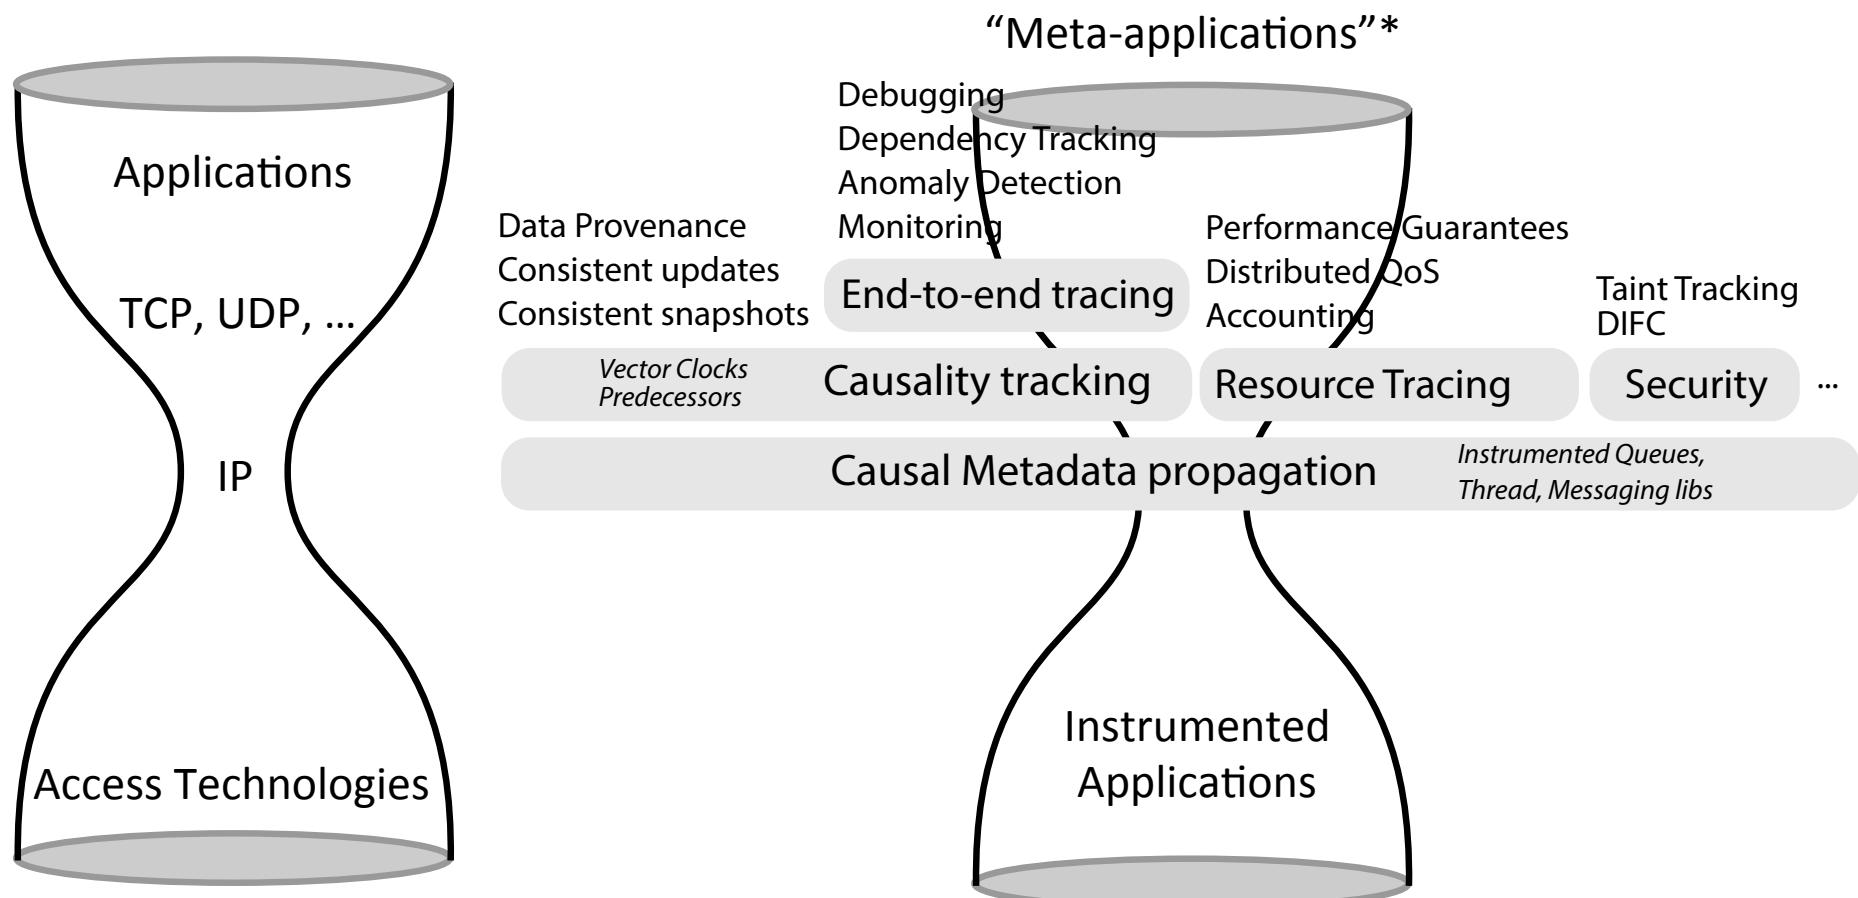


IP

- **Packet switching had been proven**
 - ARPANET, X.25, NPL, ...
- **Multiple incompatible networks in operation**
- **TCP/IP designed to connect all of them**
- **IP as the “narrow waist”**
 - Common format
 - (Later) minimal assumptions, no unnecessary burden on upper layers



Obligatory ugly hourglass picture



*Causeway (Chanda et al., Middleware 2005) used this term

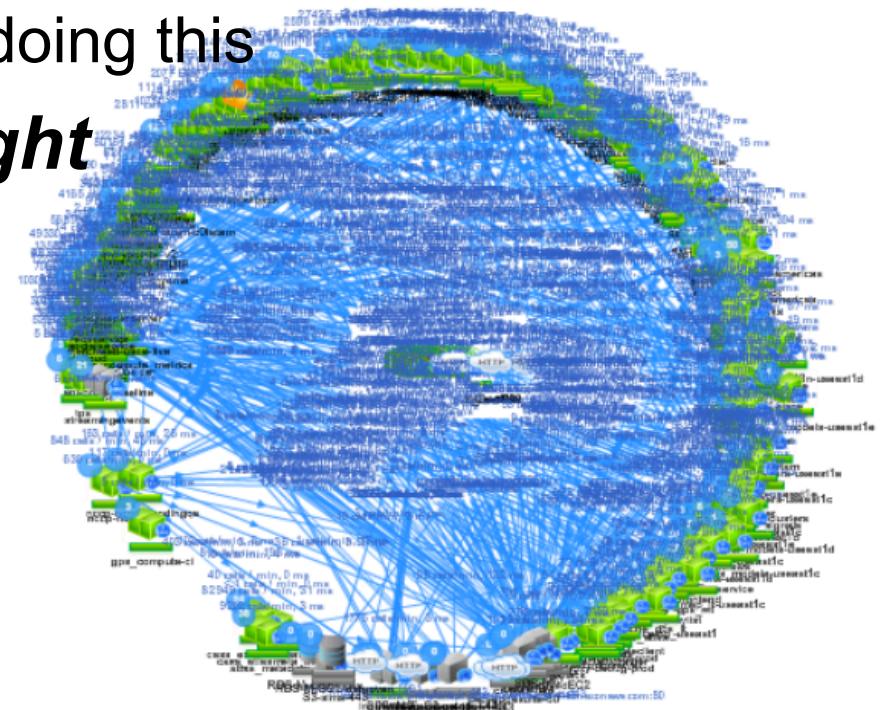
Proposal: Baggage

- API and guidelines for causal metadata propagation
- Separate propagation from semantics of data
- Instrument systems once, “baggage compliant”
- Allow multiple meta-applications



Why now?

- We are losing track...
- Huge momentum (Zipkin, HTrace, ...)
 - People care and ARE doing this
- Right time to do it *right*



Baggage API

- **PACK, UNPACK**
 - Data is key-value pairs
- **SERIALIZE, DESERIALIZE**
 - Uses protocol buffers for serialization
- **SPLIT, JOIN**
 - Apply when forking / joining
 - Use Interval Tree Clocks to correctly keep track of data



Paulo Sérgio Almeida, Carlos Baquero, and Victor Fonte. Interval tree clocks: a logical clock for dynamic systems. In *Opodis '08*.

Big Open Questions

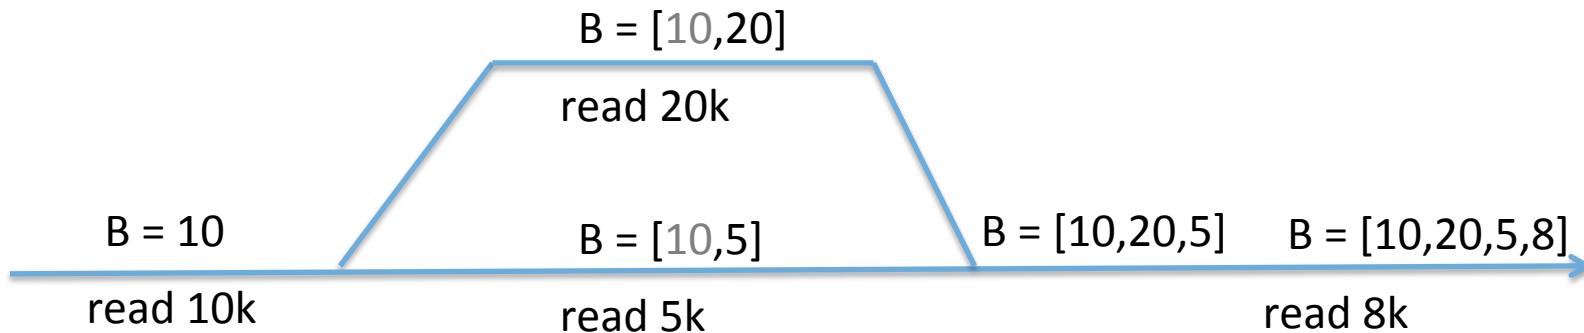
- **Is this feasible?**
 - Is the **propagation logic** the same for all/most of the meta applications?
 - Can fork/join logic be data-agnostic? Use helpers?
- **This is not just an API**
 - How to formalize the rules of propagation?
 - How to distinguish bugs in the application vs bugs in the propagation?
- **How to get broad support?**





Thank you

Example Split / Join



- **We use Interval Tree Clocks for an efficient implementation**



Paulo Sérgio Almeida, Carlos Baquero, and Victor Fonte. Interval tree clocks: a logical clock for dynamic systems. In *Opodis '08*.