



Towards General-Purpose Resource Management in Shared Cloud Services

Jonathan Mace, *Brown University*

Peter Bodik, *MSR Redmond*

Rodrigo Fonseca, *Brown University*

Madanlal Musuvathi, *MSR Redmond*



Shared-tenant cloud services

Processes service requests from multiple clients

- ✓ Great for cost and efficiency
- ✗ Performance is a challenge

Aggressive tenants and system maintenance tasks



Resource starvation and bottlenecks



Degraded performance, Violated SLOs, system outages

Shared-tenant cloud services

Ideally

manage resources to provide end-to-end guarantees and isolation

Challenge

OS/hypervisor mechanisms insufficient

- ✗ Shared threads & processes
- ✗ Application-level resource bottlenecks (locks, queues)
- ✗ Resources across multiple processes and machines


Today

lack of guarantees, isolation
some ad-hoc solutions


This paper

- 5 design principles for resource policies in shared-tenant systems
- *Retro* – prototype for **principled** resource management
- Preliminary demonstration of Retro in HDFS

Hadoop Distributed File System (HDFS)




Filesystem metadata




Replicated block storage


Hadoop Distributed File System (HDFS)























Principle 1: Consider all resources and request types

- Fine-grained resources within processes
- Resources shared between processes (disk, network)
- Many different API calls
- Bottlenecks can crop up in many places
 - hardware resources: *disk, network, cpu, ...*
 - software resources: *locks, queues, ...*
 - data structures: *transaction logs, shared batches, ...*














Principle 2: *Distinguish between tenants*


- Tenants might send different types of requests
- Tenants might be utilizing different machines
- If a policy is *efficient*, it should be able to *target* the cause of contention
 - e.g.,
 - if a tenant is causing contention, throttle otherwise leave the tenant alone











Principle 5:

*Schedule early,
schedule often*


```
while (!Thread.isInterrupted()) {  
    rate_limit();  
    sendPacket();  
}
```

Resource Management Design Principles

1. Consider all request types and all resources
2. Distinguish between tenants
3. Treat foreground and background tasks uniformly
4. Estimate resource usage at runtime
5. Schedule early, schedule often


Retro – prototype for **principled** resource management in shared-tenant systems

Retro: end-to-end tracing




 Tenants

Retro: end-to-end tracing



 Tenants

Retro: application-level resource interception




Tenants




Resources


Retro: aggregation and centralized reporting




Retro: application-level enforcement





Retro: distributed scheduling






Retro: distributed scheduling





Early Results



**HDFS NNBench
benchmark**
0.01% to 2%
average overhead
on end-to-end
latency, throughput







Retrospective

Thus far:

- Per-tenant identification
- Resource measurements
- Schedule enforcement

Next steps:

- Abstractions for writing simplified high-level policies
- Low-level enforcement mechanisms
- Policies to monitor system, find bottlenecks, provide guarantees