

# Retro

## Targeted Resource Management in Multi-tenant Distributed Systems

**Jonathan Mace**  
*Brown University*



**Peter Bodik**  
*MSR Redmond*



**Rodrigo Fonseca**  
*Brown University*



**Madanlal Musuvathi**  
*MSR Redmond*



BROWN UNIVERSITY

Microsoft®  
**Research**

# *Resource Management in Multi-Tenant Systems*

# *Resource Management in Multi-Tenant Systems*



Failure in availability zone **cascades to shared control plane**, causes thread pool starvation for all zones

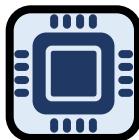
- April 2011 – Amazon EBS Failure

# *Resource Management in Multi-Tenant Systems*



Failure in availability zone **cascades to shared control plane**, causes thread pool starvation for all zones

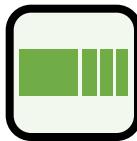
- April 2011 – Amazon EBS Failure



Aggressive **background task** responds to increased hardware capacity with deluge of warnings and logging

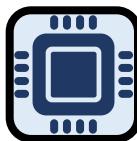
- Aug. 2012 – Azure Storage Outage

# *Resource Management in Multi-Tenant Systems*



Failure in availability zone **cascades to shared control plane**, causes thread pool starvation for all zones

- April 2011 – Amazon EBS Failure



Aggressive **background task** responds to increased hardware capacity with deluge of warnings and logging

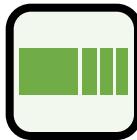
- Aug. 2012 – Azure Storage Outage



Code change increases database usage, shifts bottleneck to unmanaged **application-level lock**

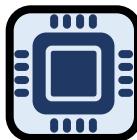
- Nov. 2014 – Visual Studio Online outage

# *Resource Management in Multi-Tenant Systems*



Failure in availability zone **cascades to shared control plane**, causes thread pool starvation for all zones

- April 2011 – Amazon EBS Failure



Aggressive **background task** responds to increased hardware capacity with deluge of warnings and logging

- Aug. 2012 – Azure Storage Outage



Code change increases database usage, shifts bottleneck to unmanaged **application-level lock**

- Nov. 2014 – Visual Studio Online outage



**Shared storage layer** bottlenecks circumvent resource management layer

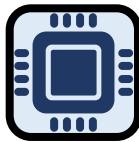
- 2014 – Communication with Cloudera

# *Resource Management in Multi-Tenant Systems*



Failure in availability zone **cascades to shared control plane**, causes thread pool starvation for all zones

- April 2011 – Amazon EBS Failure



Aggressive **background task** responds to increased hardware capacity with deluge of warnings and logging

- Aug. 2012 – Azure Storage Outage



Code change increases database usage, shifts bottleneck to unmanaged **application-level lock**

- Nov. 2014 – Visual Studio Online outage

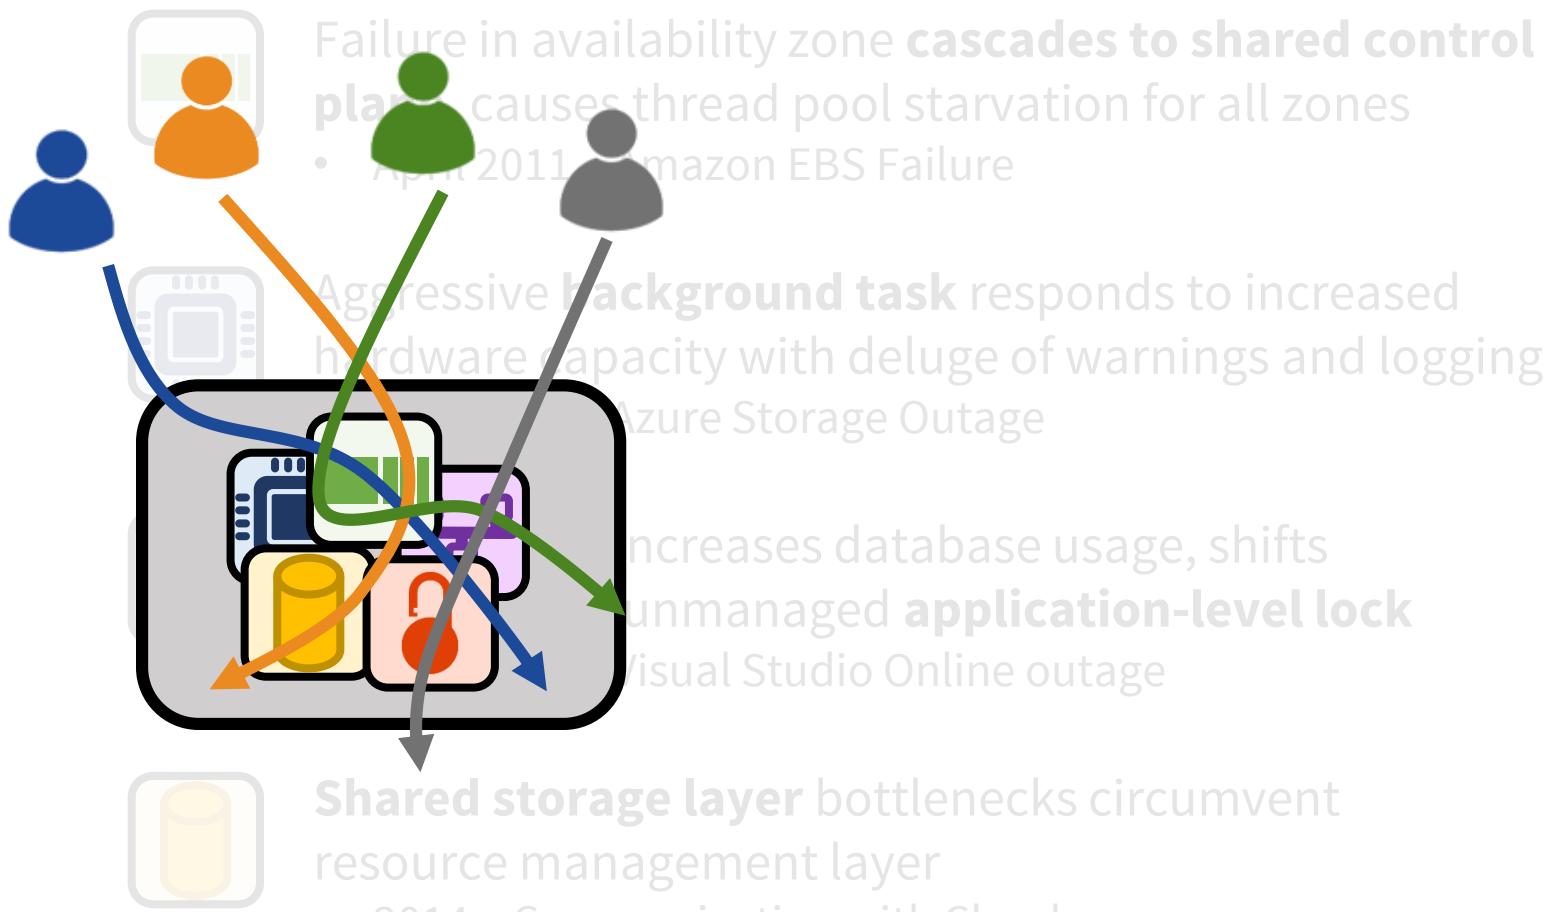


**Shared storage layer** bottlenecks circumvent resource management layer

- 2014 – Communication with Cloudera

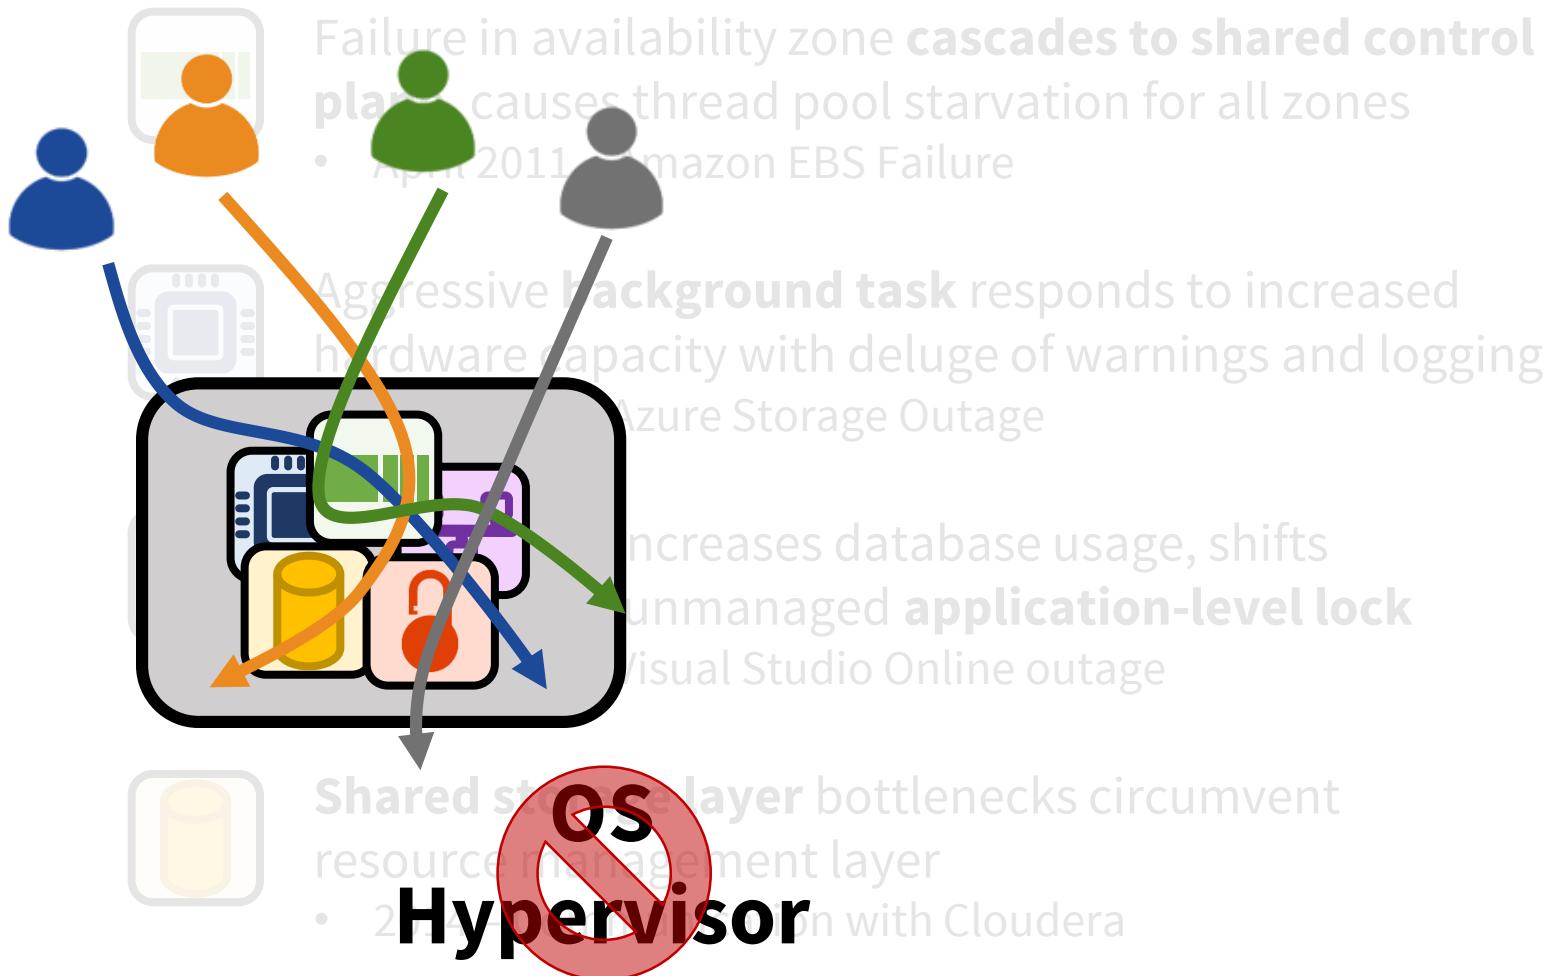
Degraded performance, Violated SLOs, system outages

# Resource Management in Multi-Tenant Systems



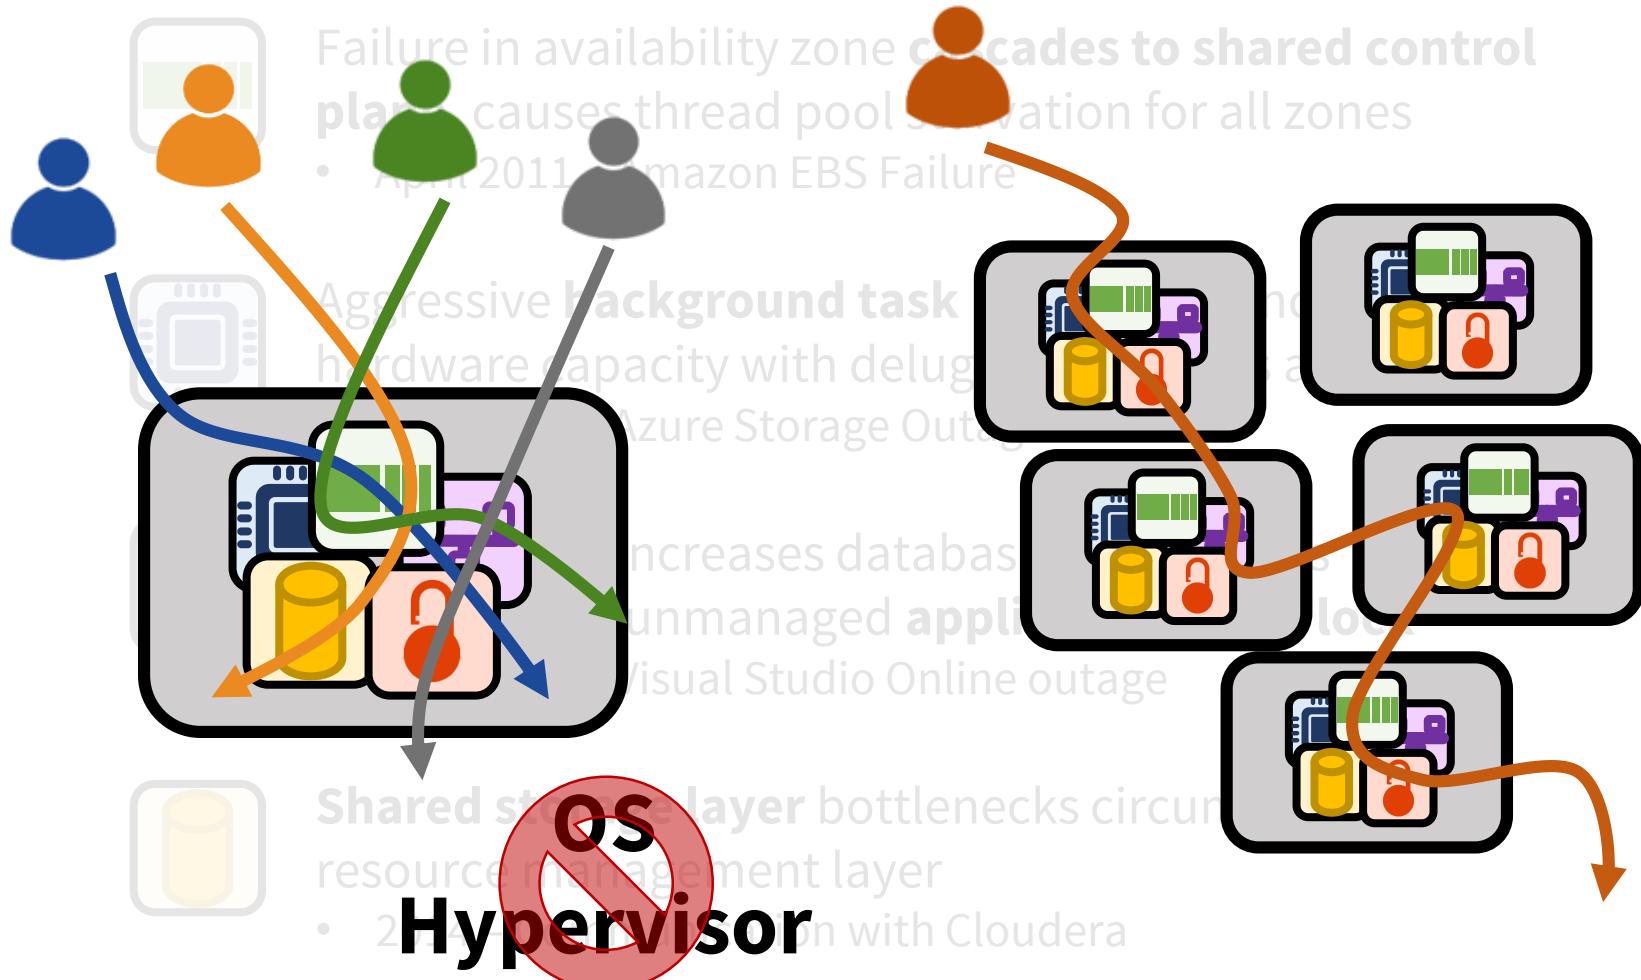
Degraded performance, Violated SLOs, system outages

# Resource Management in Multi-Tenant Systems



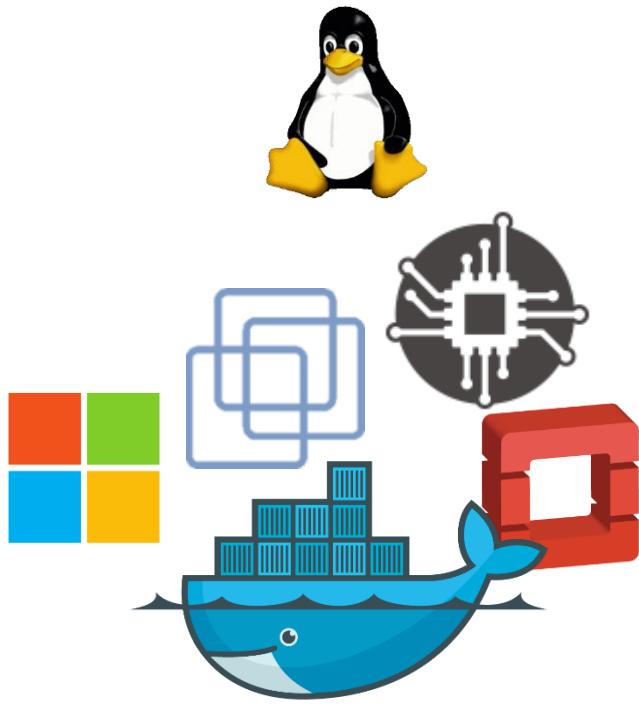
Degraded performance, Violated SLOs, system outages

# Resource Management in Multi-Tenant Systems

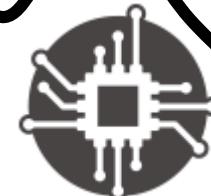


Degraded performance, Violated SLOs, system outages

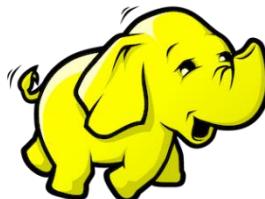
## Containers / VMs



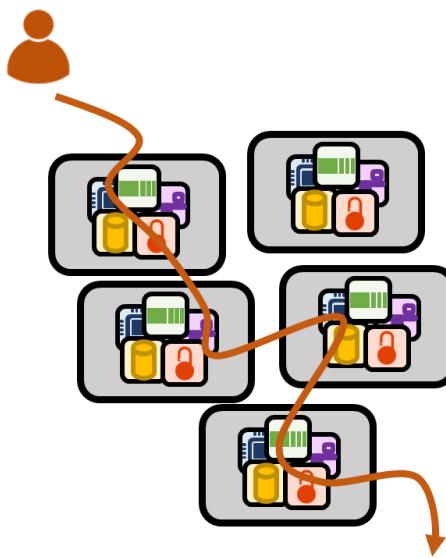
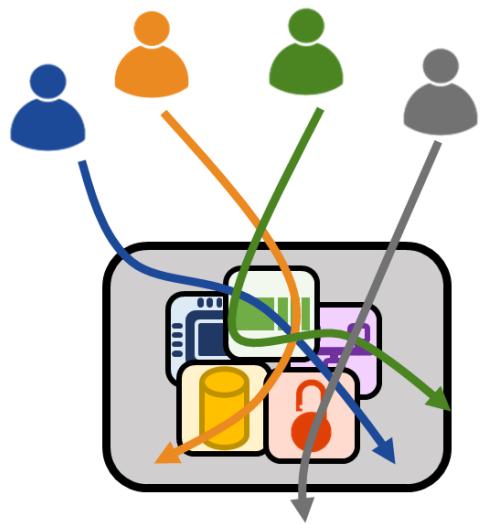
Containers / VMs



**HBASE**



Shared Systems:  
Storage, Database,  
Queueing, etc.







Monitors resource usage of each tenant in near real-time  
Actively schedules tenants and activities



Monitors resource usage of each tenant in near real-time  
Actively schedules tenants and activities

High-level, centralized policies:  
Encapsulates resource management logic



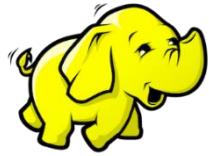
Monitors resource usage of each tenant in near real-time  
Actively schedules tenants and activities

High-level, centralized policies:

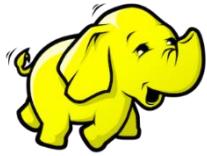
- Encapsulates resource management logic

- Abstractions – not specific to resource type, system

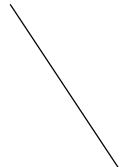
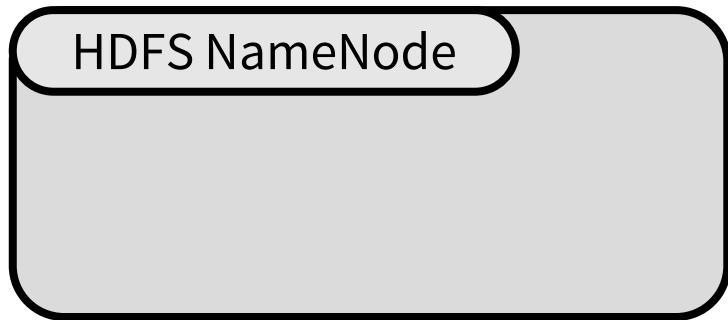
- Achieve different goals: guarantee average latencies, fair share a resource, etc.



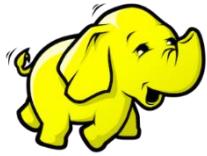
# Hadoop Distributed File System (HDFS)



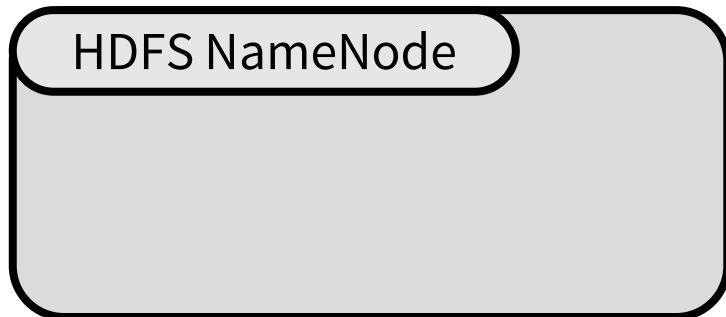
# Hadoop Distributed File System (HDFS)



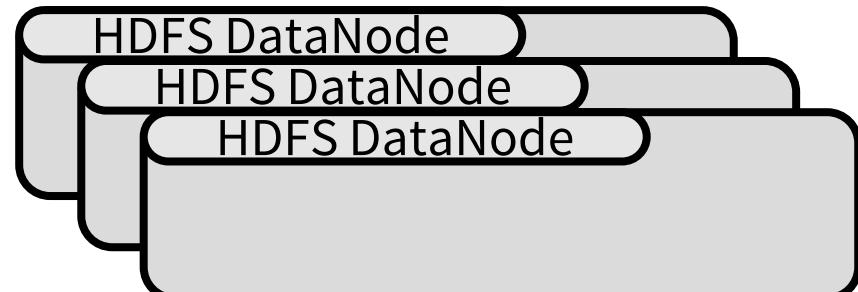
Filesystem metadata



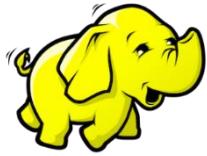
# Hadoop Distributed File System (HDFS)



Filesystem metadata

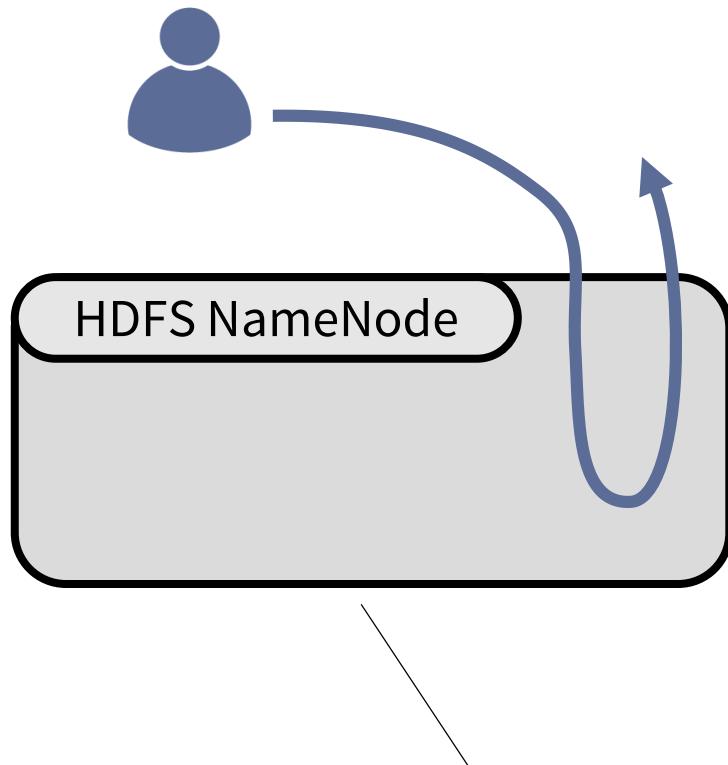


Replicated block storage

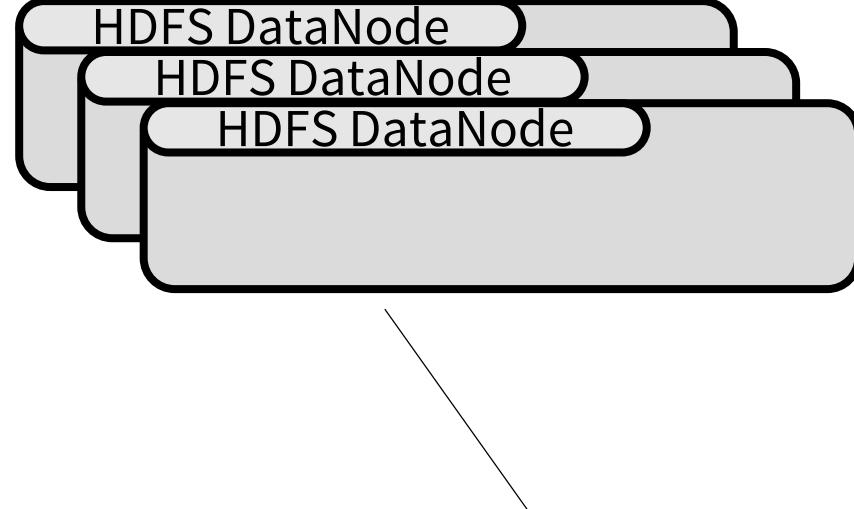


# Hadoop Distributed File System (HDFS)

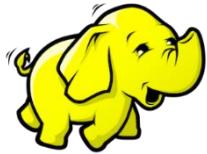
Rename



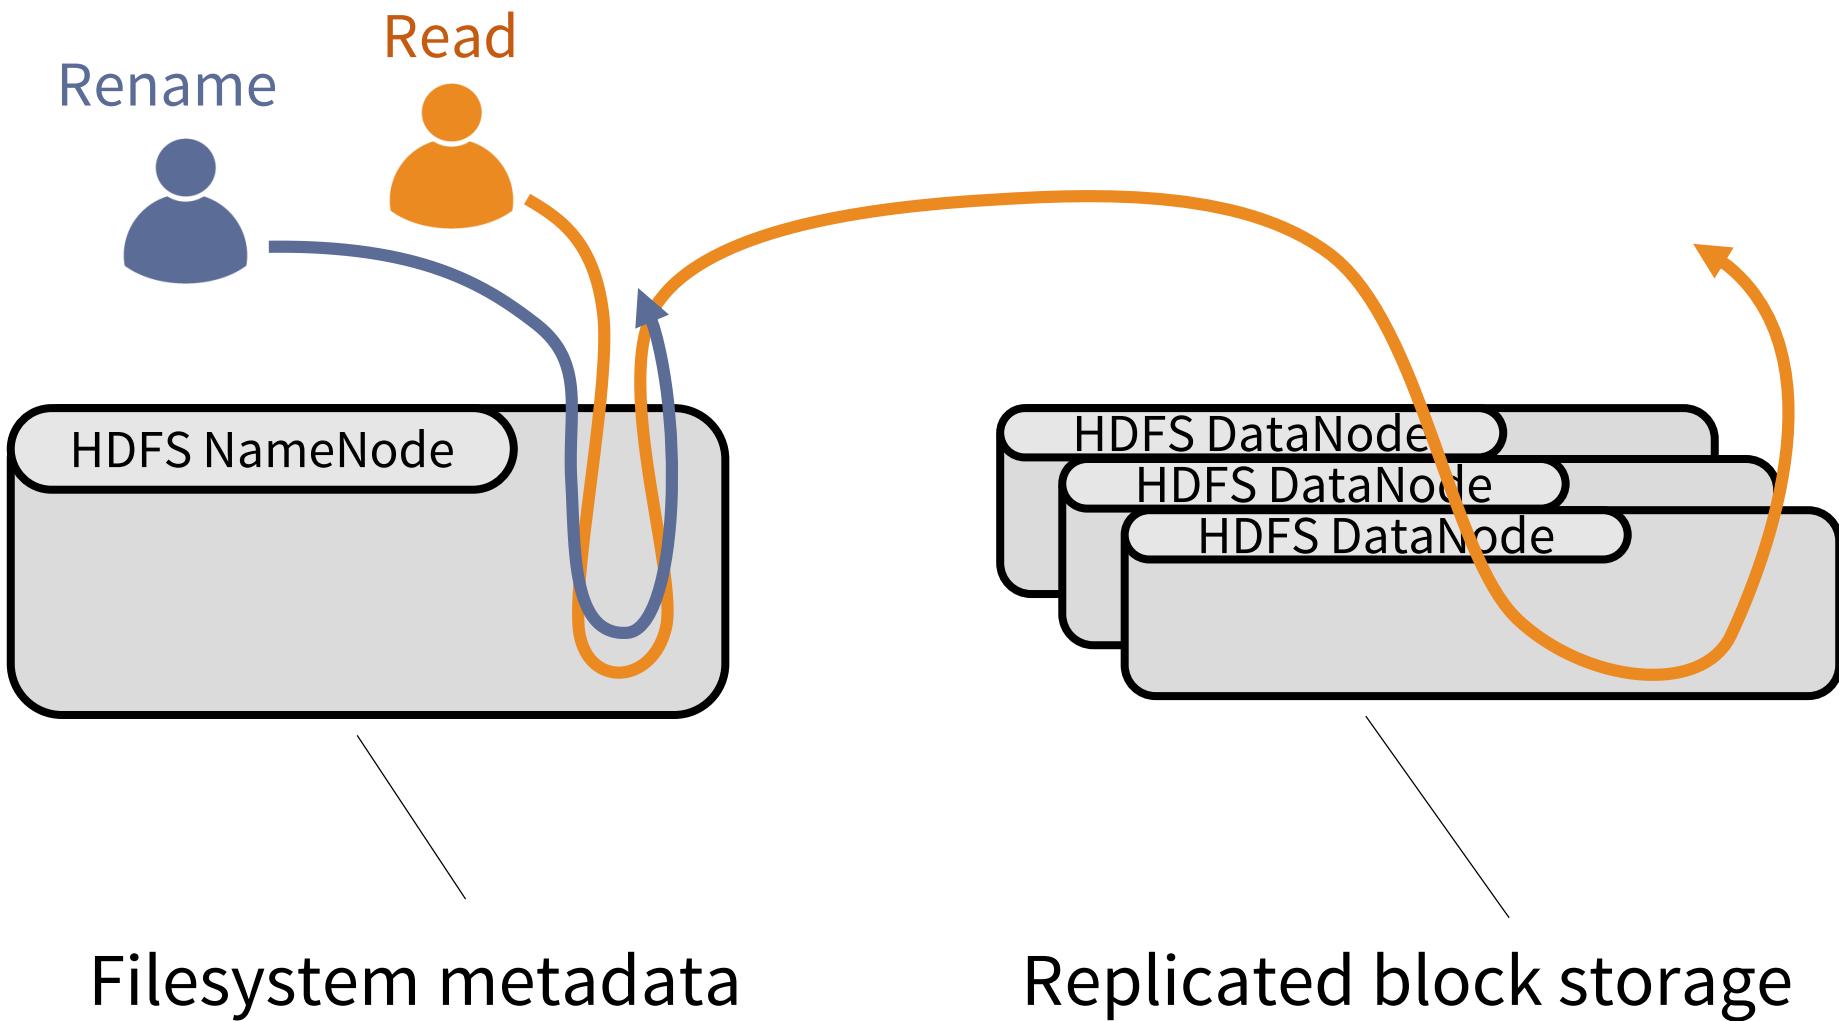
Filesystem metadata

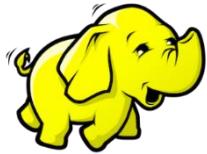


Replicated block storage

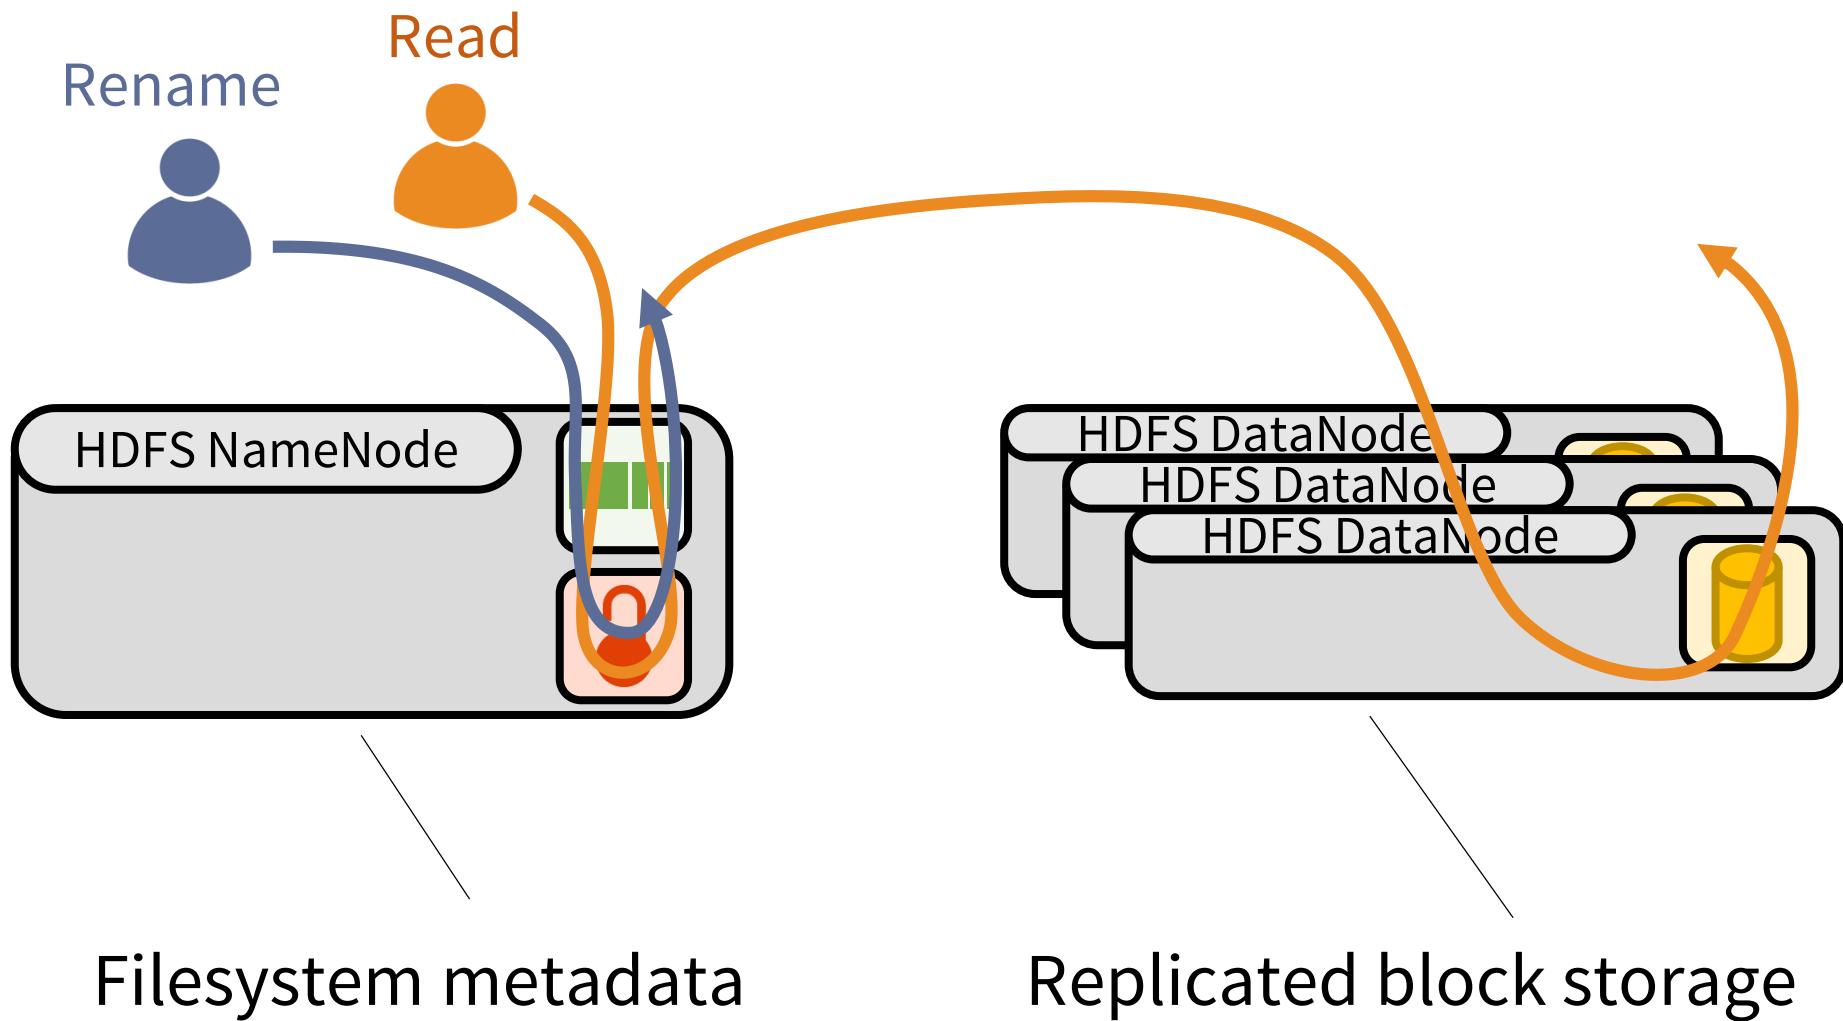


# Hadoop Distributed File System (HDFS)

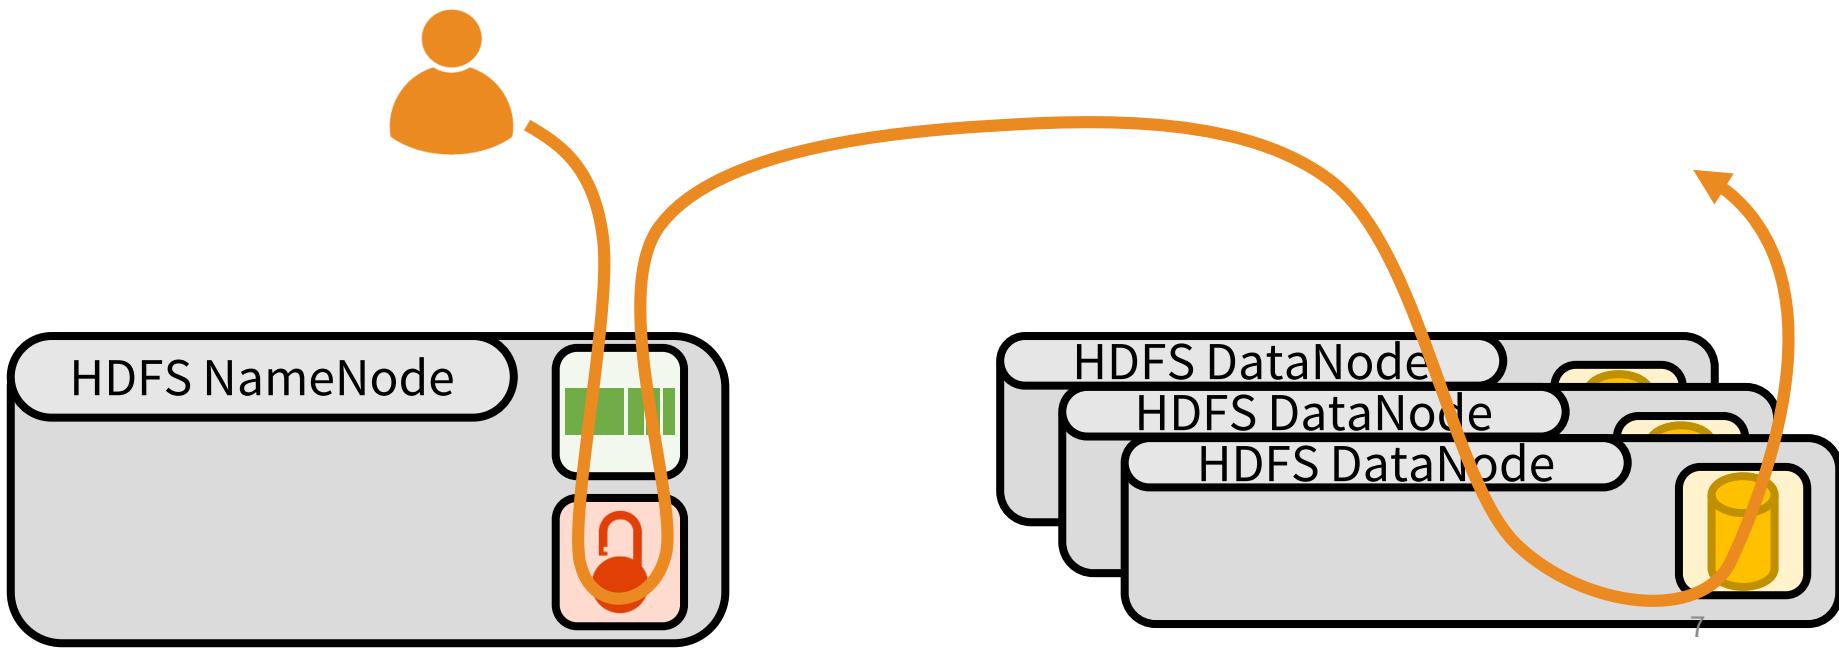


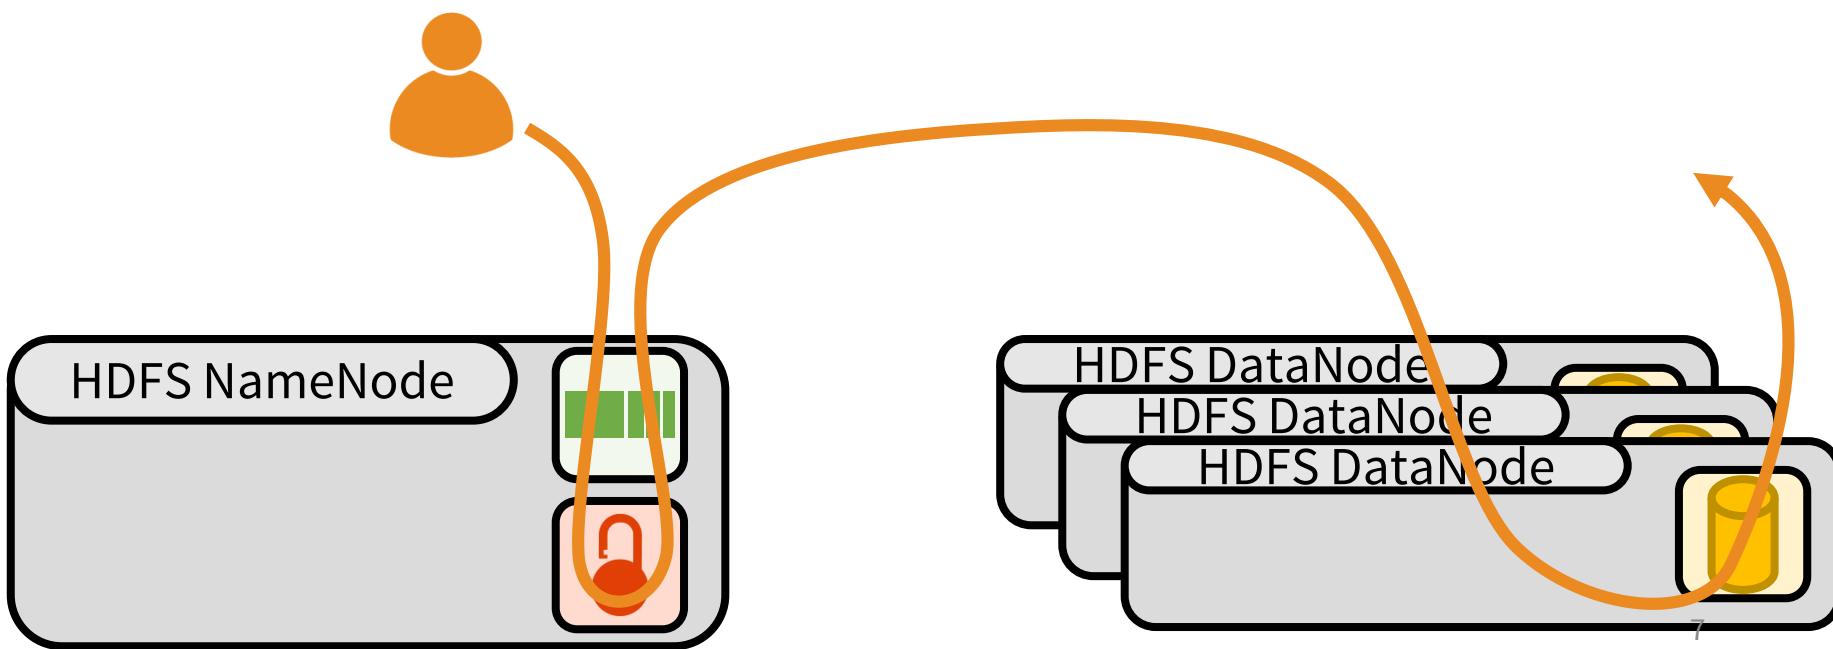
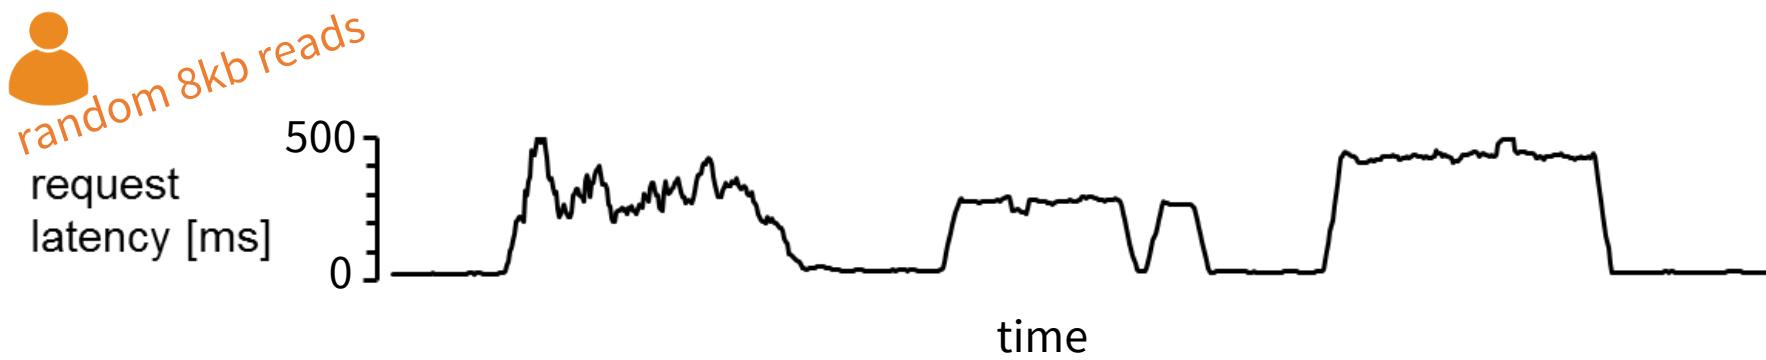


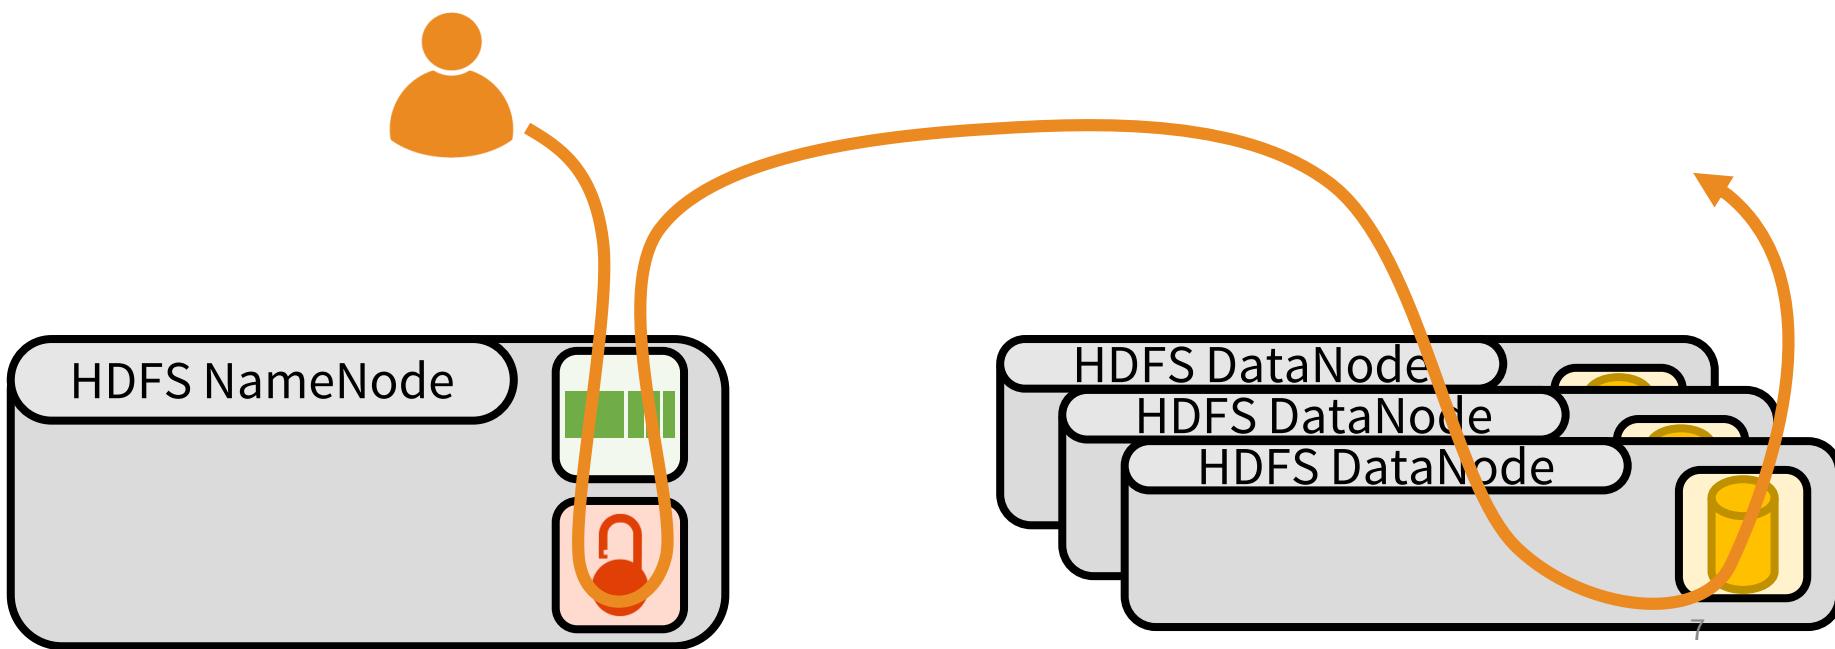
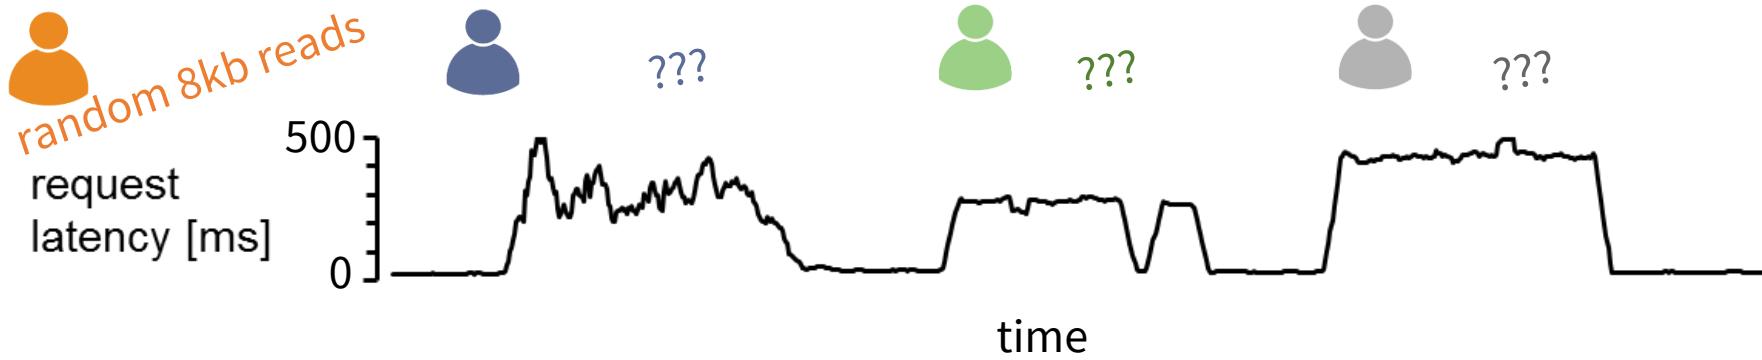
# Hadoop Distributed File System (HDFS)

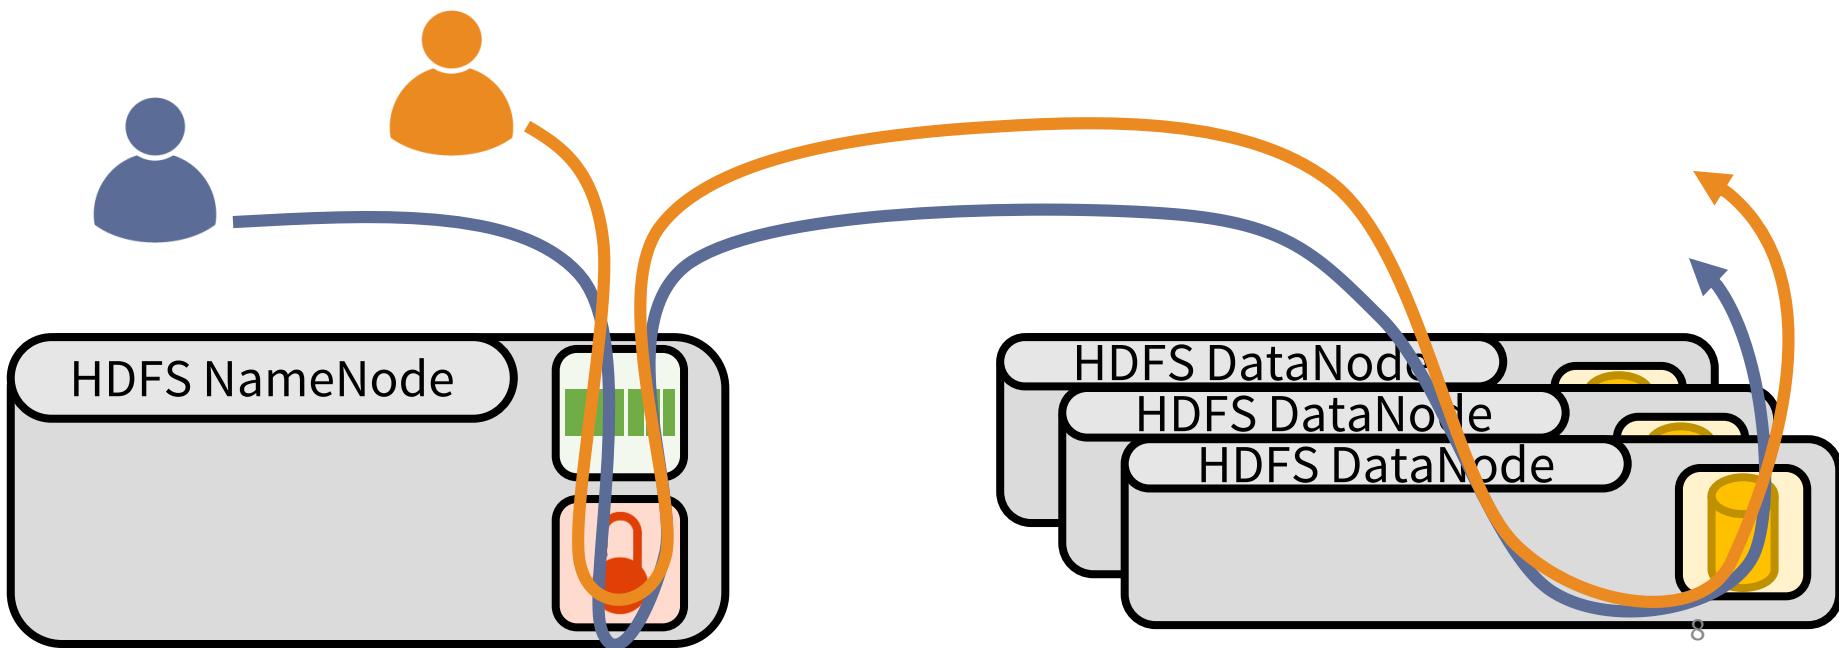
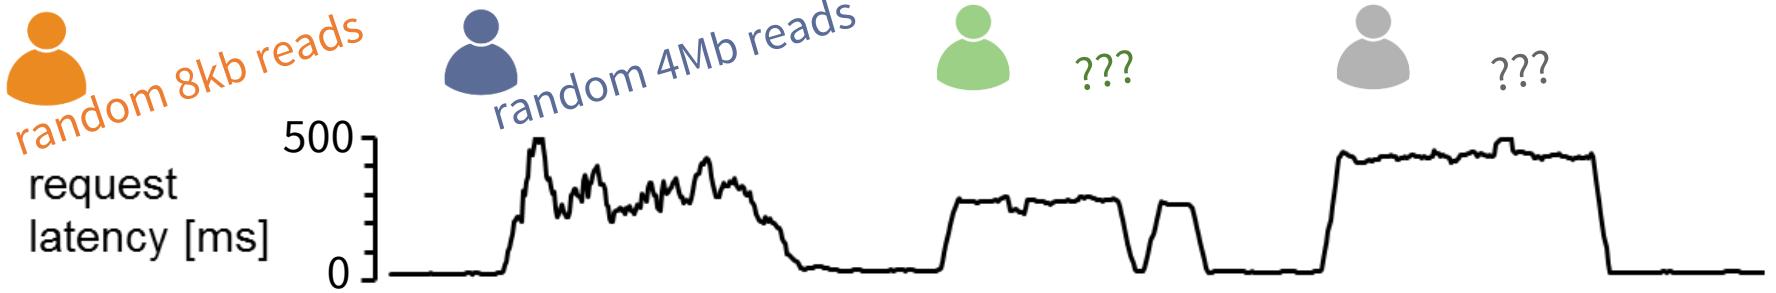


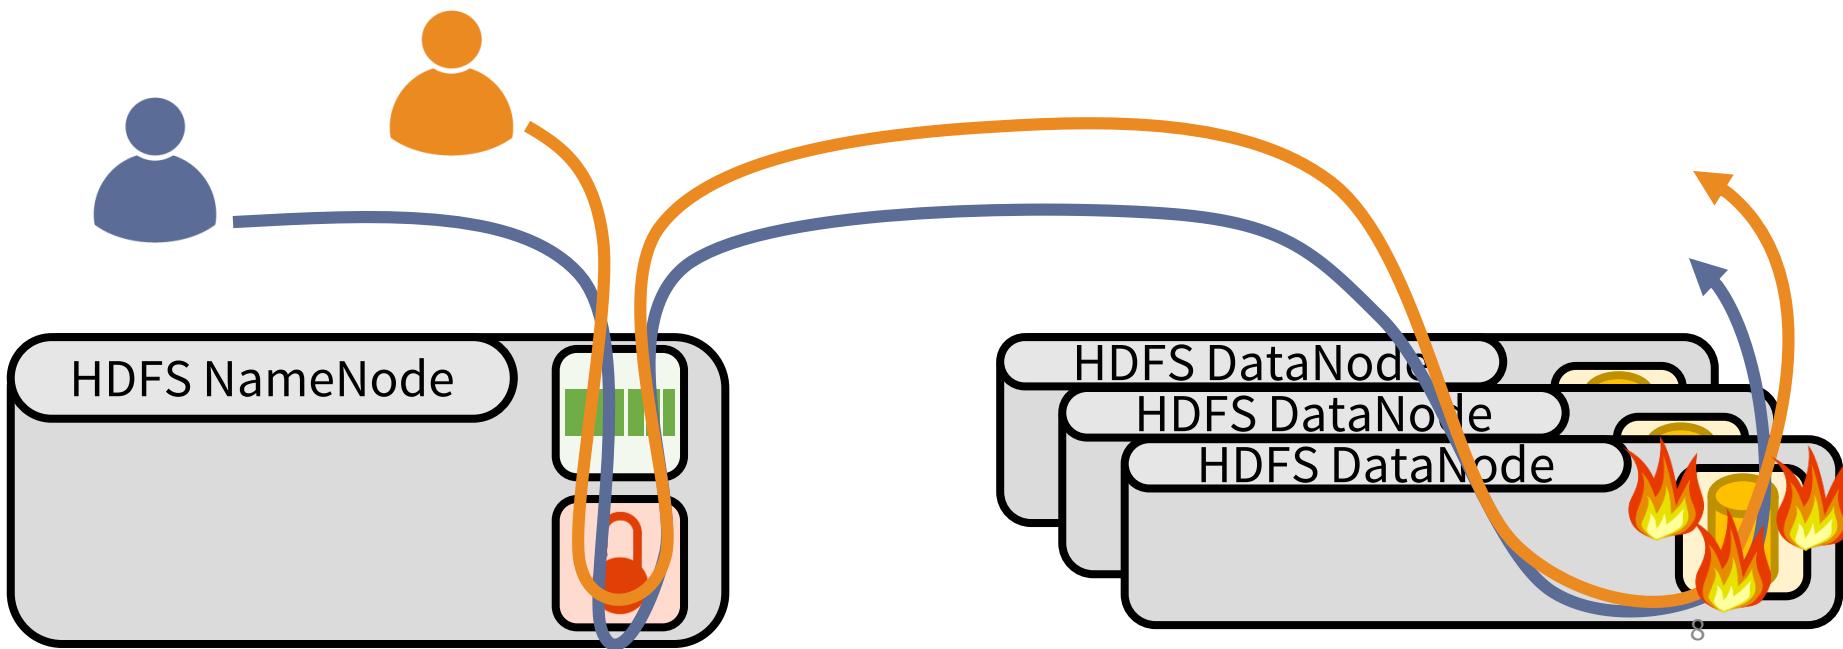
random 8kb reads

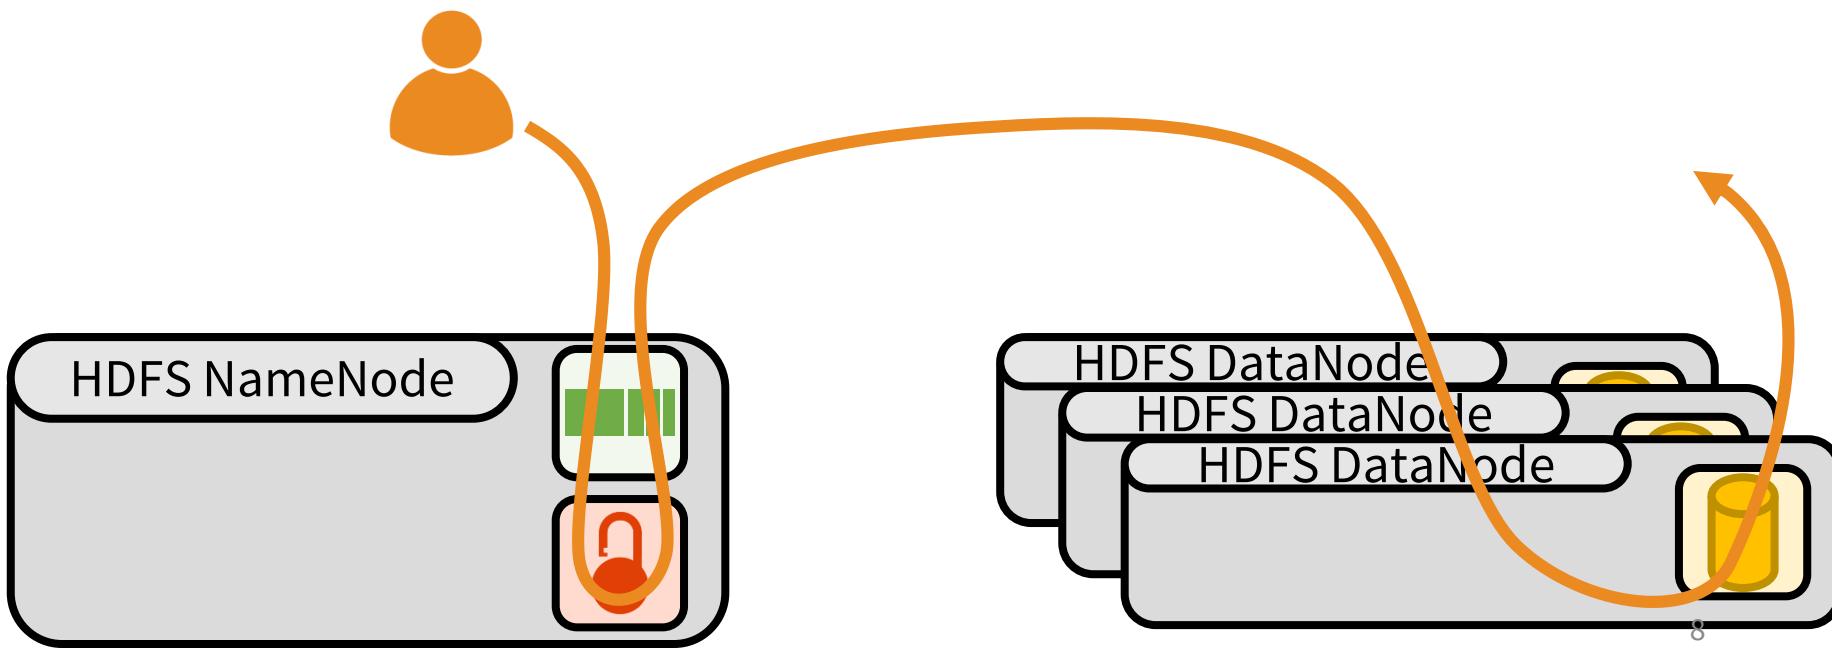


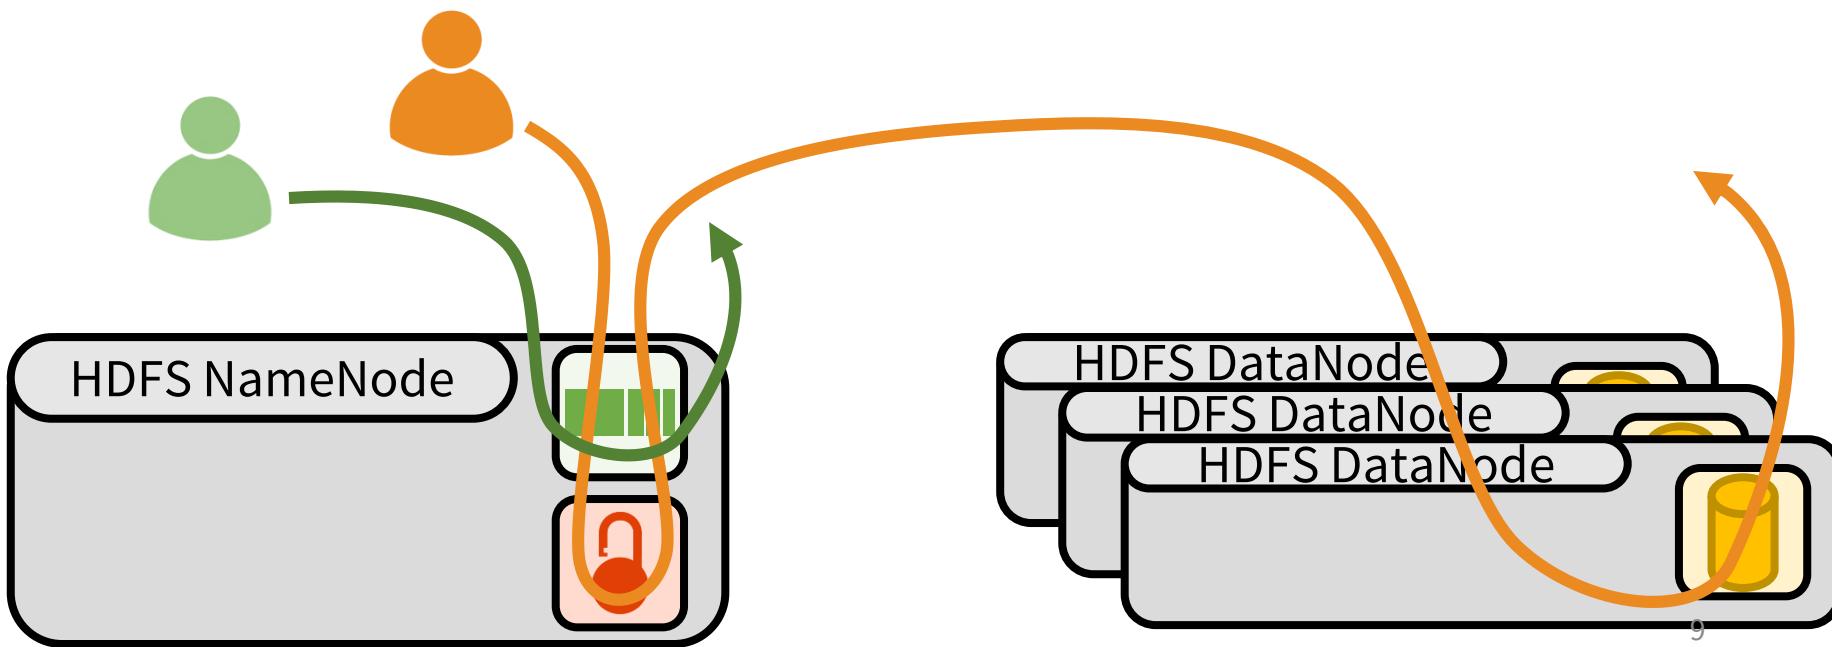


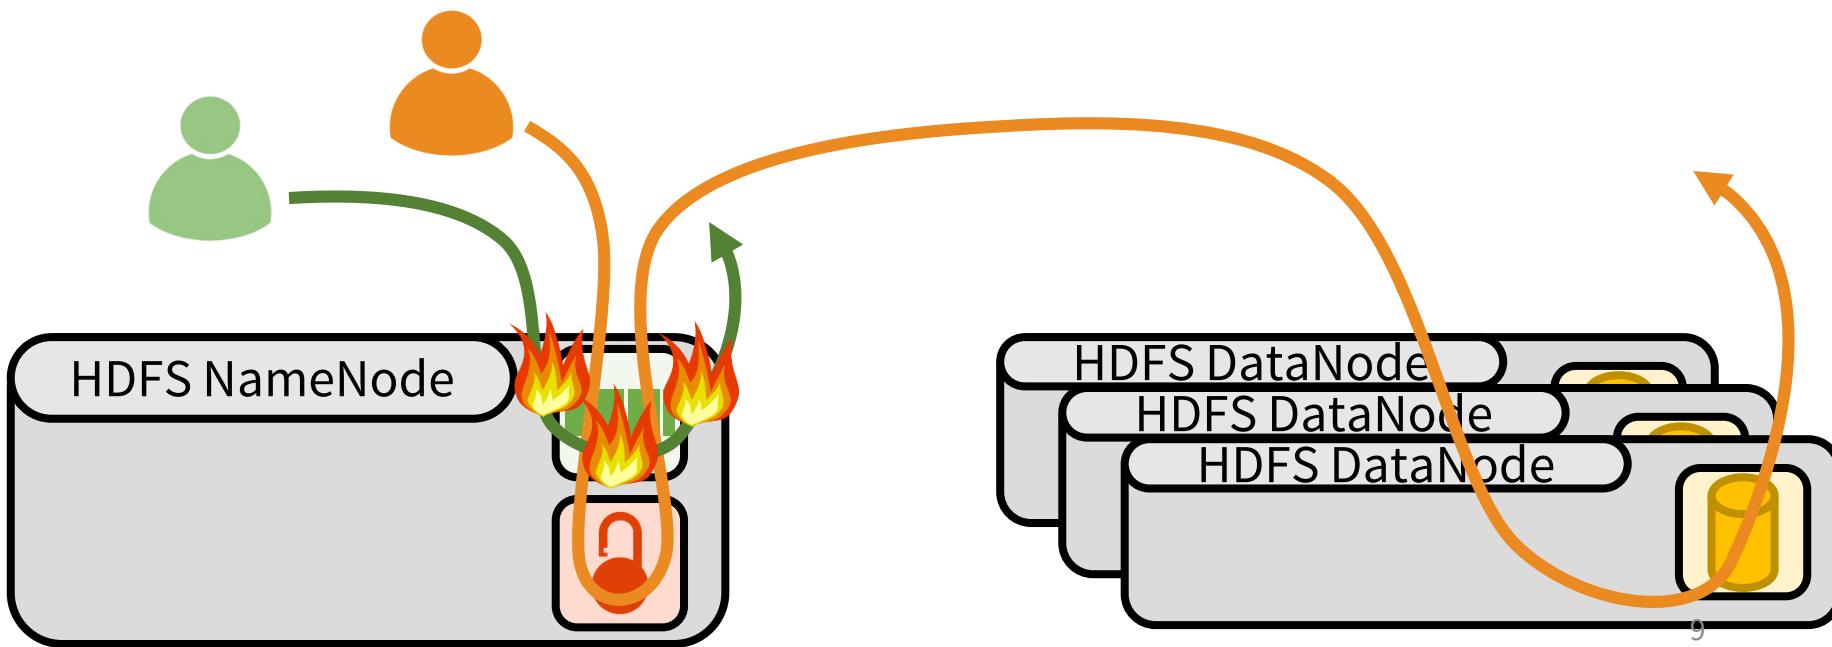
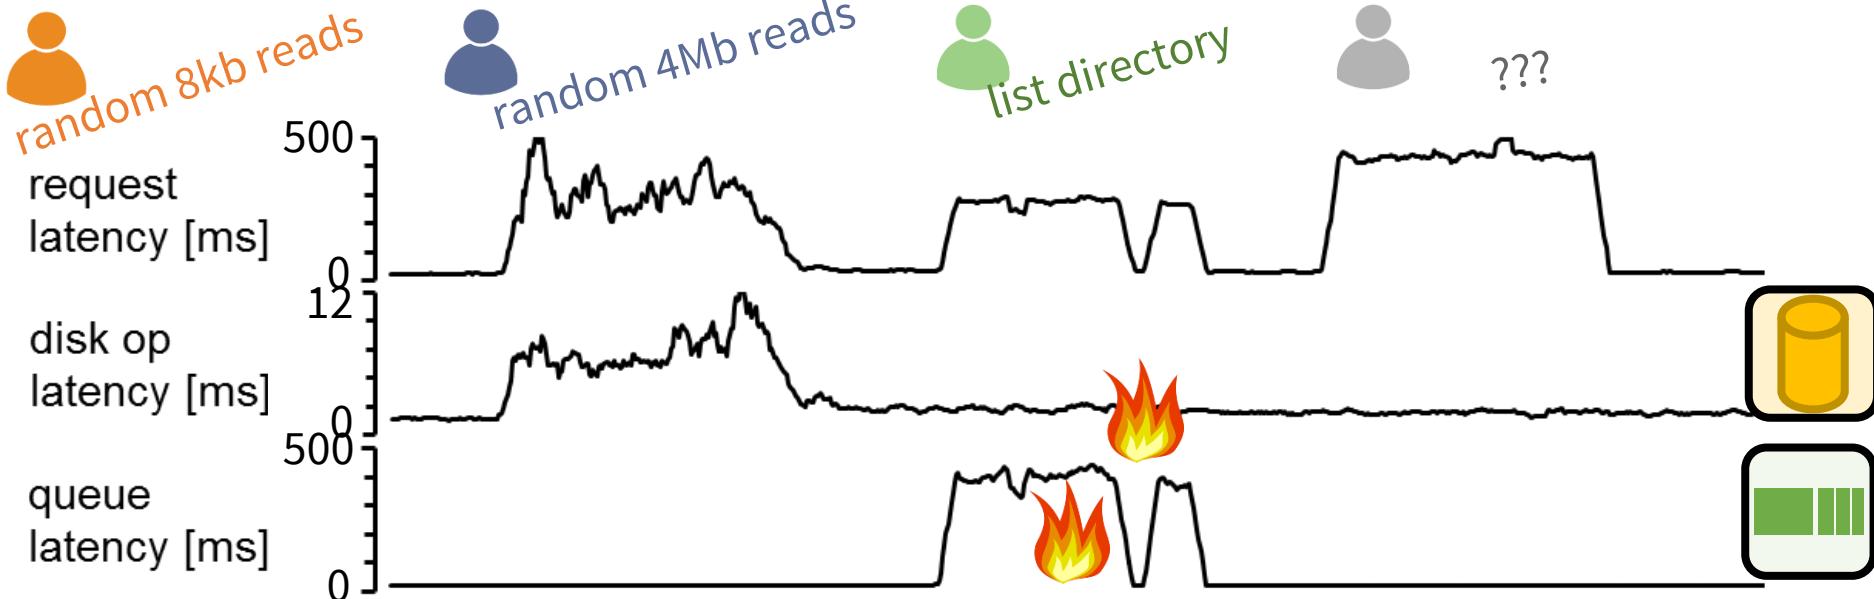


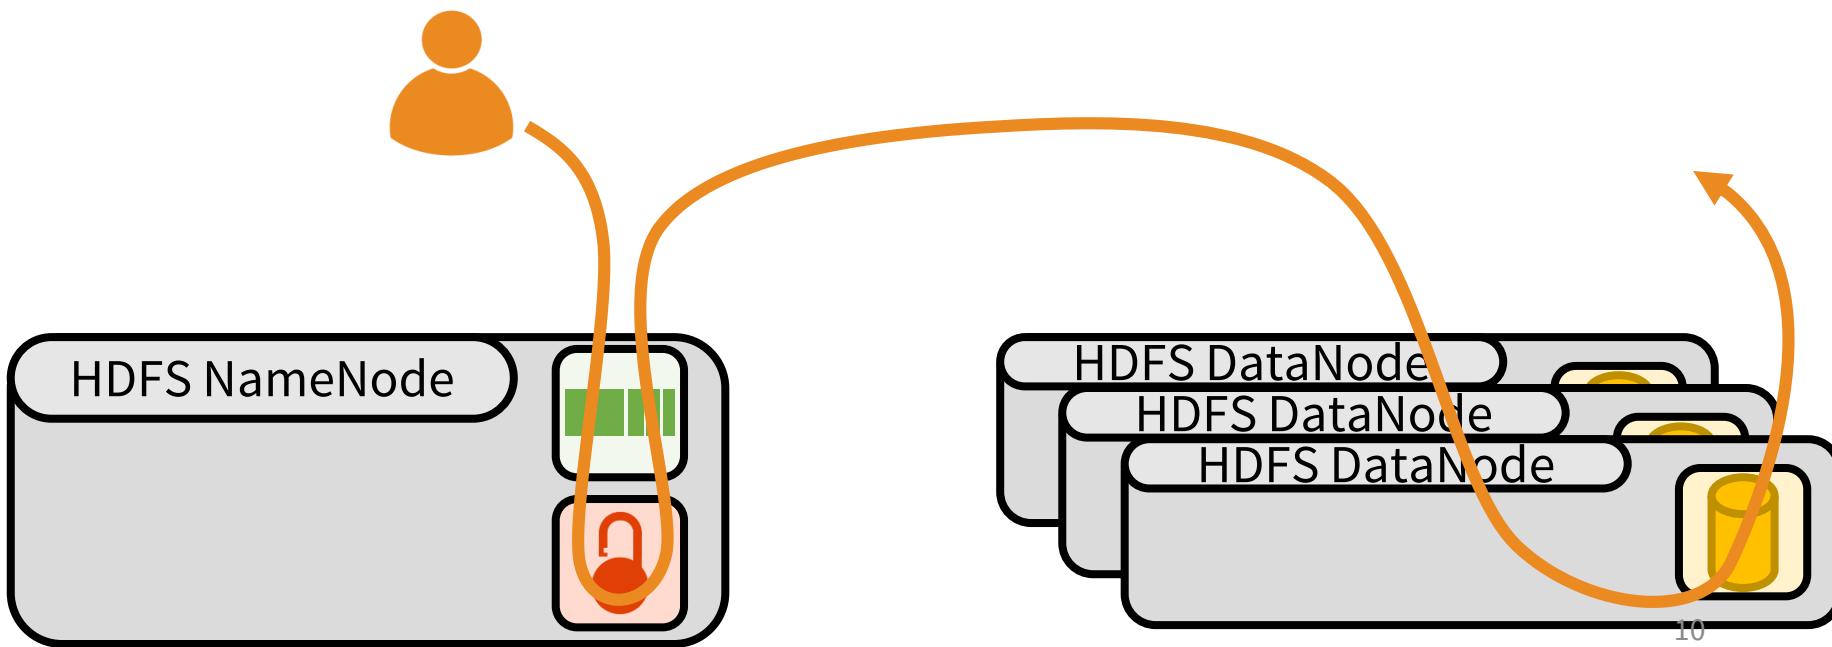
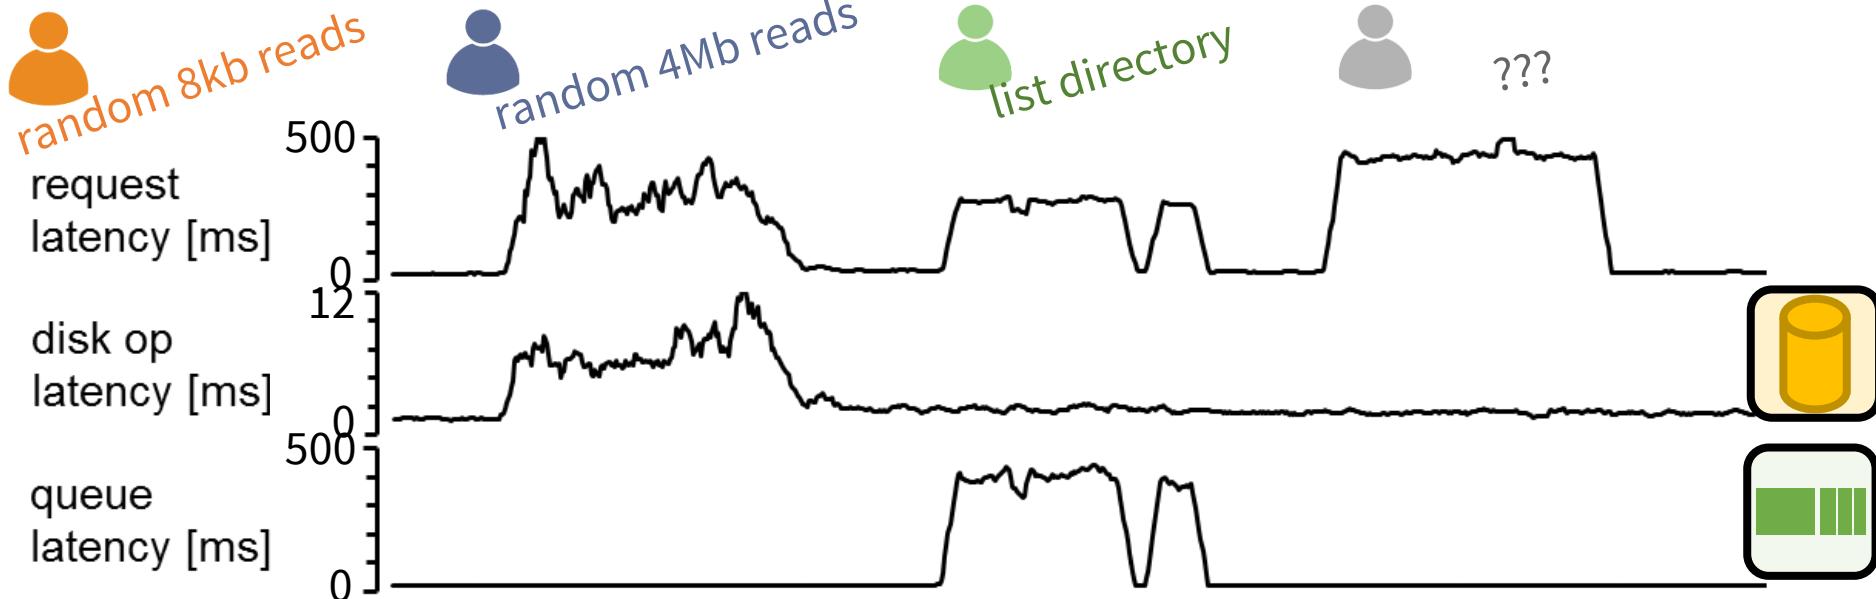


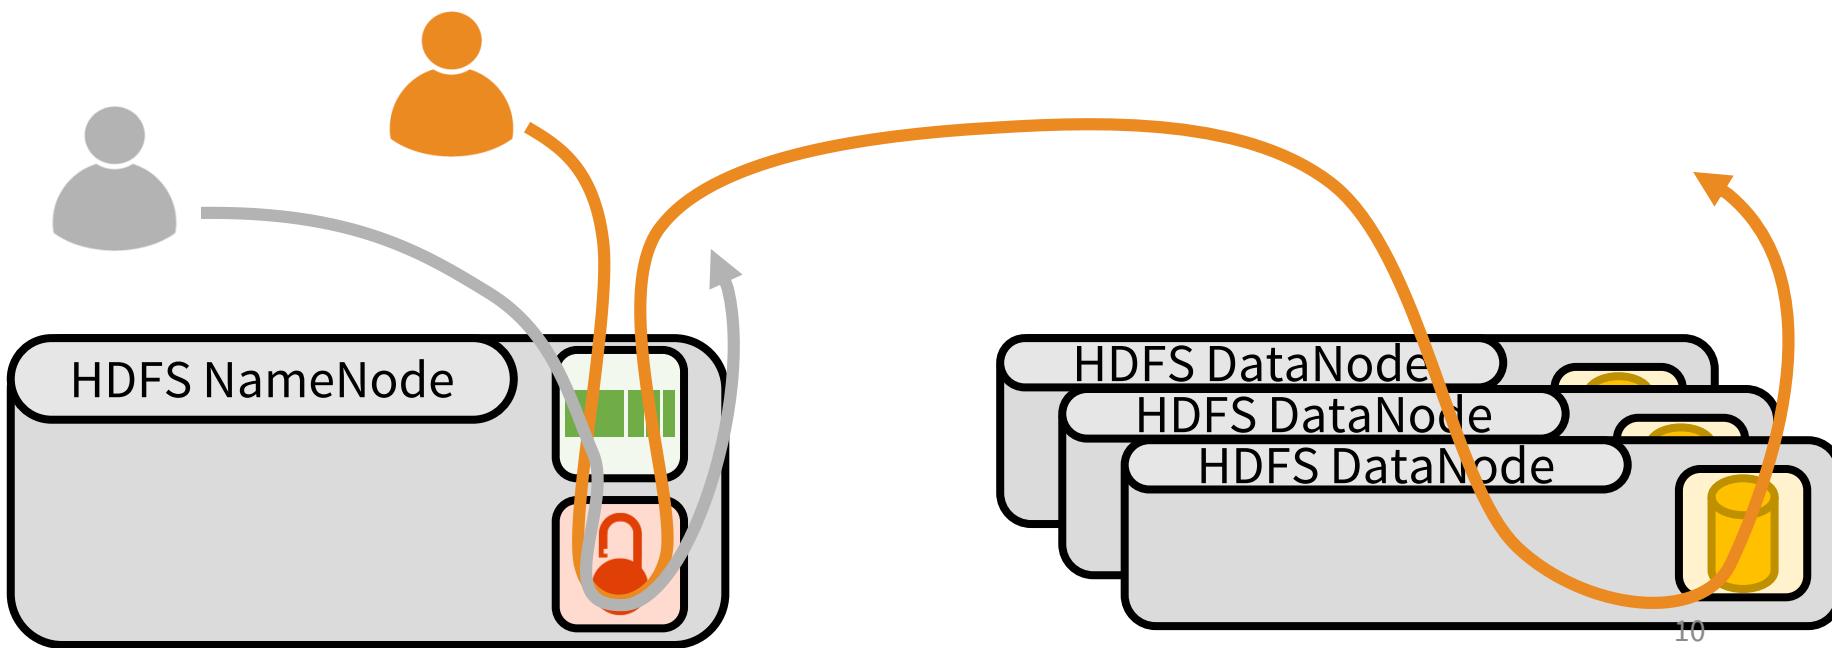
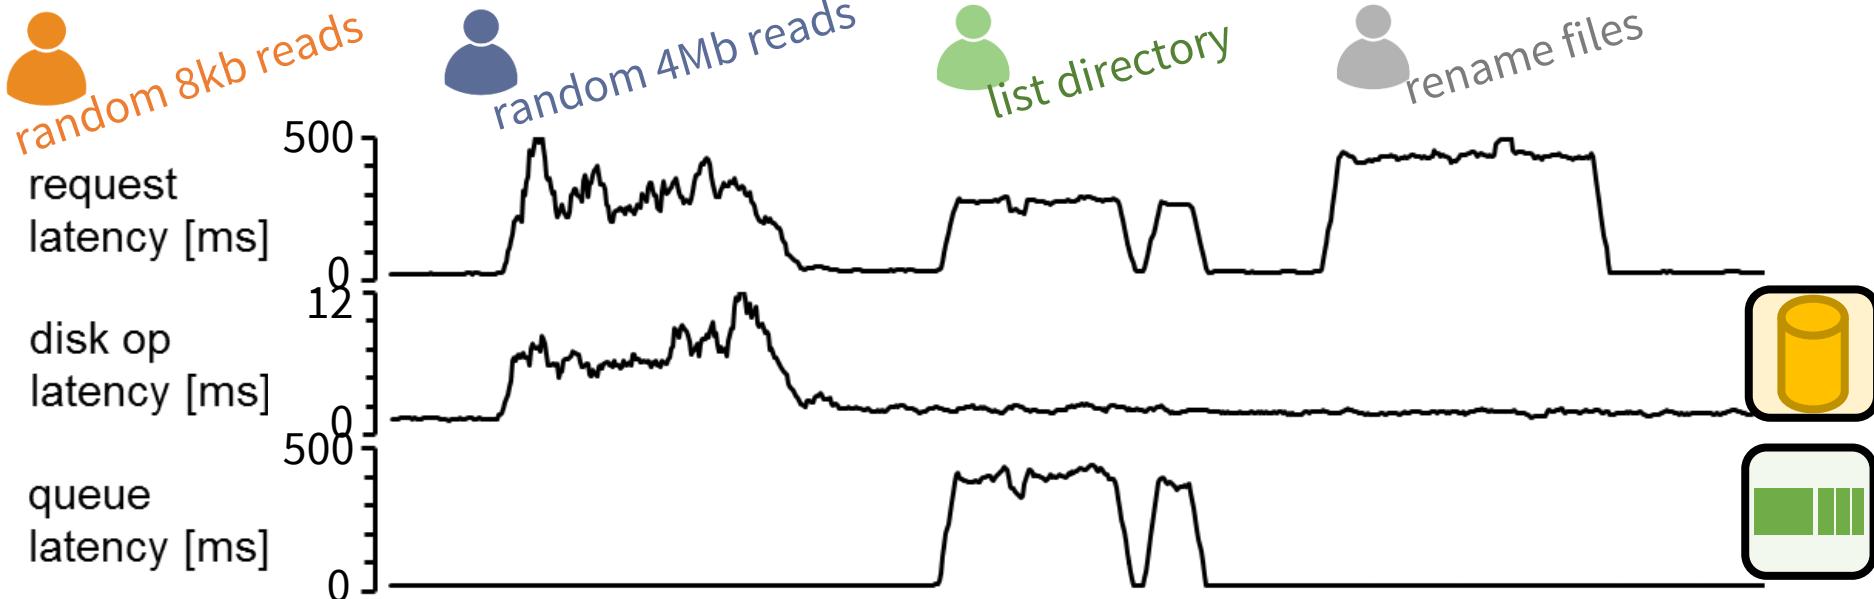


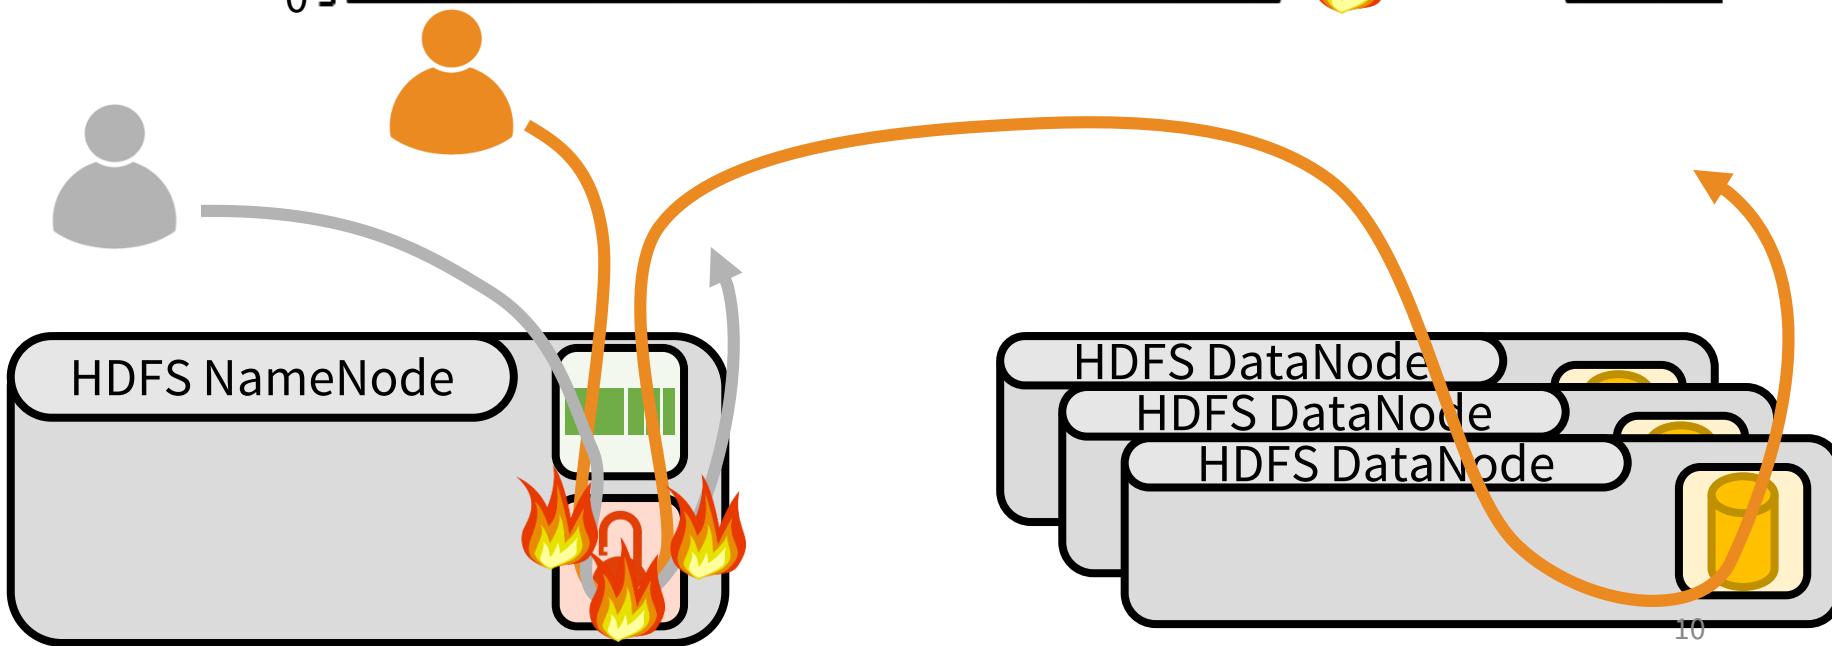
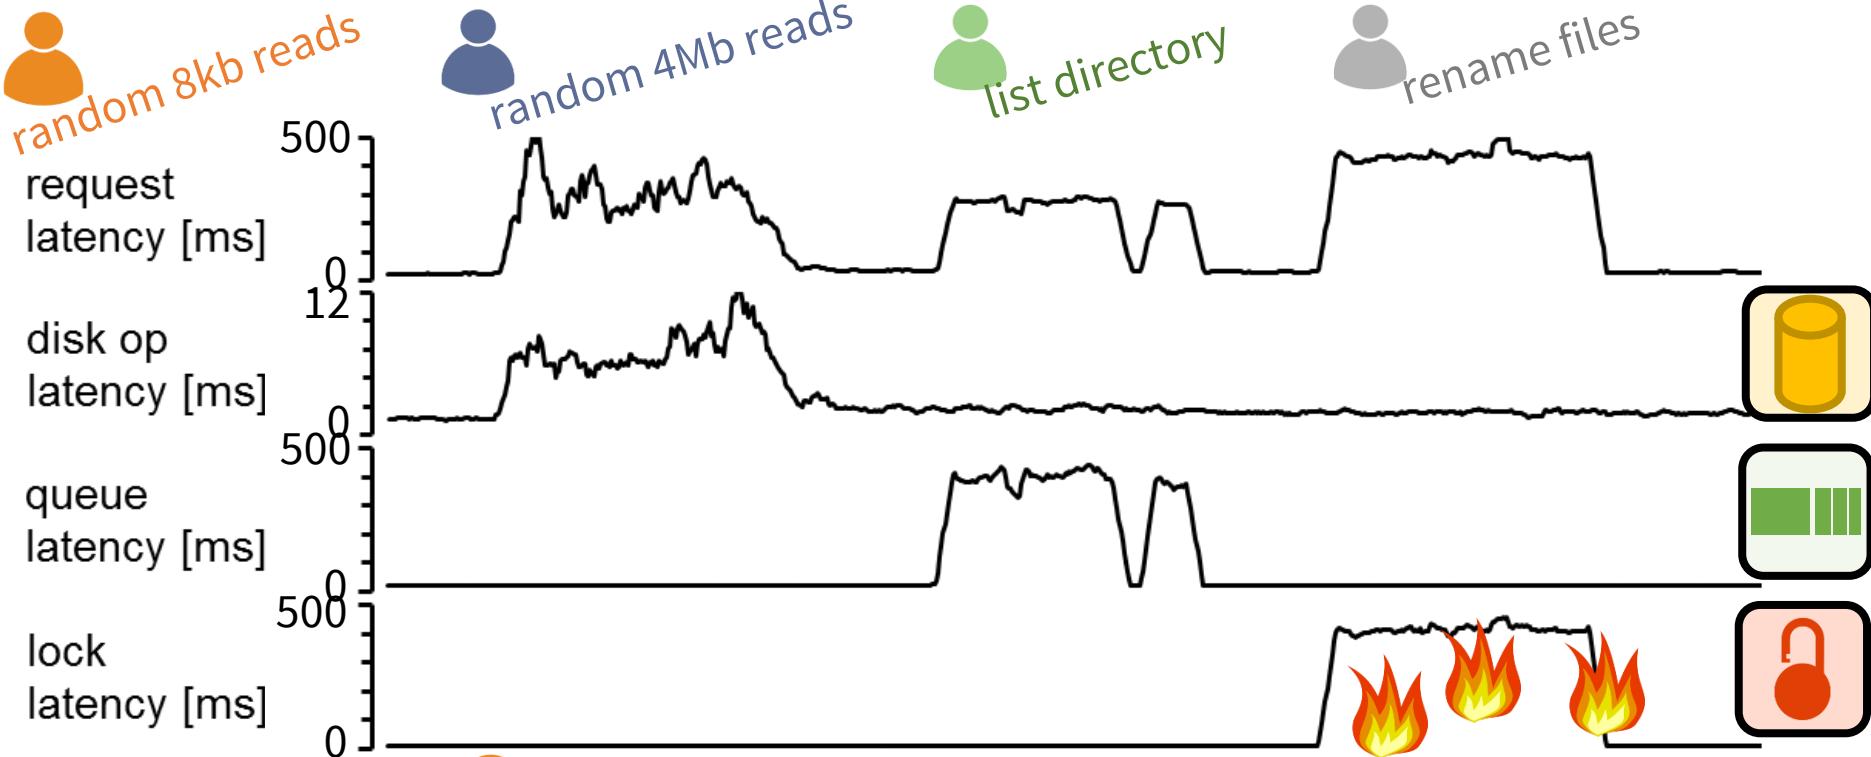


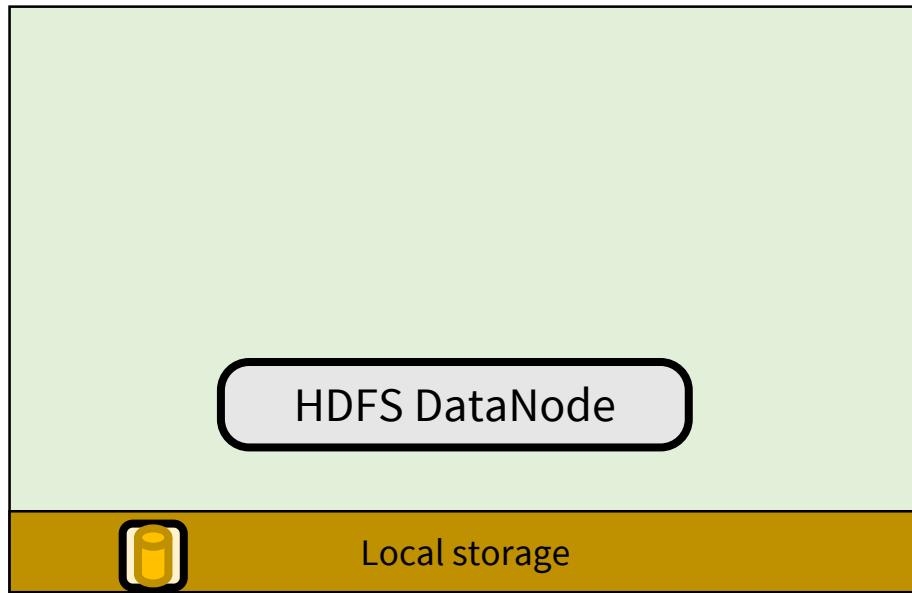


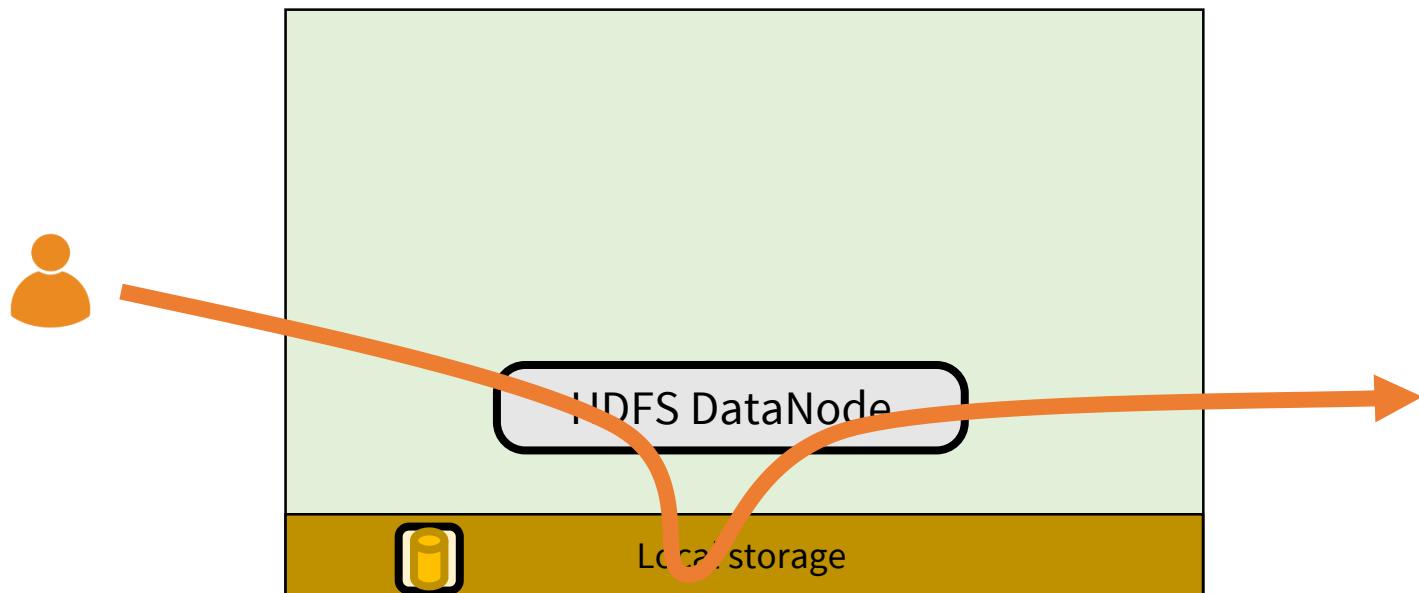


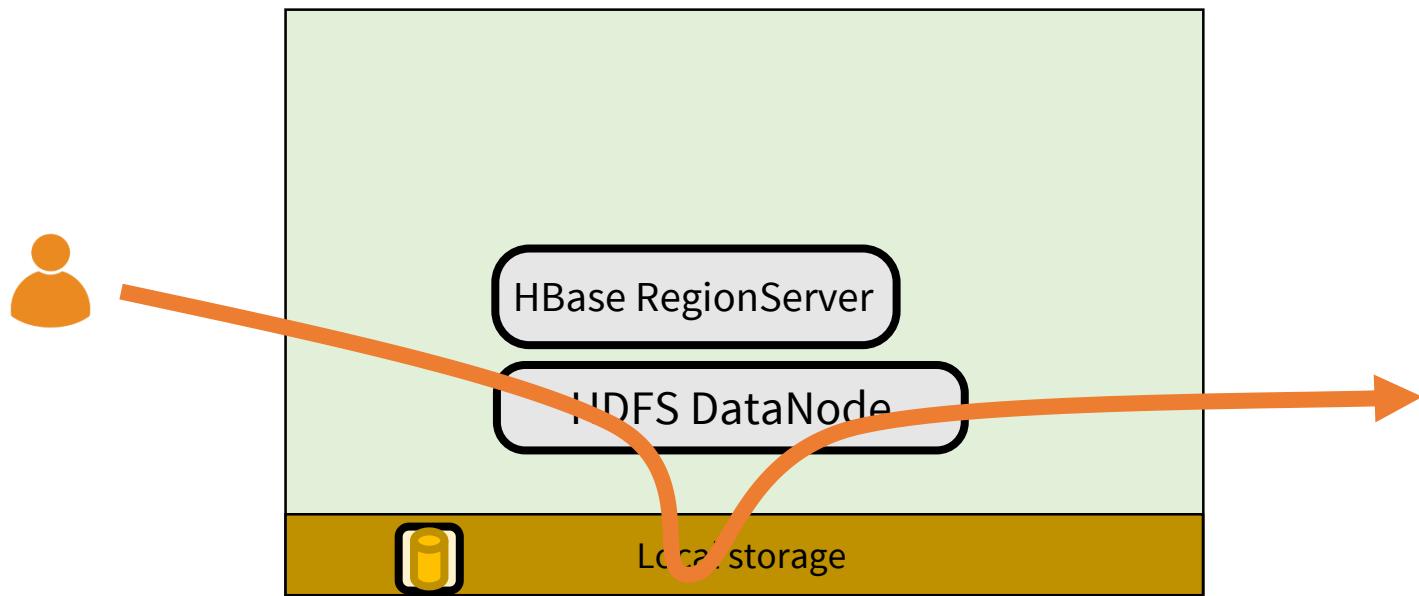


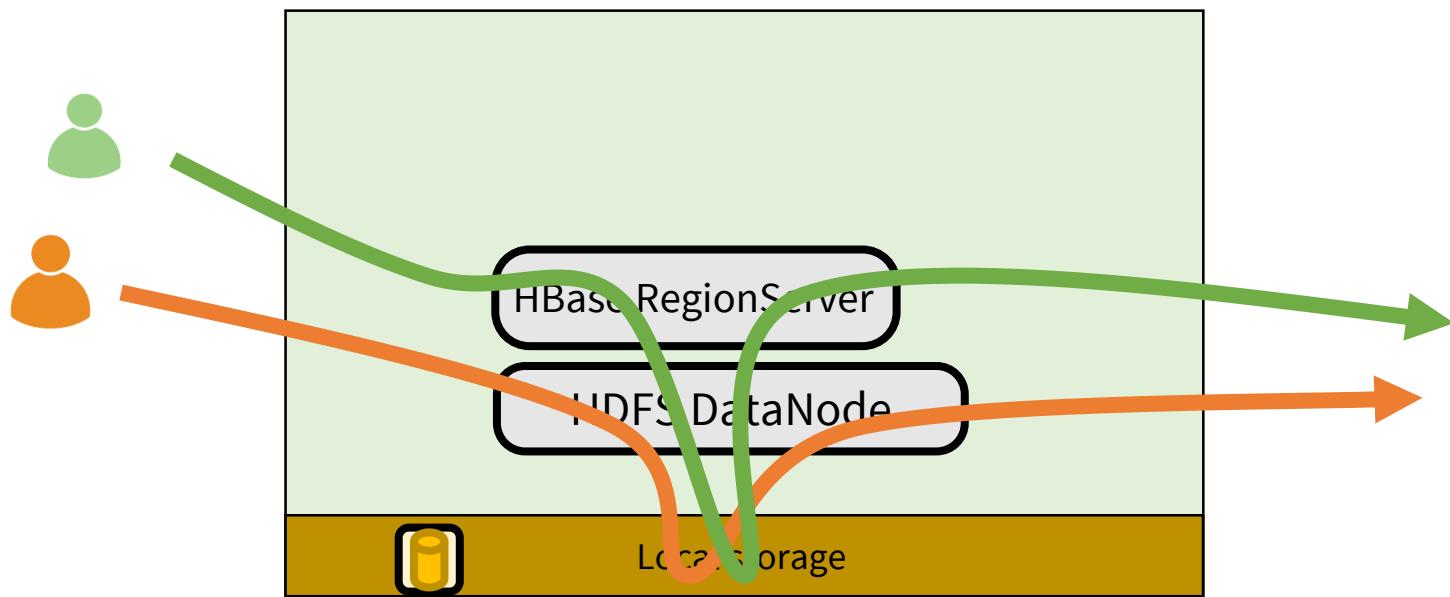


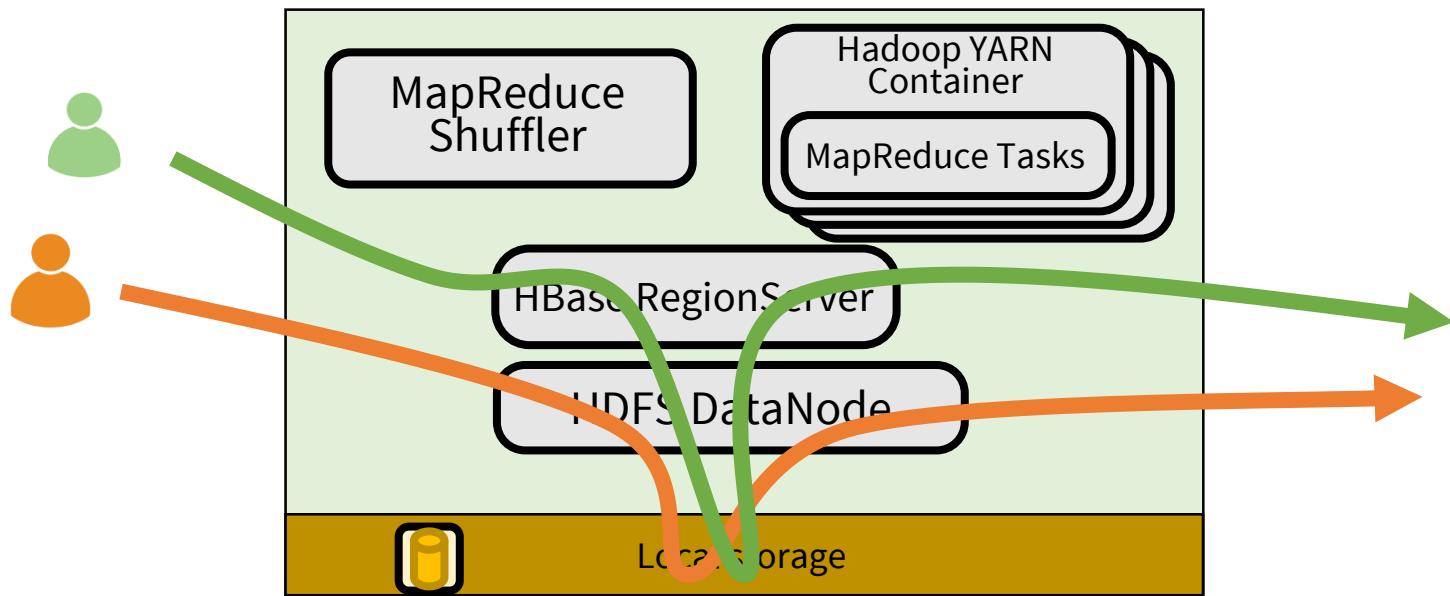


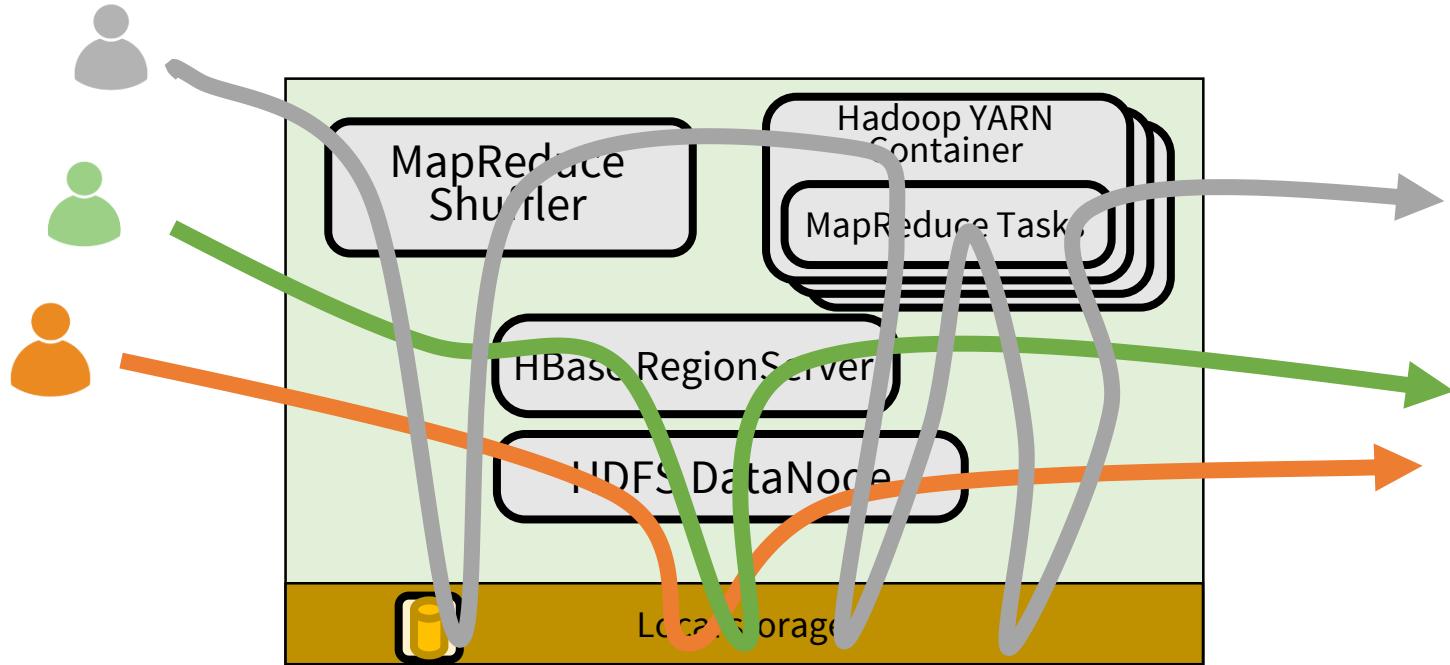


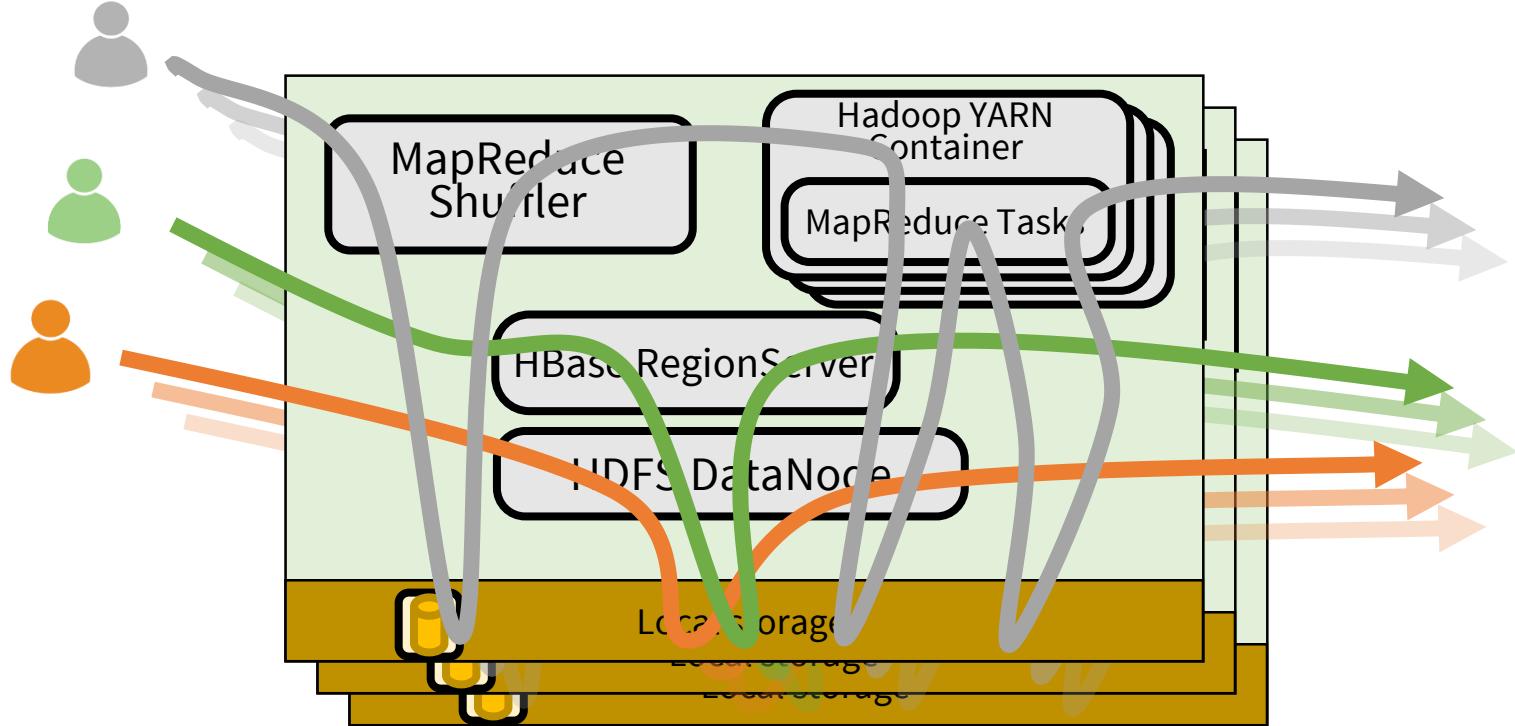












# *Retro* **Goals**

# *Retro* **Goals**

Co-ordinated control across processes, machines, and services

# *Retro* **Goals**

Co-ordinated control across processes, machines, and services

Handle system and application level resources

# *Retro* **Goals**

Co-ordinated control across processes, machines, and services

Handle system and application level resources

Principals: tenants, background tasks

# *Retro* **Goals**

Co-ordinated control across processes, machines, and services

Handle system and application level resources

Principals: tenants, background tasks

Real-time and reactive

# *Retro* **Goals**

Co-ordinated control across processes, machines, and services

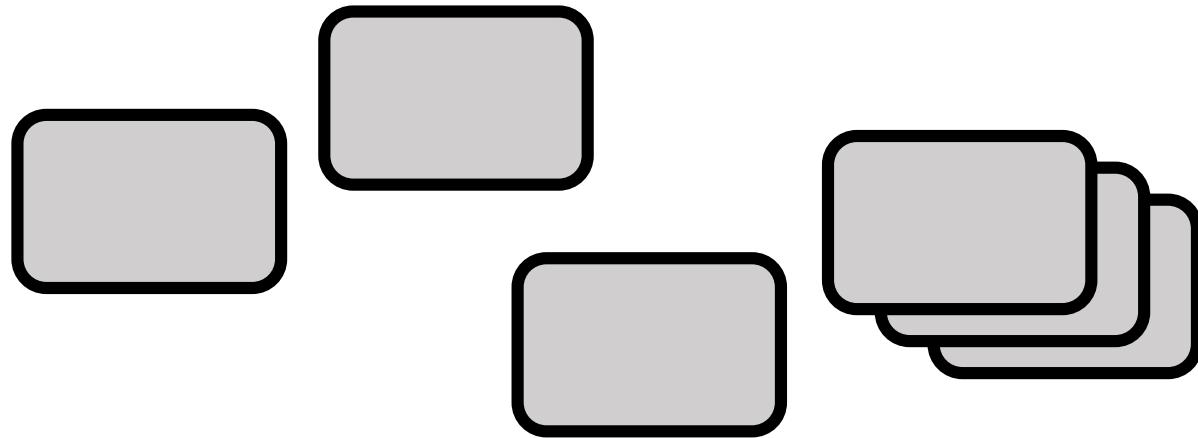
Handle system and application level resources

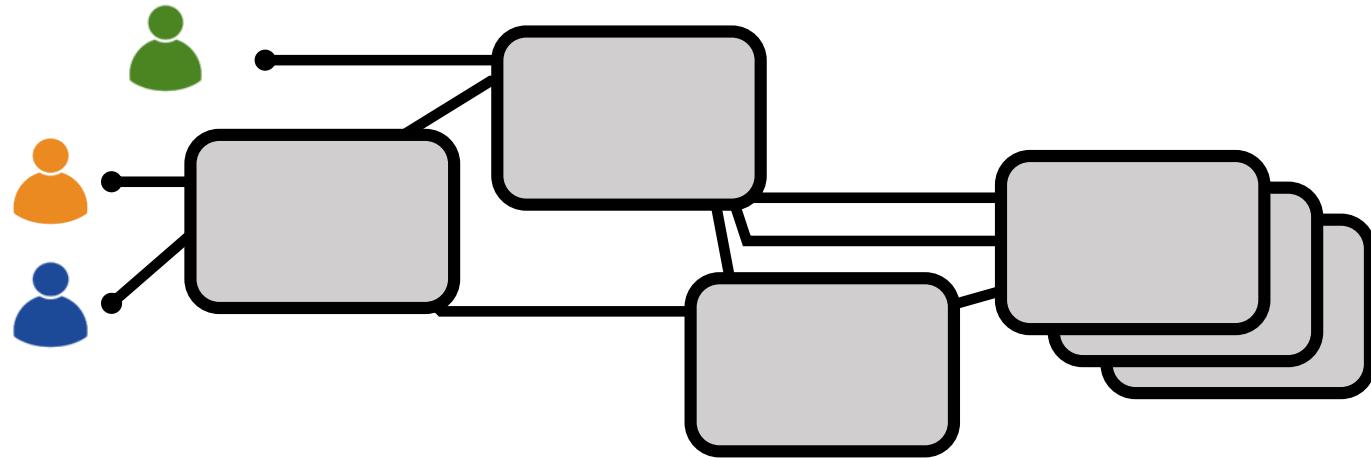
Principals: tenants, background tasks

Real-time and reactive

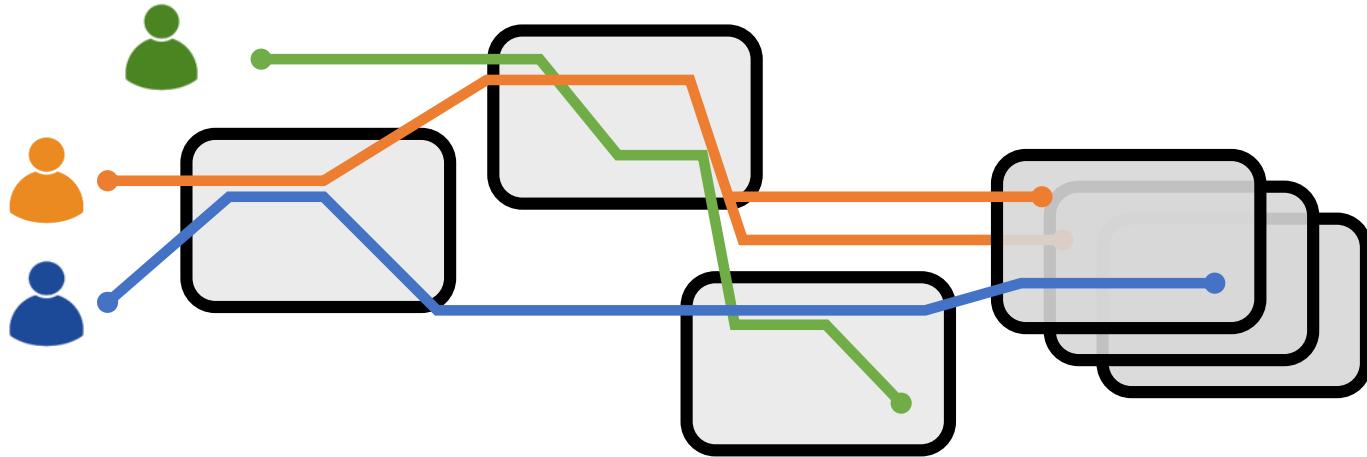
Efficient: Only control what is needed

# *Retro* **Architecture**





 Tenant Requests

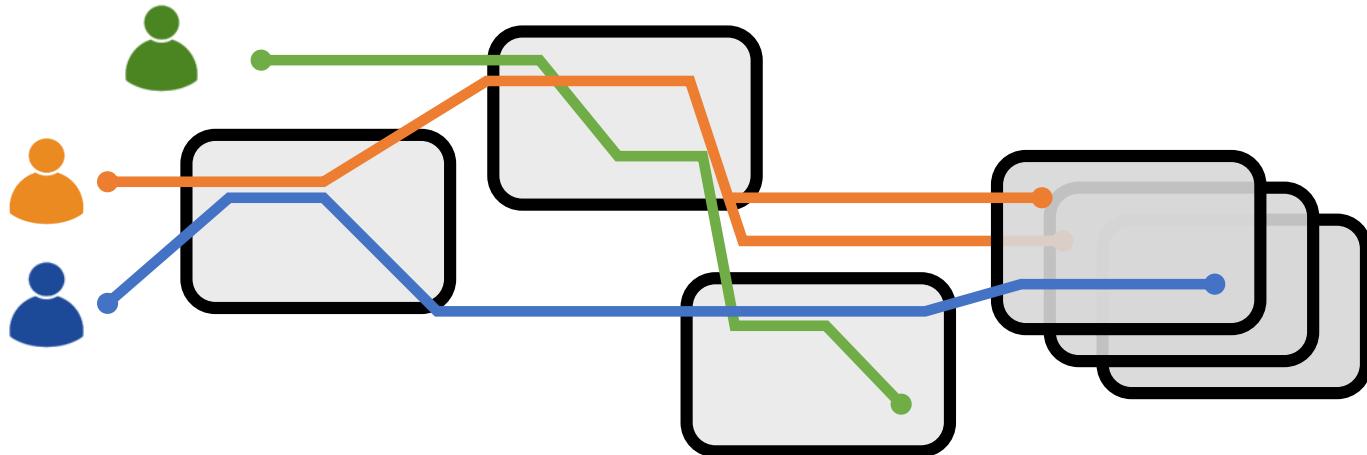


### Workflows

Purpose: identify requests from different users, background activities

eg, all requests from a tenant over time

eg, data balancing in HDFS



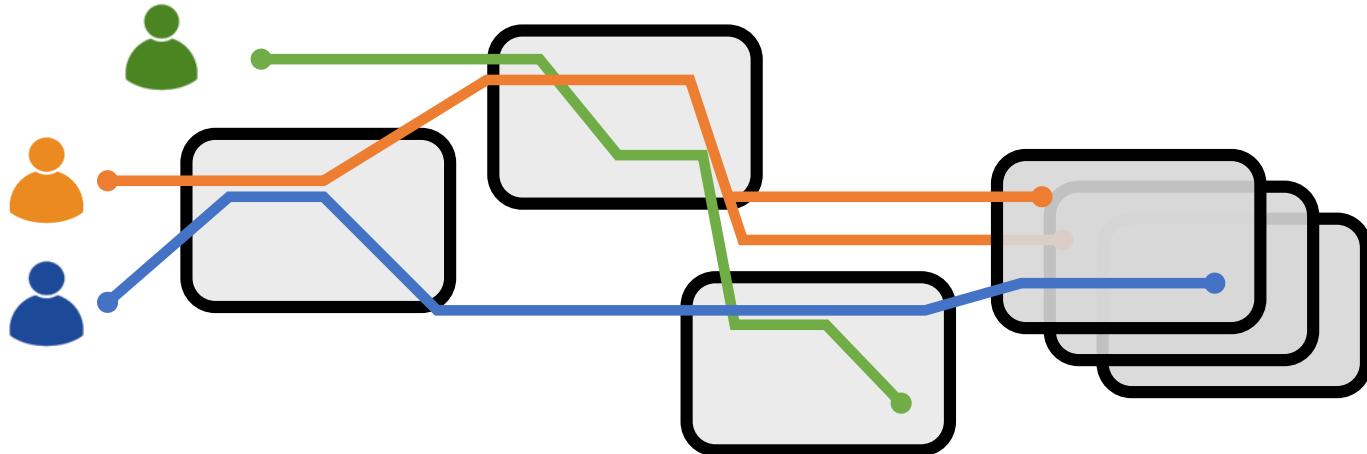
### Workflows

Purpose: identify requests from different users, background activities

eg, all requests from a tenant over time

eg, data balancing in HDFS

Unit of resource measurement, attribution, and enforcement



### Workflows

Purpose: identify requests from different users, background activities

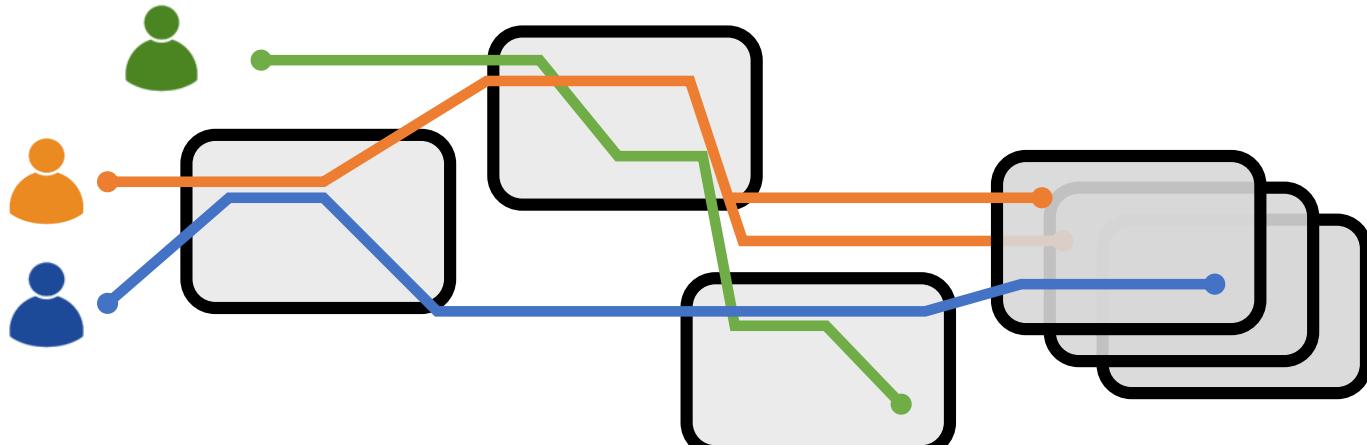
eg, all requests from a tenant over time

eg, data balancing in HDFS

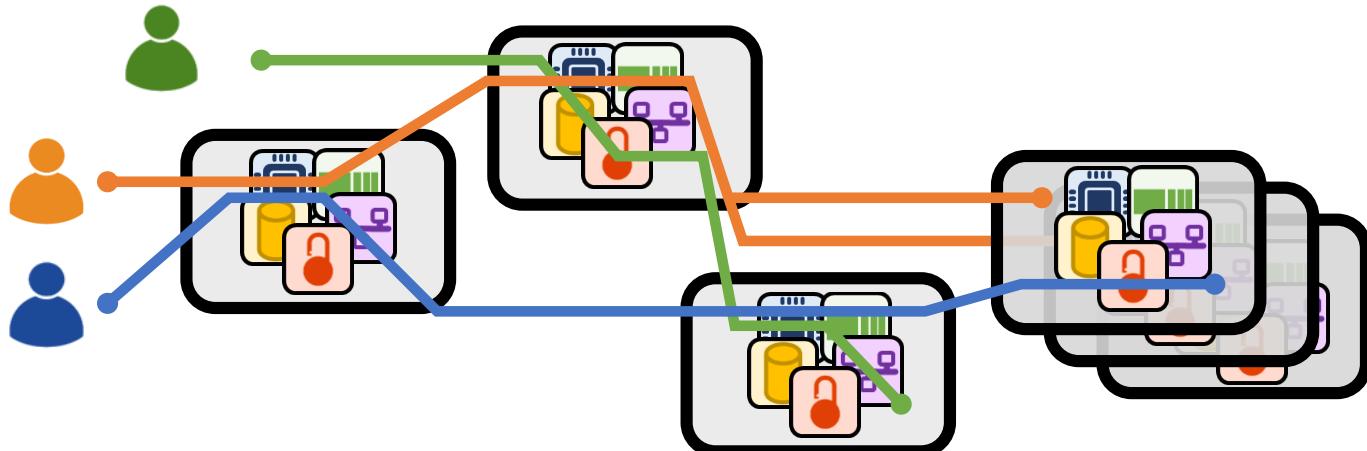
Unit of resource measurement, attribution, and enforcement

Tracks a request across varying levels of granularity

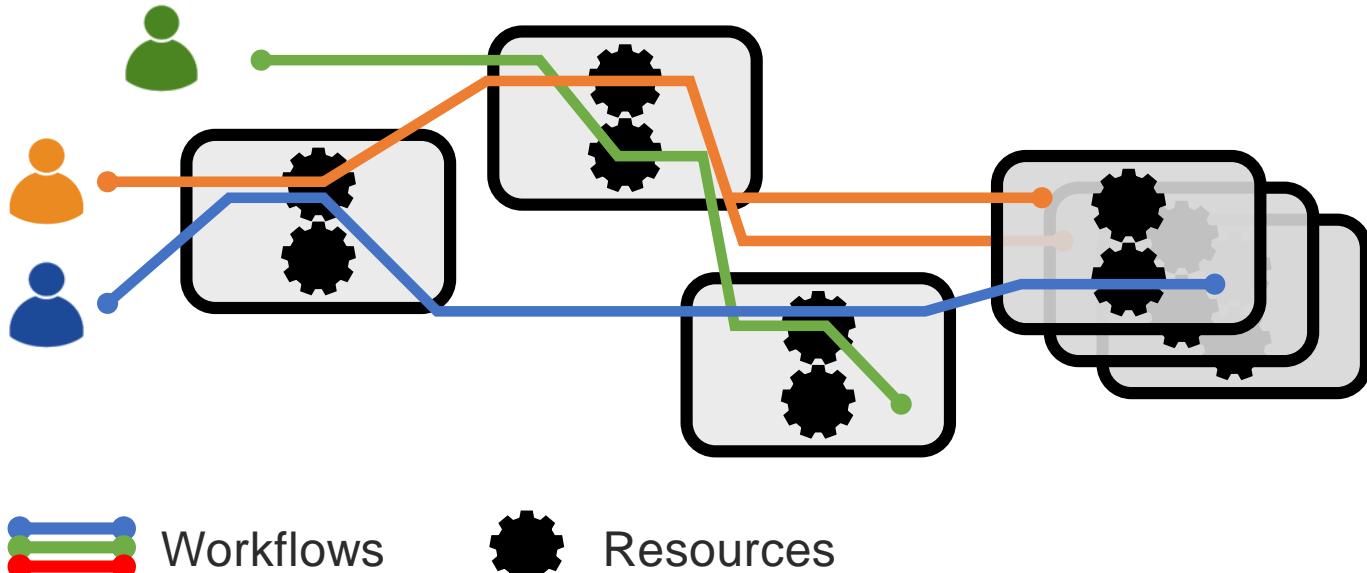
Orthogonal to threads, processes, network flows, etc.



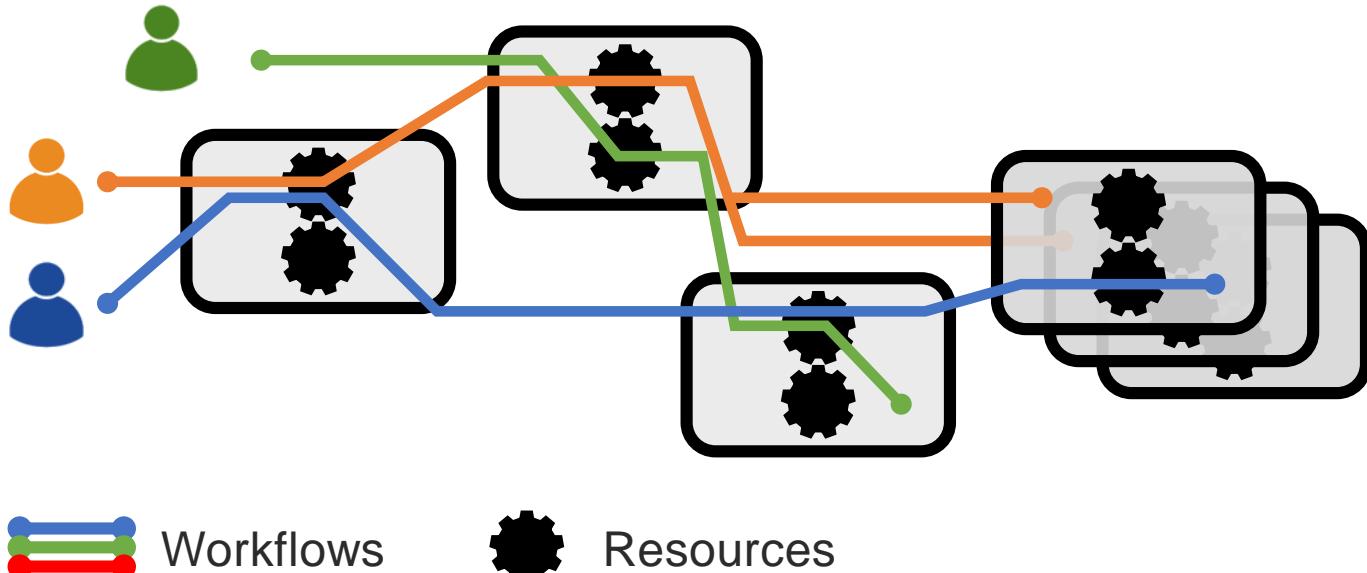
 Workflows



 Workflows



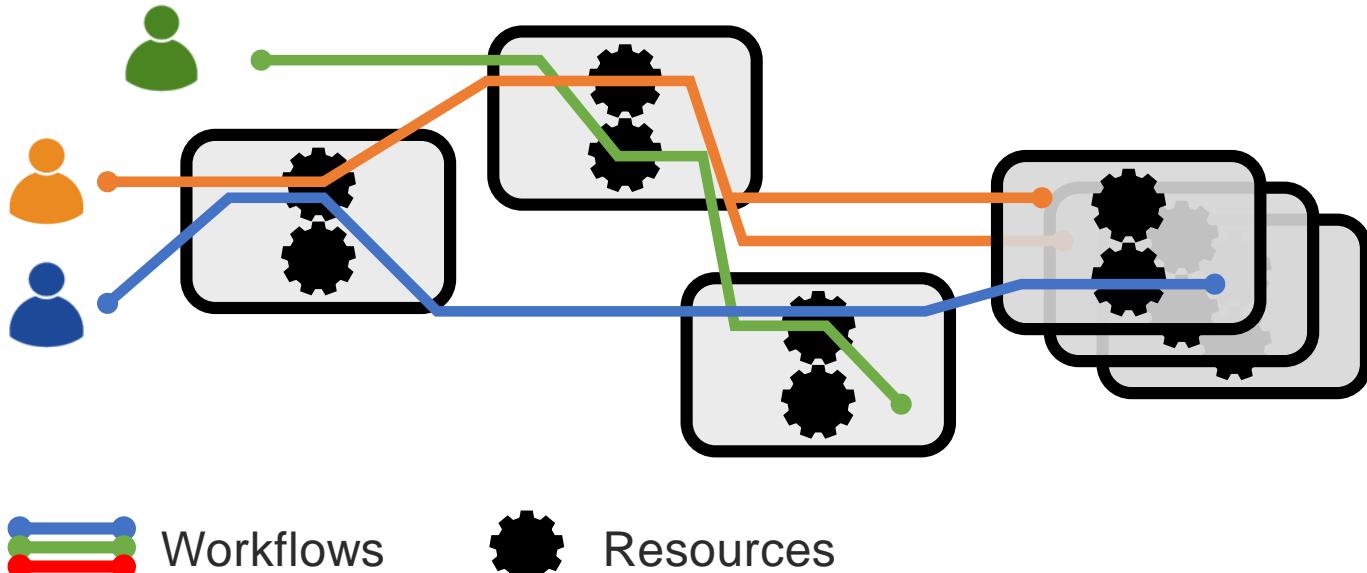
Purpose: cope with diversity of resources



Purpose: cope with diversity of resources

What we need:

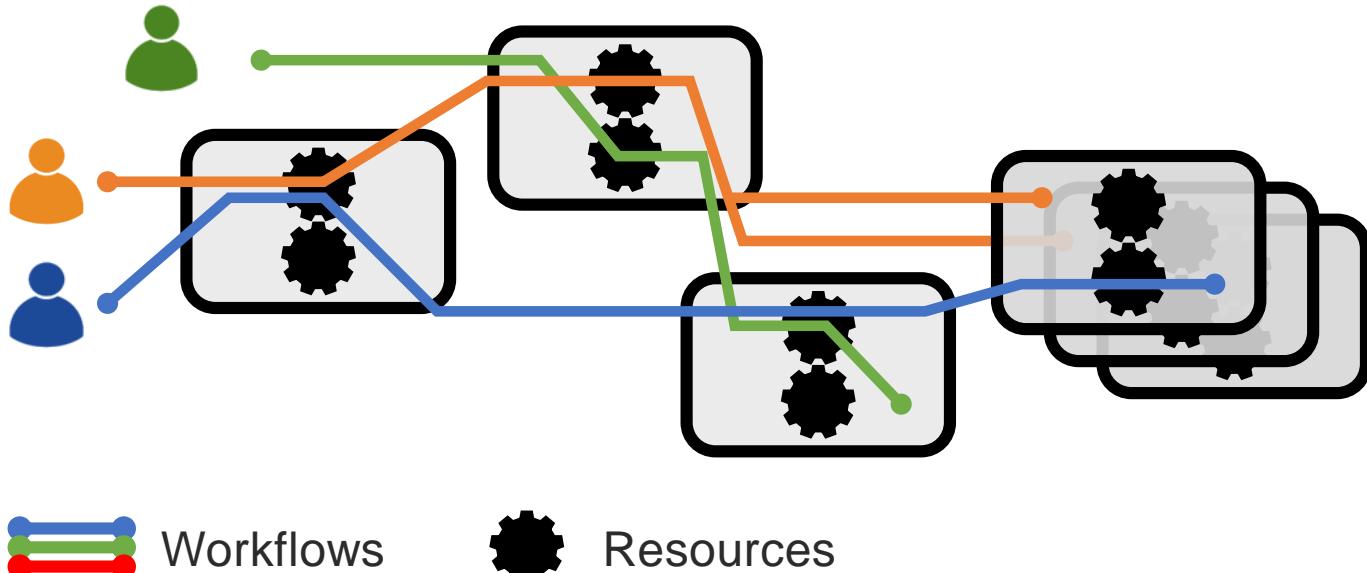
1. Identify overloaded resources



Purpose: cope with diversity of resources

What we need:

1. Identify overloaded resources
  
2. Identify culprit workflows



Purpose: cope with diversity of resources

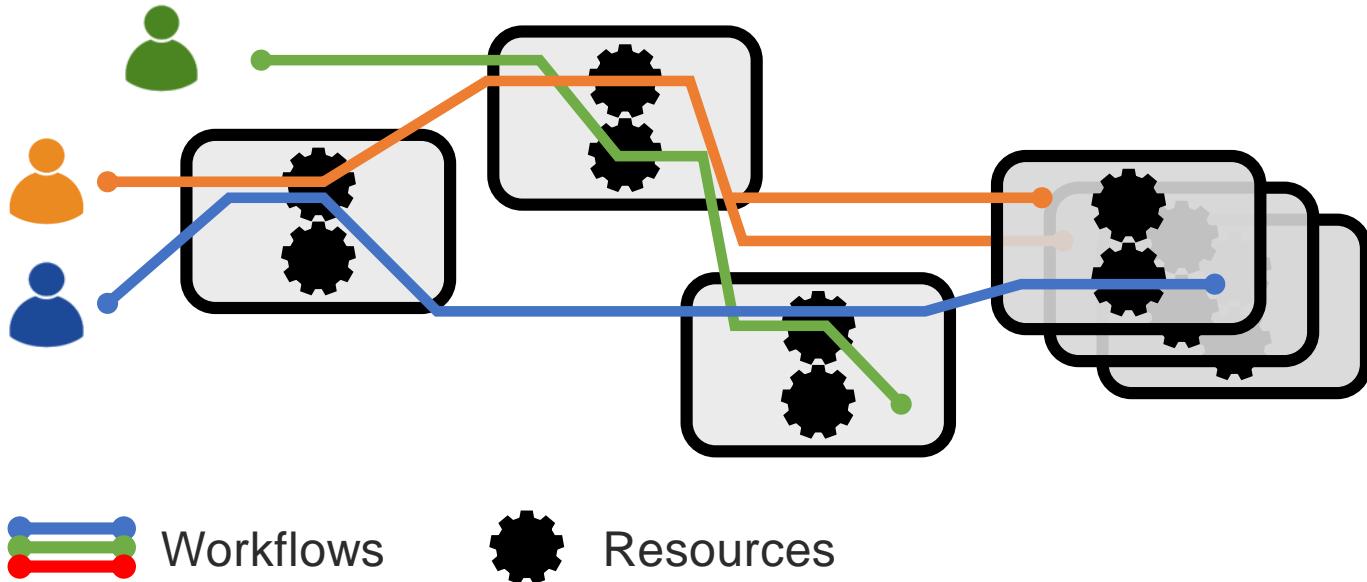
What we need:

1. Identify overloaded resources

**Slowdown**

Ratio of how slow the resource is now compared to its baseline performance with no contention.

2. Identify culprit workflows



Purpose: cope with diversity of resources

What we need:

1. Identify overloaded resources

**Slowdown**

Ratio of how slow the resource is now compared to its baseline performance with no contention.

2. Identify culprit workflows

**Load**

Fraction of current utilization that we can attribute to each workflow



Resources

Purpose: cope with diversity of resources

What we need:

1. Identify overloaded resources

**Slowdown**

Ratio of how slow the resource is now compared to its baseline performance with no contention.

2. Identify culprit workflows

**Load**

Fraction of current utilization that we can attribute to each workflow



Workflows



Resources

Purpose: cope with diversity of resources

What we need:

1. Identify overloaded resources

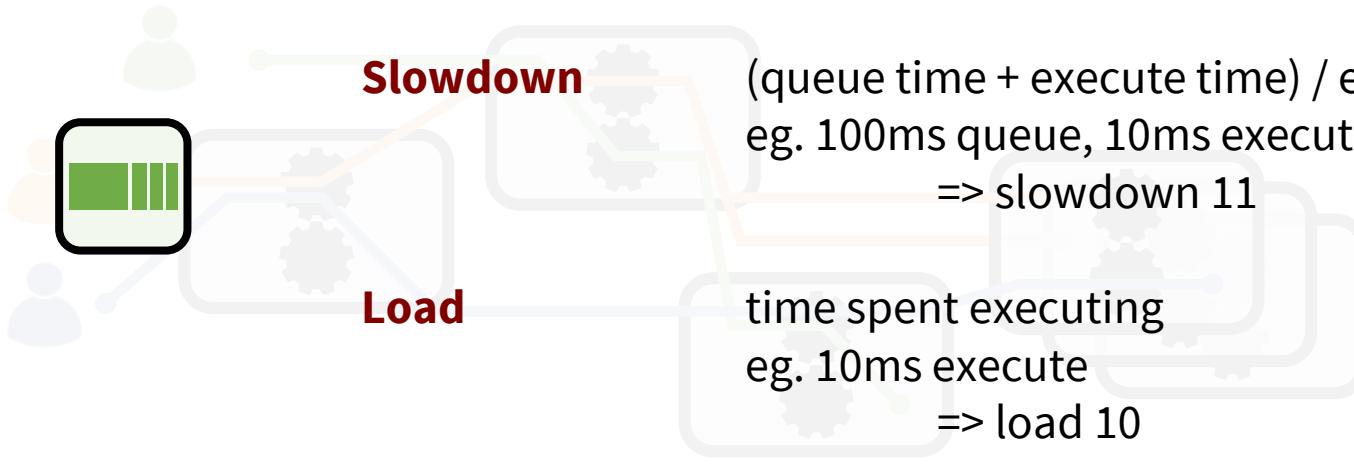
**Slowdown**

Ratio of how slow the resource is now compared to its baseline performance with no contention.

2. Identify culprit workflows

**Load**

Fraction of current utilization that we can attribute to each workflow



Resources

Purpose: cope with diversity of resources

What we need:

1. Identify overloaded resources

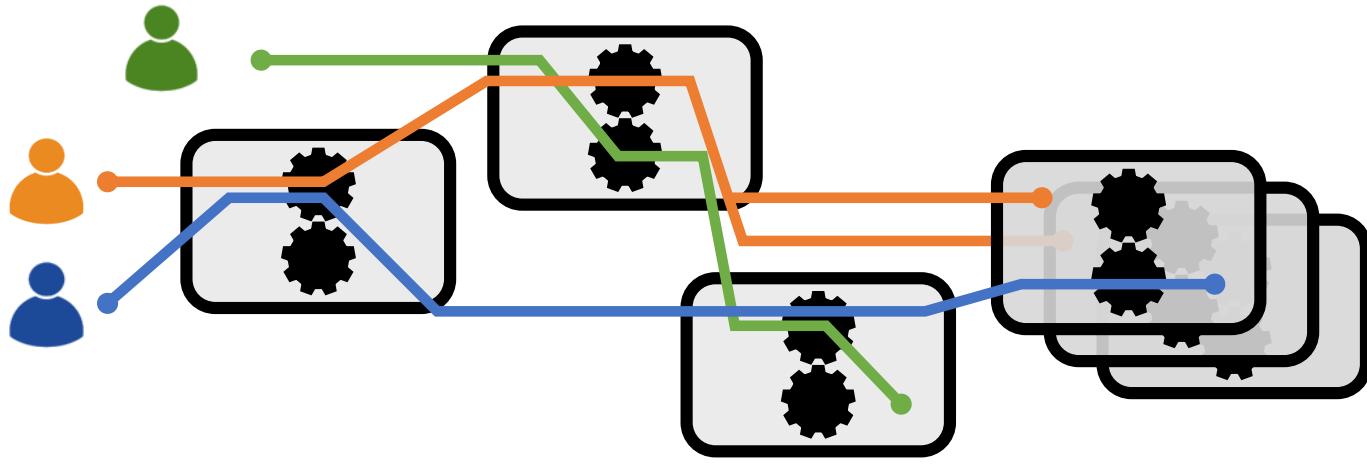
**Slowdown**

Ratio of how slow the resource is now compared to its baseline performance with no contention.

2. Identify culprit workflows

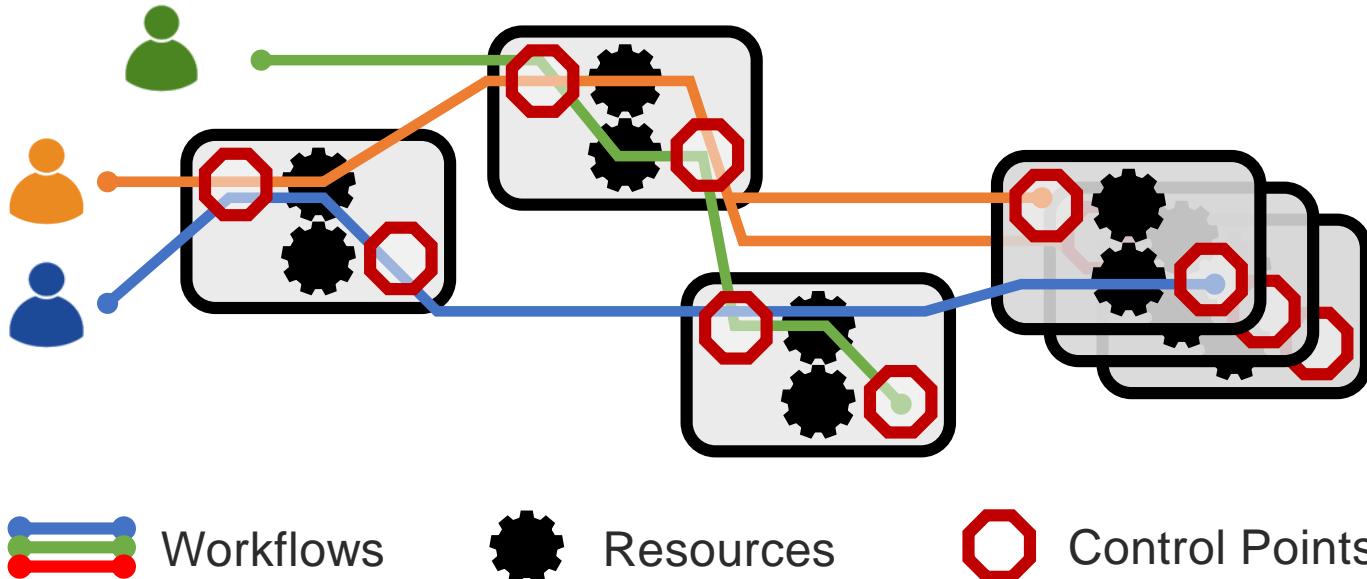
**Load**

Fraction of current utilization that we can attribute to each workflow

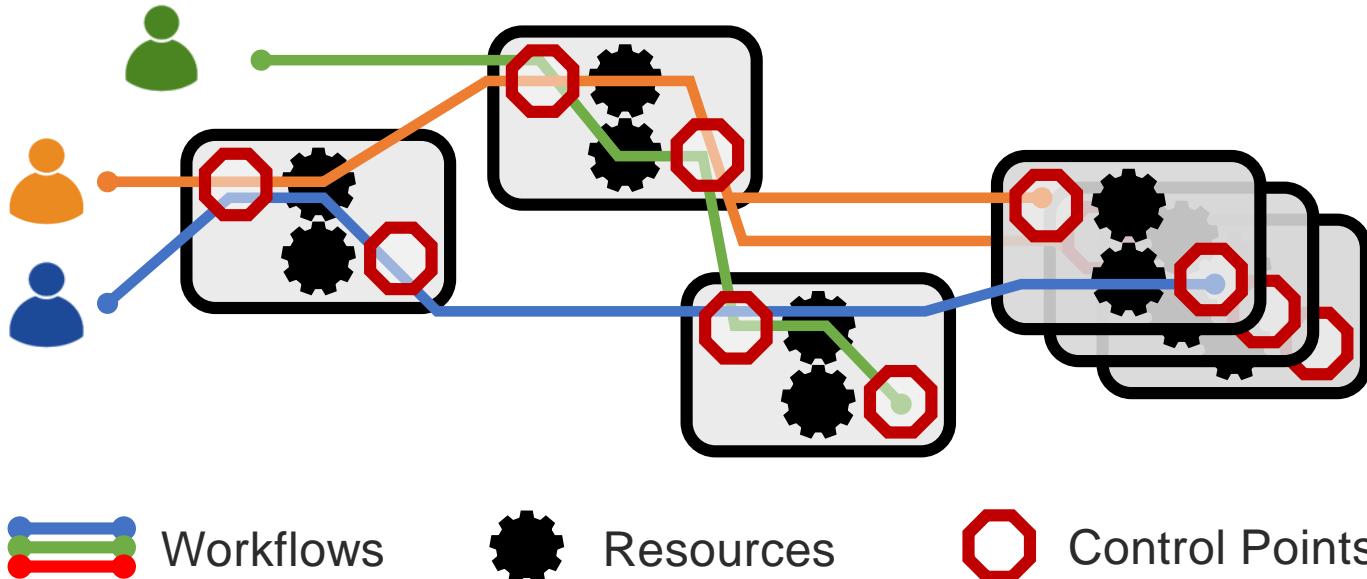


Workflows

Resources



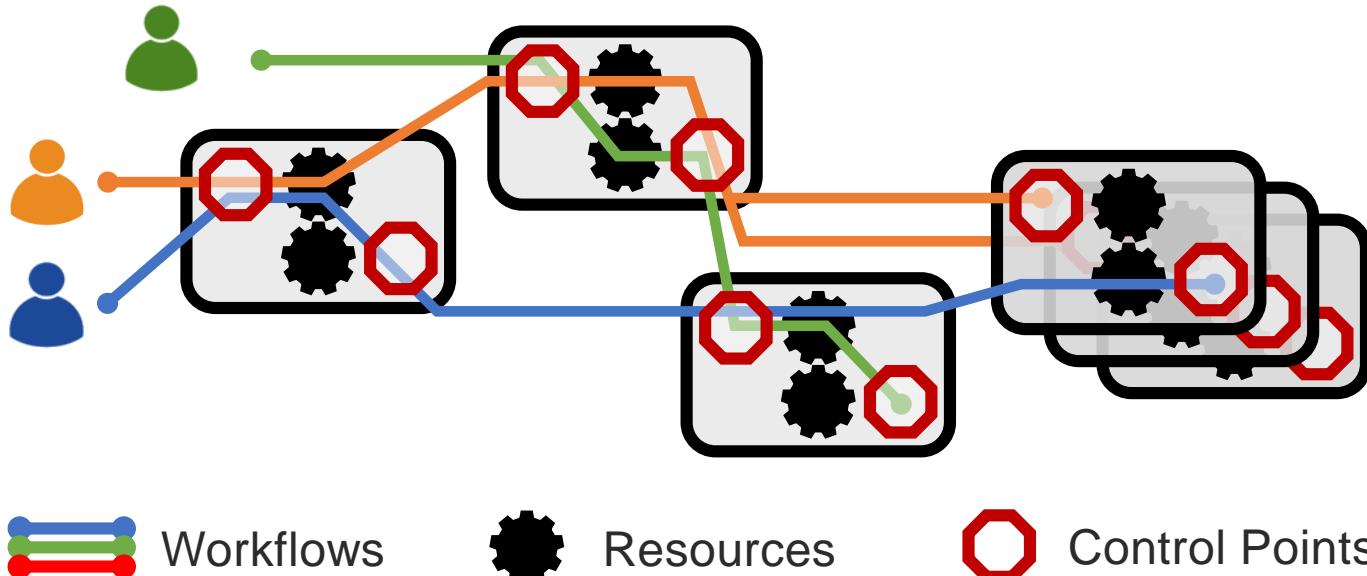
Goal: enforce resource management decisions



Goal: enforce resource management decisions

Decoupled from resources

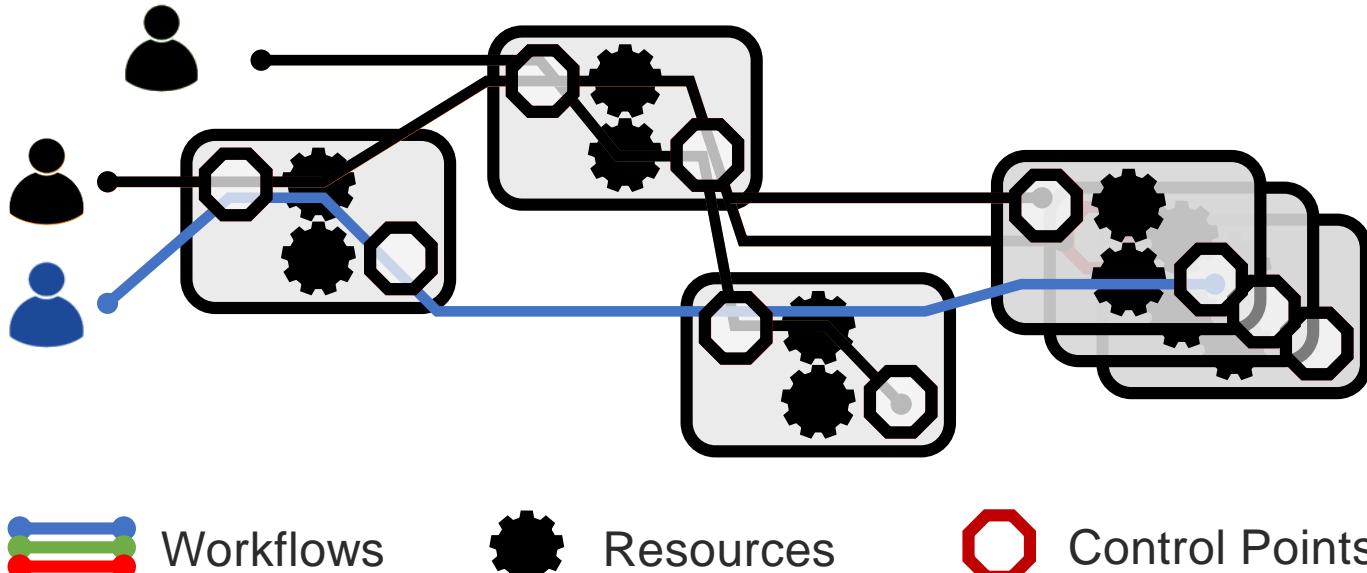
Rate-limits workflows, agnostic to underlying implementation e.g.,  
token bucket  
priority queue



Goal: enforce resource management decisions

Decoupled from resources

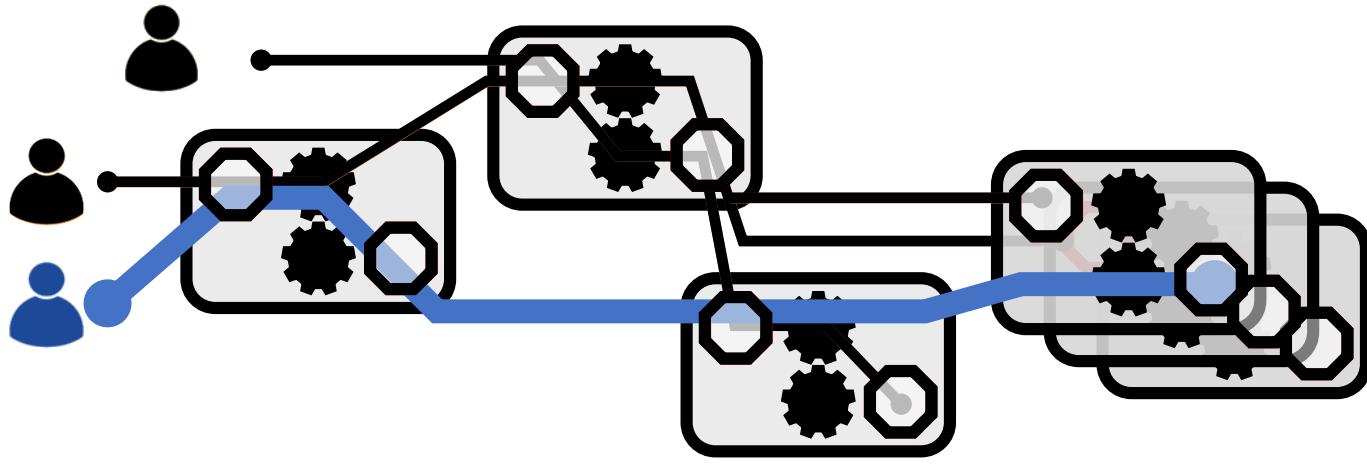
Rate-limits workflows, agnostic to underlying implementation e.g.,  
token bucket  
priority queue



Goal: enforce resource management decisions

Decoupled from resources

Rate-limits workflows, agnostic to underlying implementation e.g.,  
token bucket  
priority queue



Workflows



Resources

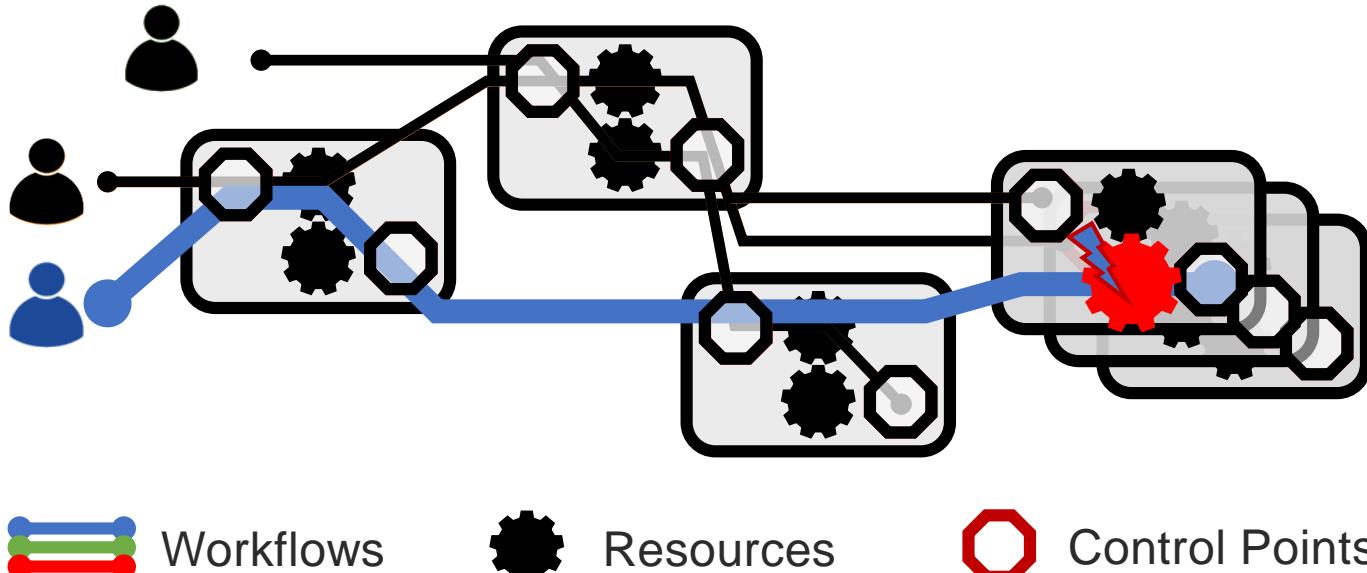


Control Points

Goal: enforce resource management decisions

Decoupled from resources

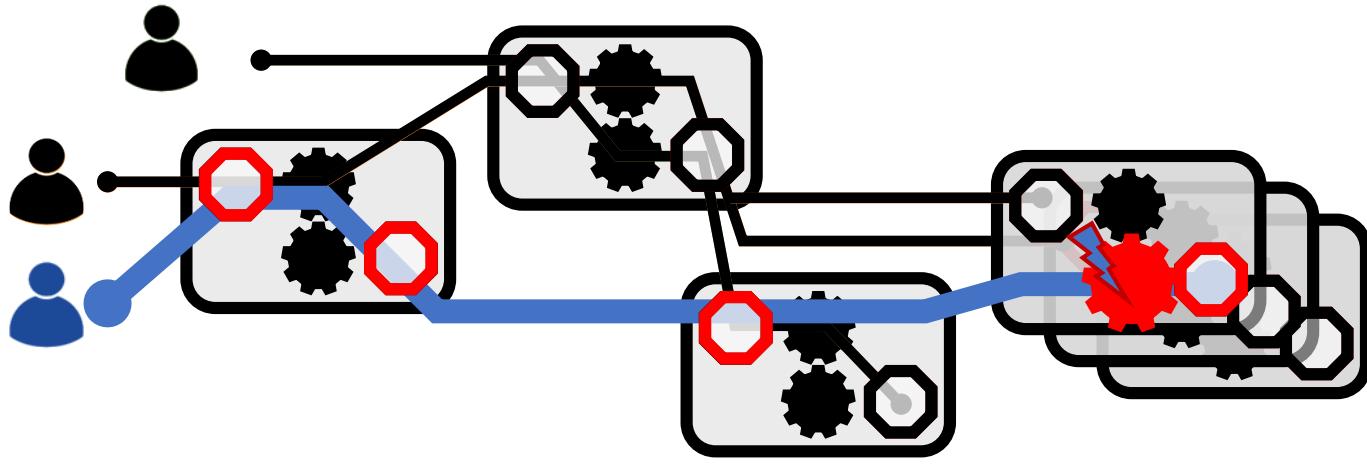
Rate-limits workflows, agnostic to underlying implementation e.g.,  
token bucket  
priority queue



Goal: enforce resource management decisions

Decoupled from resources

Rate-limits workflows, agnostic to underlying implementation e.g.,  
token bucket  
priority queue

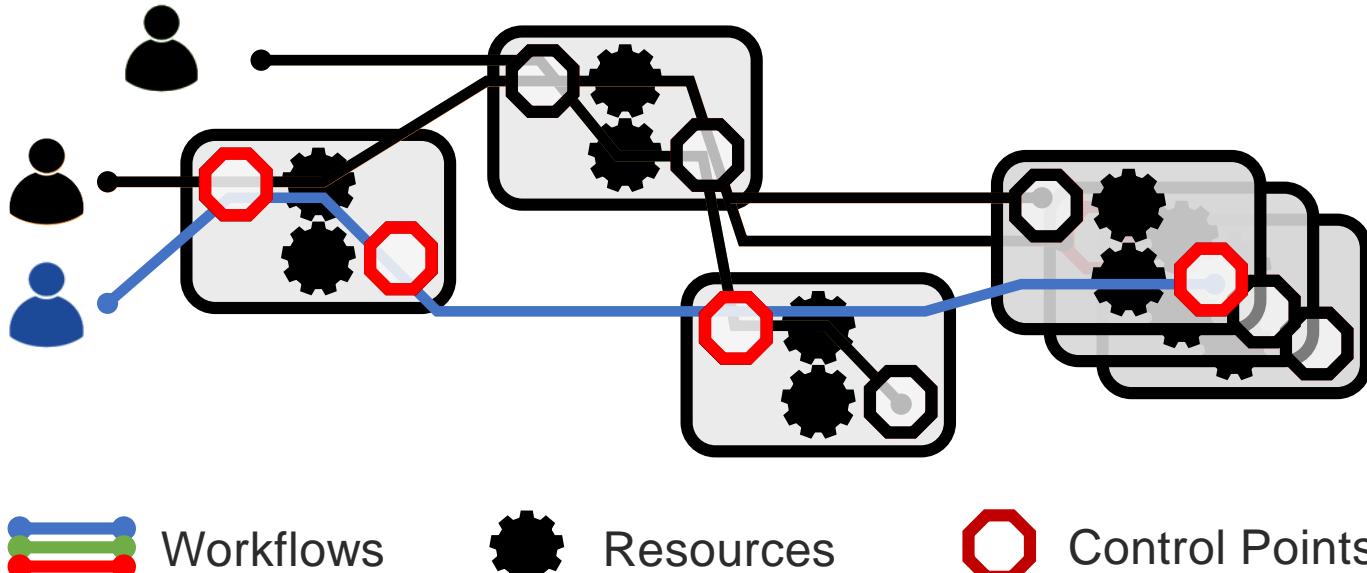


 Workflows     Resources     Control Points

Goal: enforce resource management decisions

Decoupled from resources

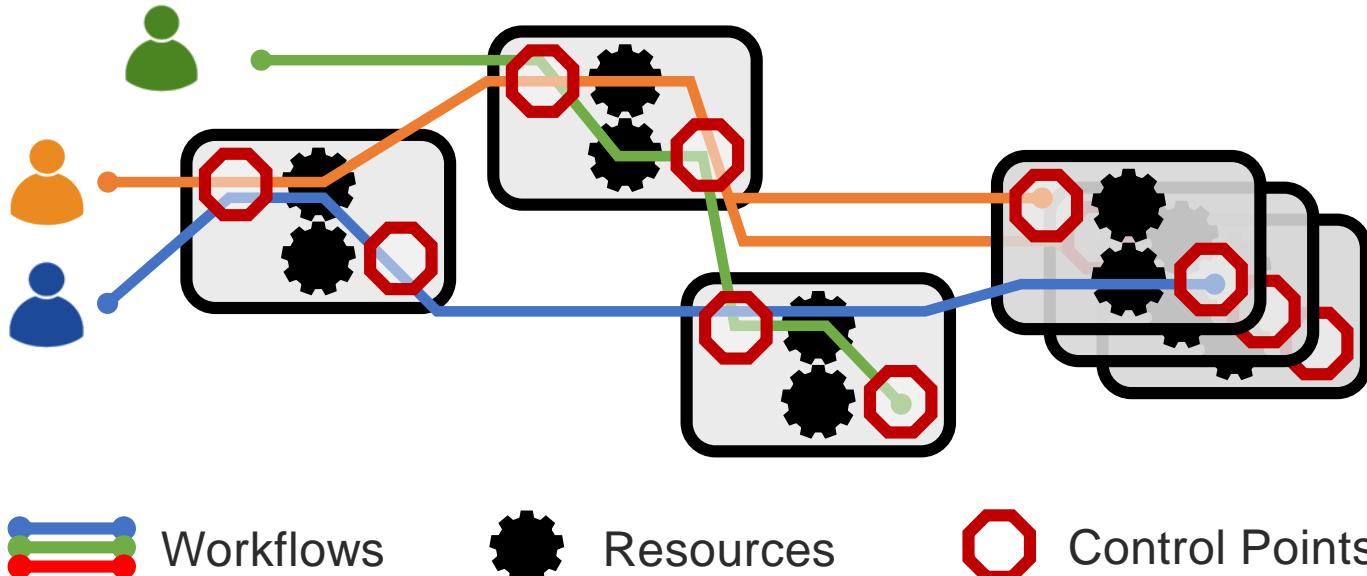
Rate-limits workflows, agnostic to underlying implementation e.g.,  
token bucket  
priority queue



Goal: enforce resource management decisions

Decoupled from resources

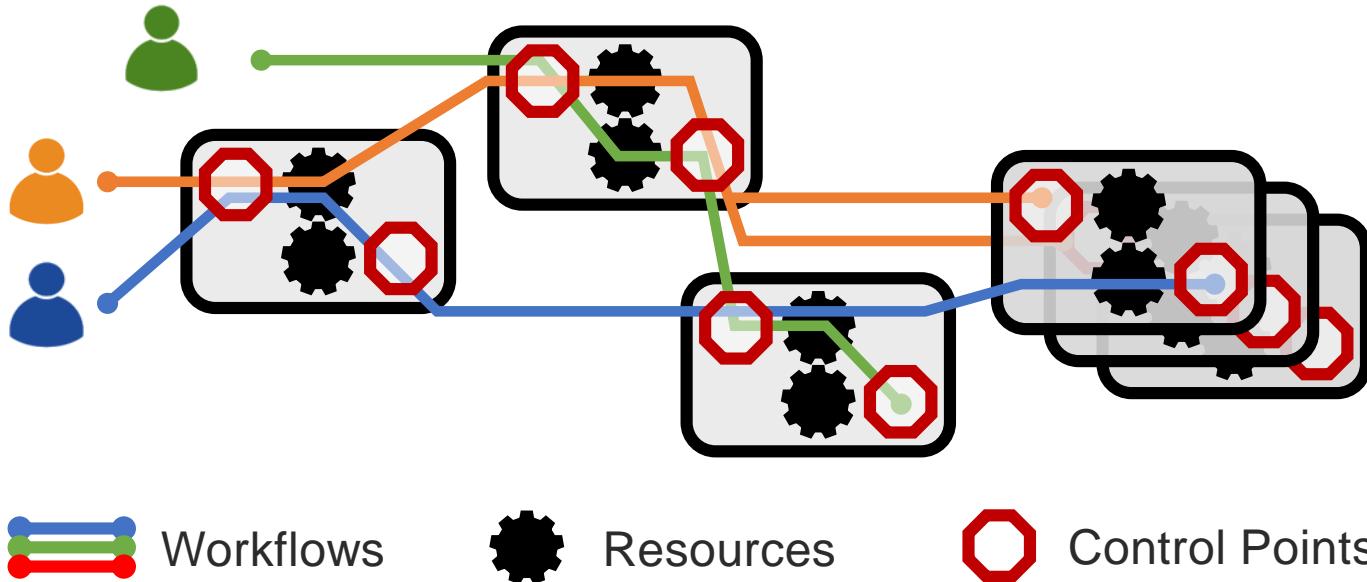
Rate-limits workflows, agnostic to underlying implementation e.g.,  
token bucket  
priority queue

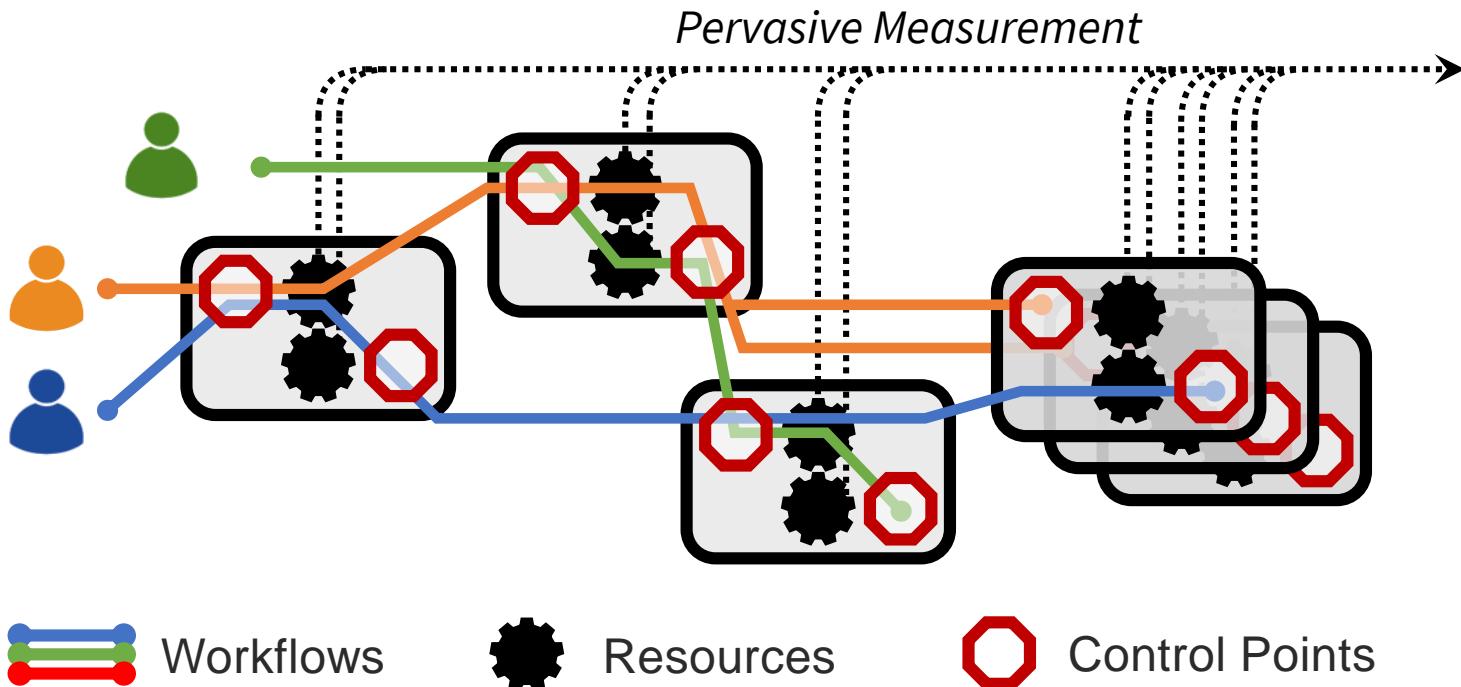


Goal: enforce resource management decisions

Decoupled from resources

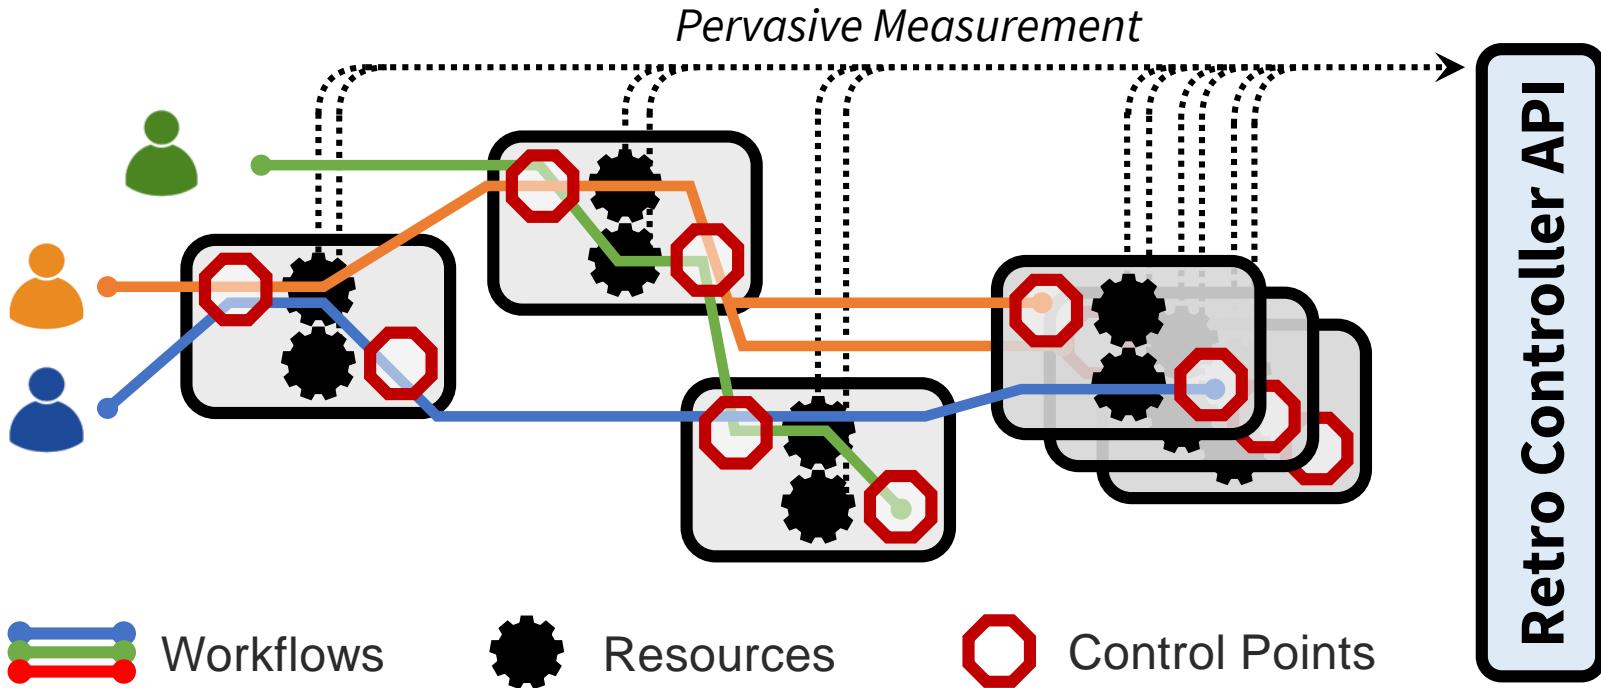
Rate-limits workflows, agnostic to underlying implementation e.g.,  
token bucket  
priority queue





## 1. Pervasive Measurement

Aggregated locally then reported centrally once per second

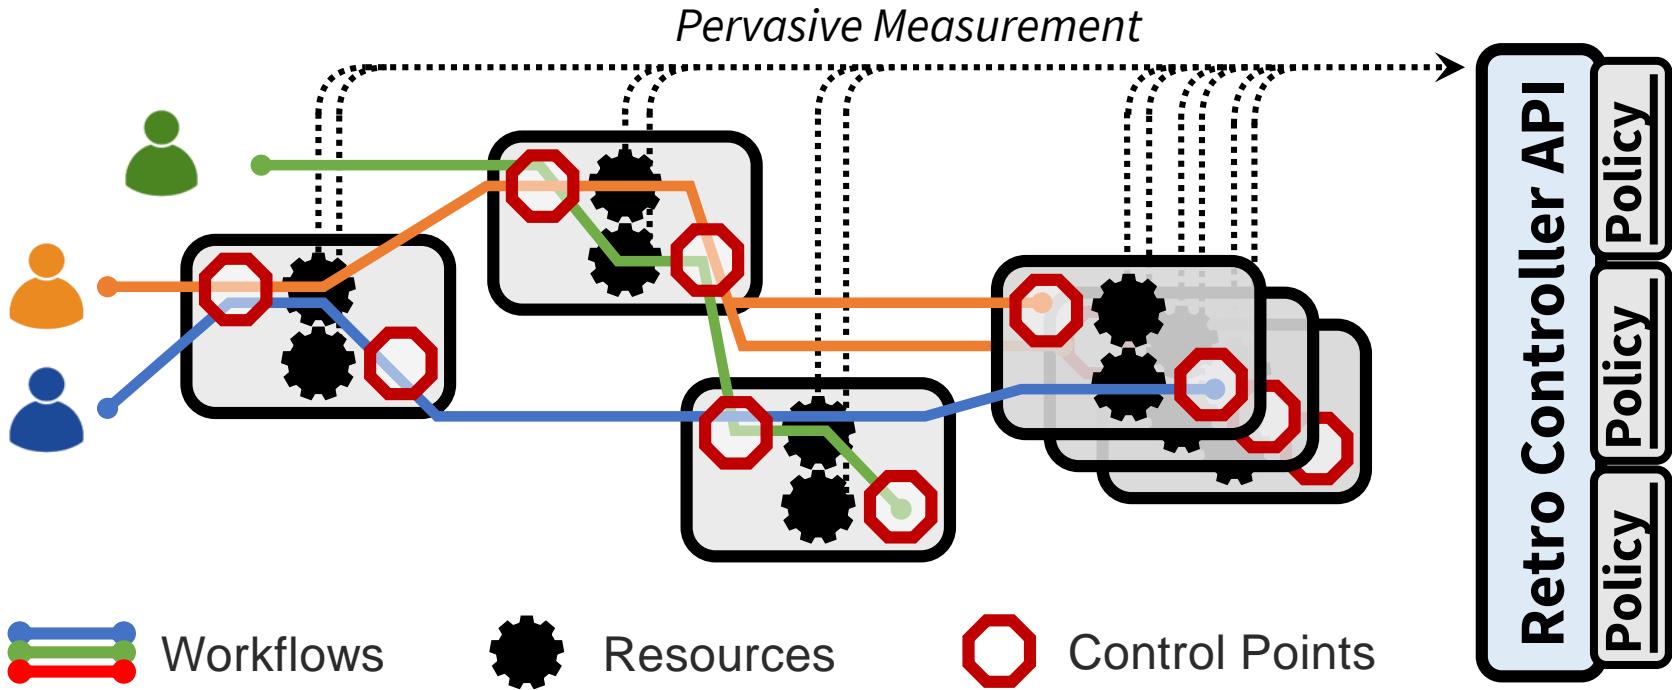


## 1. Pervasive Measurement

Aggregated locally then reported centrally once per second

## 2. Centralized Controller

Global, abstracted view of the system



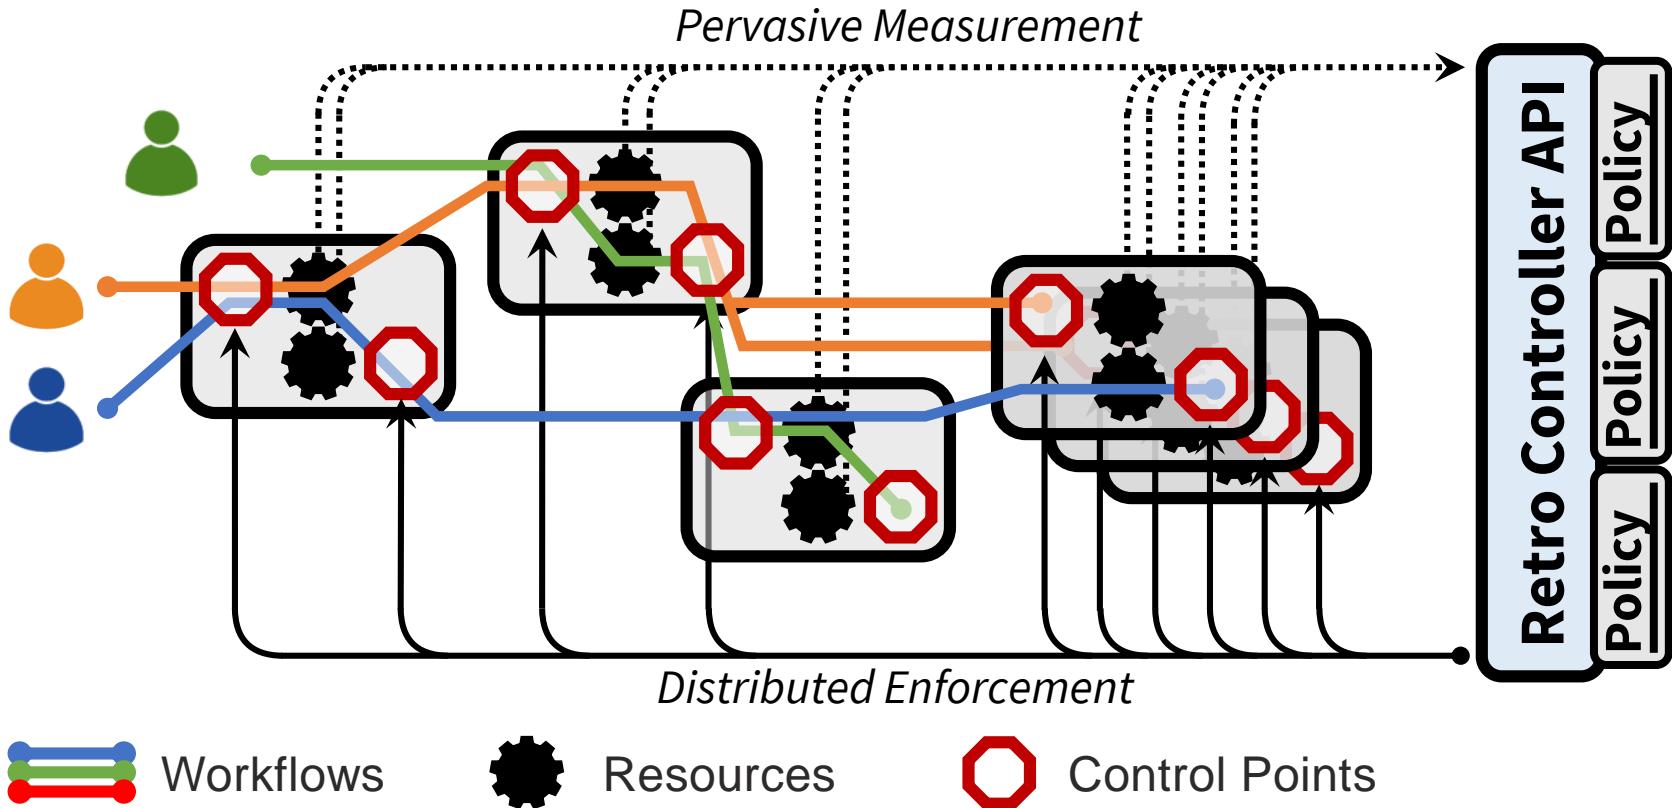
## 1. Pervasive Measurement

Aggregated locally then reported centrally once per second

## 2. Centralized Controller

Global, abstracted view of the system

Policies run in continuous control loop



### 1. Pervasive Measurement

Aggregated locally then reported centrally once per second

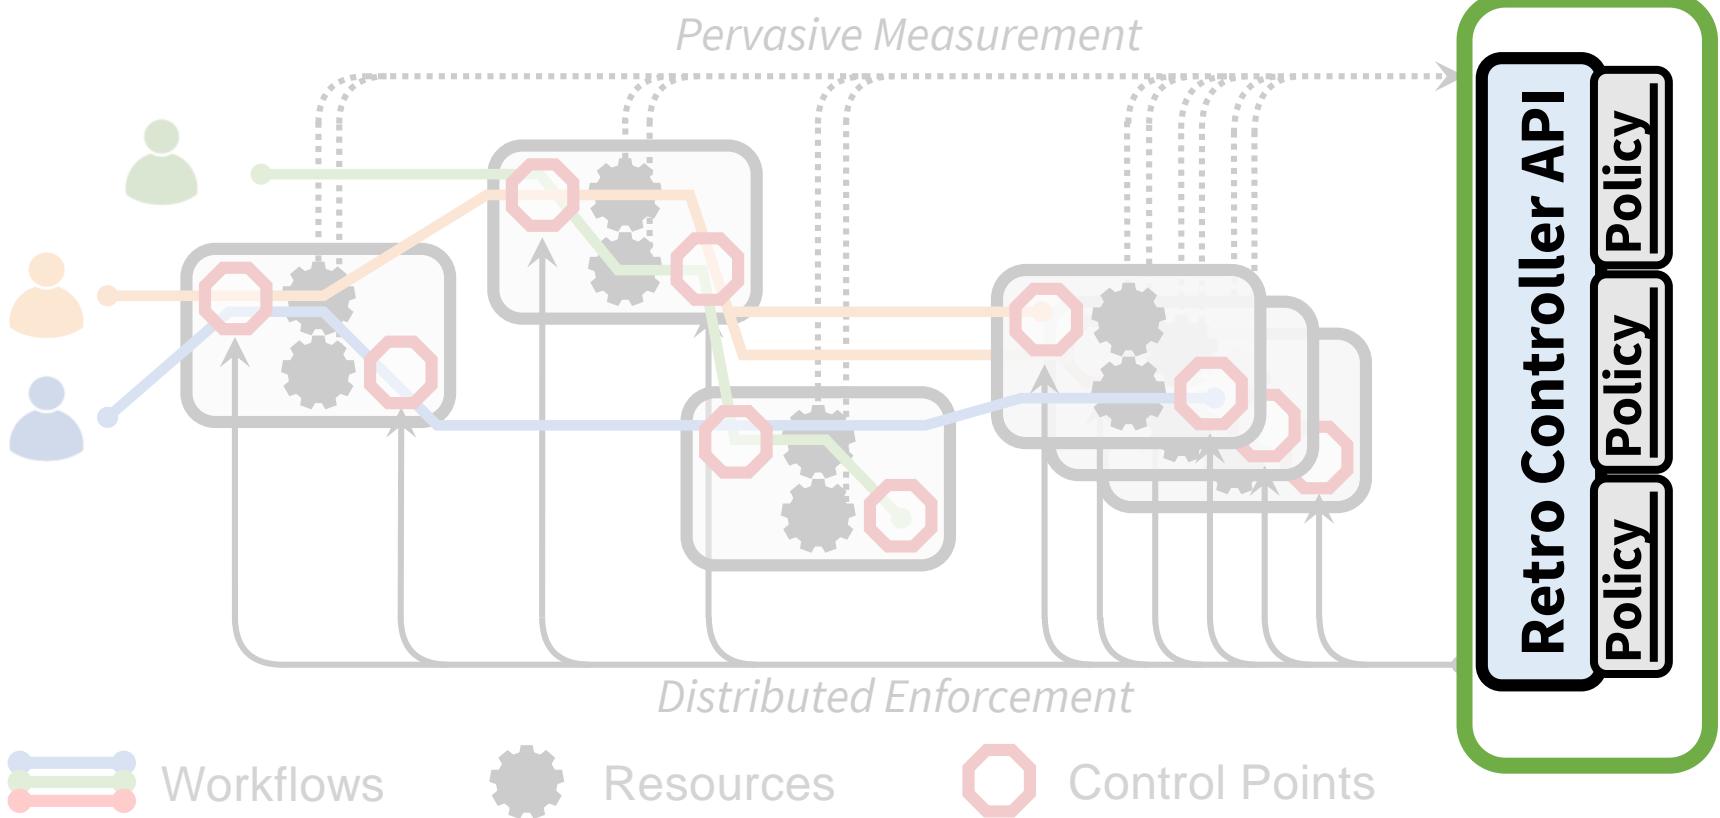
### 2. Centralized Controller

Global, abstracted view of the system

Policies run in continuous control loop

### 3. Distributed Enforcement

Co-ordinates enforcement using distributed token bucket



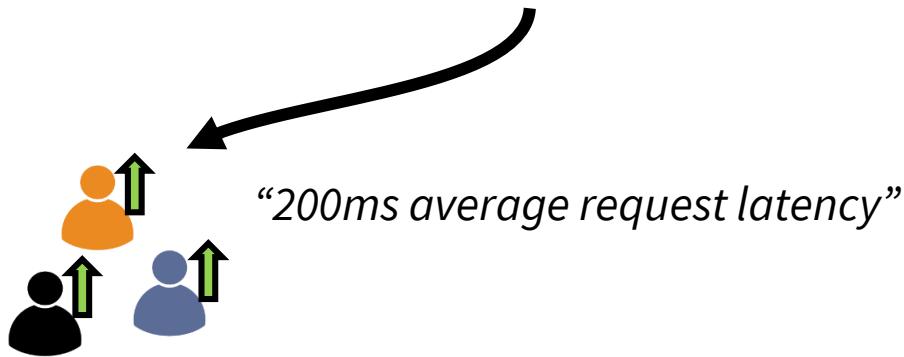
“Control Plane” for resource management  
 Global, abstracted view of the system  
 Easier to write  
 Reusable

**Policy**

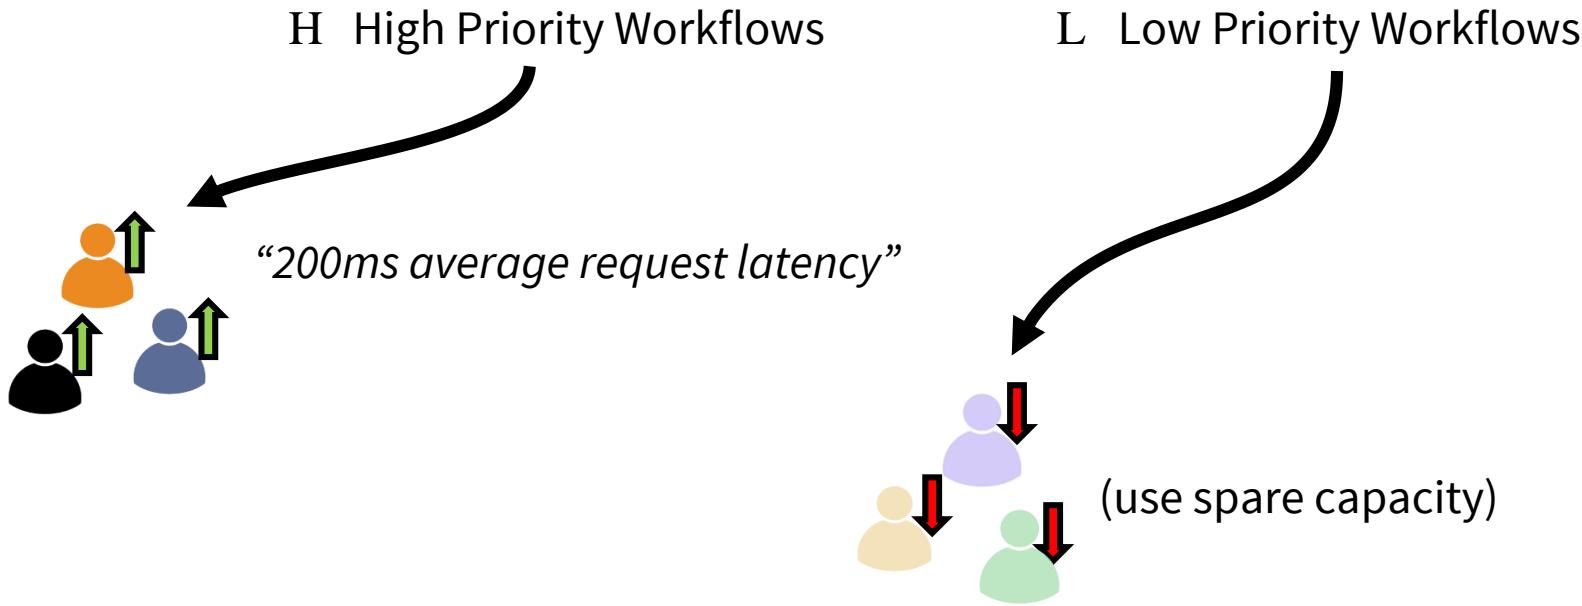
# Example: LatencySLO

# Example: LatencySLO

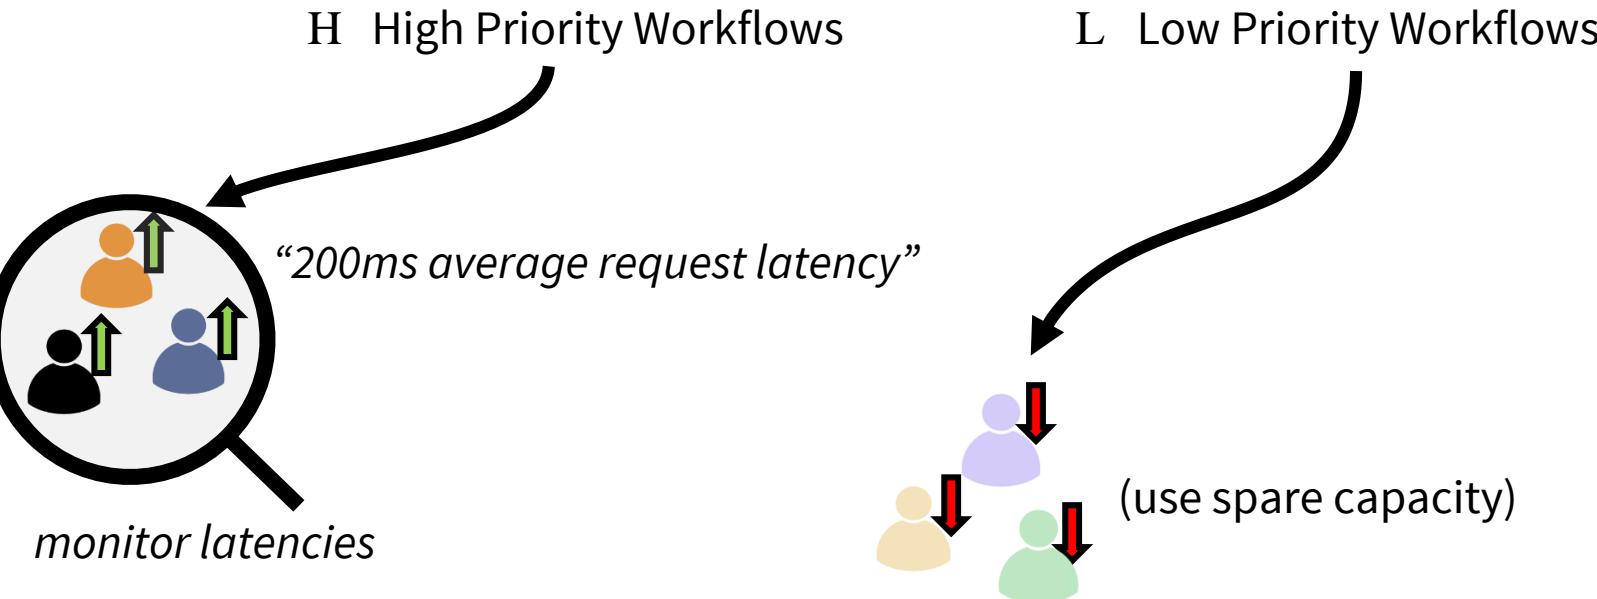
H High Priority Workflows



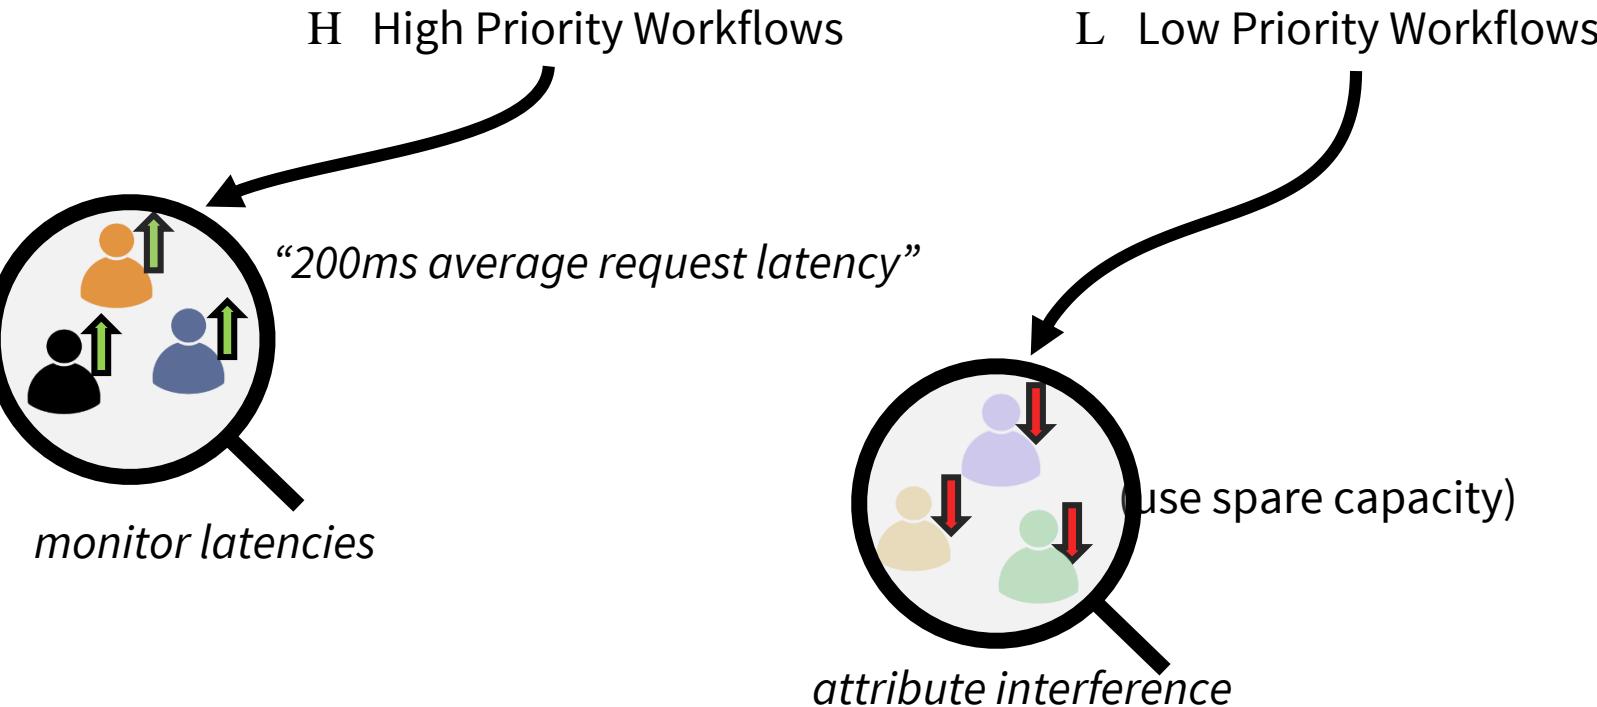
# Example: LatencySLO



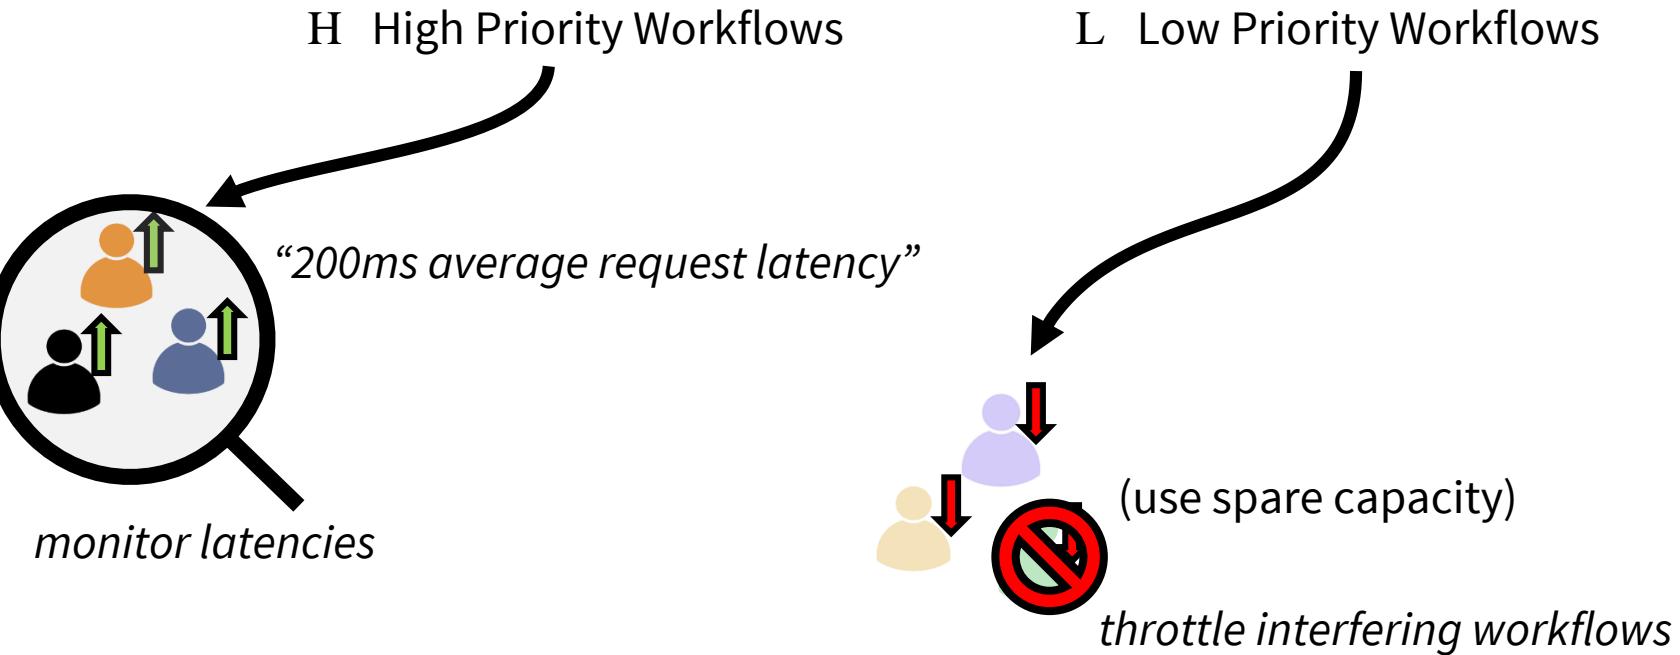
# Example: LatencySLO



# Example: LatencySLO



# Example: LatencySLO



# Example: LatencySLO

H High Priority Workflows

L Low Priority Workflows

```

1  foreach candidate in H
2      miss[candidate] = latency(candidate) / guarantee[candidate]
3  W = candidate in H with max miss[candidate]

4  foreach rsrc in resources()    // calculate importance of each resource for hipri
5      importance[rsrc] = latency(W, rsrc) * log(slowdown(rsrc))

6  foreach lopri in L          // calculate low priority workflow interference
7      interference[lopri] =  $\sum_{rsrc}$  importance[rsrc] * load(lopri, rsrc) / load(rsrc)

8  foreach lopri in L          // normalize interference
9      interference[lopri] /=  $\sum_k$  interference[k]

10 foreach lopri in L
11     if miss[W] > 1 // throttle
12         scalefactor = 1 -  $\alpha$  * (miss[W] - 1) * interference[lopri]
13     else                // release
14         scalefactor = 1 +  $\beta$ 

15 foreach cpoint in controlpoints() // apply new rates
16     set_rate(cpoint, lopri, scalefactor * get_rate(cpoint, lopri)

```

# Example: LatencySLO

H High Priority Workflows

L Low Priority Workflows

```
1  foreach candidate in H
2      miss[candidate] = latency(candidate) / guarantee[candidate]
3  W = candidate in H with max miss[candidate]
```

```
4  foreach rsrc in resources()    // calculate importance of each resource for hipri
5      importance[rsrc] = latency(W, rsrc) * log(slowdown(rsrc))

6  foreach lopri in L            // calculate low priority workflow interference
7      interference[lopri] = Σrsrc importance[rsrc] * load(lopri, rsrc) / load(rsrc)
```

```
8  foreach lopri in L          // normalize interference
9      interference[lopri] /= Σk interference[k]

10 foreach lopri in L
11     if miss[W] > 1 // throttle
12         scalefactor = 1 - α * (miss[W] - 1) * interference[lopri]
13     else             // release
14         scalefactor = 1 + β

15 foreach cpoint in controlpoints() // apply new rates
16     set_rate(cpoint, lopri, scalefactor * get_rate(cpoint, lopri))
```

# Example: LatencySLO

H High Priority Workflows

L Low Priority Workflows

```

1  foreach candidate in H
2      miss[candidate] = get_rate(cpoint, lopri) - get_rate(cpoint, candidate)
3  W = candidate in H with max miss[candidate]

```

```

4  foreach rsrc in resources()    // calculate importance of each resource for hipri
5      importance[rsrc] = latency(W, rsrc) * log(slowdown(rsrc))

6  foreach lopri in L            // calculate low priority workflow interference
7      interference[lopri] = Σrsrc importance[rsrc] * load(lopri, rsrc) / load(rsrc)

```

```

8  foreach lopri in L          // normalize interference
9      interference[lopri] /= Σk interference[k]

10 foreach lopri in L
11     if miss[W] > 1 // throttle
12         scalefactor = 1 - α * (miss[W] - 1) * interference[lopri]
13     else             // release
14         scalefactor = 1 + β

15 foreach cpoint in controlpoints() // apply new rates
16     set_rate(cpoint, lopri, scalefactor * get_rate(cpoint, lopri))

```

# Example: LatencySLO

H High Priority Workflows

L Low Priority Workflows

```

1  foreach candidate in H
2      miss[candidate] = get_miss_rate(candidate)
3  W = candidate in H with max miss[candidate]

```

**Select the high priority workflow  $W$  with worst performance**

```

4  foreach rsrc in resources() // calculate importance of each resource for hipri
5      importance[rsrc] = latency(rsrc) * log(1 + load(rsrc))
6  foreach lopri in L          // calculate low priority workflow interference
7      interference[lopri] = Σrsrc importance[rsrc] * load(lopri, rsrc) / load(rsrc)

```

**Weight low priority workflows by their interference with  $W$**

```

8  foreach lopri in L          // normalize interference
9      interference[lopri] /= Σk interference[k]

10 foreach lopri in L
11     if miss[W] > 1 // throttle
12         scalefactor = 1 - α * (miss[W] - 1) * interference[lopri]
13     else             // release
14         scalefactor = 1 + β

15 foreach cpoint in controlpoints() // apply new rates
16     set_rate(cpoint, lopri, scalefactor * get_rate(cpoint, lopri))

```

# Example: LatencySLO

H High Priority Workflows

L Low Priority Workflows

```

1  foreach candidate in H
2      miss[candidate] = get_miss_rate(candidate)
3  W = candidate in H with max miss[candidate]

```

**Select the high priority workflow  $W$  with worst performance**

```

4  foreach rsrc in resources()    // calculate importance of each resource for hipri
5      importance[rsrc] = latency(rsrc) * log(1 + load(rsrc))
6  foreach lopri in L            // calculate low priority workflow interference
7      interference[lopri] = Σrsrc importance[rsrc] * load(lopri, rsrc) / load(rsrc)

```

```

8  foreach lopri in L          // normalize interference
9      interference[lopri] /= Σk interference[k]

```

**Weight low priority workflows by their interference with  $W$**

```

10 foreach lopri in L
11     if miss[W] > threshold
12         scalefactor = 1 - α * (miss[W] - 1) * interference[lopri]
13     else
14         scalefactor = 1 + β

```

```

15  foreach cpoint in controlpoints() // apply new rates
16      set_rate(cpoint, lopri, scalefactor * get_rate(cpoint, lopri))

```

# Example: LatencySLO

H High Priority Workflows

L Low Priority Workflows

**Select the high priority workflow  $W$  with worst performance**

```

1  foreach candidate in H
2      miss[candidate] = latency(candidate) / guarantee[candidate]
3  W = candidate in H with max miss[candidate]
```

```

4  foreach rsrc in resources()    // calculate importance of each resource for hipri
5      importance[rsrc] = latency(W, rsrc) * log(slowdown(rsrc))
6  foreach lopri in L            // calculate low priority workflow interference
7      interference[lopri] =  $\sum_{rsrc}$  importance[rsrc] * load(lopri, rsrc) / load(rsrc)
```

```

8  foreach lopri in L          // normalize interference
9      interference[lopri] /=  $\sum_k$  interference[k]
```

```

10 foreach lopri in L
11     if miss[W] > 1 // throttle
12     else             // release
13         scalefactor = 1 +  $\beta$ 
```

**Throttle low priority workflows proportionally to their weight**

```

15  foreach cpoint in controlpoints() // apply new rates
16      set_rate(cpoint, lopri, scalefactor * get_rate(cpoint, lopri))
```

# Example: LatencySLO

H High Priority Workflows

L Low Priority Workflows

**Select the high priority workflow  $W$  with worst performance**

```

1  foreach candidate in H
2      miss[candidate] = latency(candidate) / guarantee[candidate]
3   $W = \text{candidate in } H \text{ with max miss[candidate]}$ 

```

**Weight low priority workflows by their interference with  $W$**

```

4  foreach rsrc in resources() // calculate importance of each resource for hipri
5      importance[rsrc] = latency( $W$ , rsrc) * log(slowdown(rsrc))

6  foreach lopri in L           // calculate low priority workflow interference
7      interference[lopri] =  $\sum_{rsrc}$  importance[rsrc] * load(lopri, rsrc) / load(rsrc)

```

```

8  foreach lopri in L           // normalize interference
9      interference[lopri] /=  $\sum_k$  interference[k]

```

```

10 foreach lopri in L
11     if miss[ $W$ ] > 1 // throttle
12     else                // release
13         scalefactor = 1 +  $\beta$ 

```

**Throttle low priority workflows proportionally to their weight**

```

15  foreach cpoint in controlpoints() // apply new rates
16      set_rate(cpoint, lopri, scalefactor * get_rate(cpoint, lopri))

```

# Example: LatencySLO

H High Priority Workflows

L Low Priority Workflows

**Select the high priority workflow  $W$  with worst performance**

```

1  foreach candidate in H
2      miss[candidate] = latency(candidate) / guarantee[candidate]
3   $W = \text{candidate in } H \text{ with max miss[candidate]}$ 

```

**Weight low priority workflows by their interference with  $W$**

```

4  foreach rsrc in resources() // calculate importance of each resource for hipri
5      importance[rsrc] = latency( $W$ , rsrc) * log(slowdown(rsrc))

6  foreach lopri in L           // calculate low priority workflow interference
7      interference[lopri] =  $\sum_{rsrc} \text{importance}[rsrc] * \text{load}(lopri, rsrc) / \text{load}(rsrc)$ 

```

**Throttle low priority workflows proportionally to their weight**

```

8  foreach lopri in L           // normalize interference
9      interference[lopri] /=  $\sum_k \text{interference}[k]$ 

10 foreach lopri in L
11     if miss[ $W$ ] > 1 // throttle
12         scalefactor = 1 -  $\alpha * (\text{miss}[W] - 1) * \text{interference}[lopri]$ 
13     else                  // release
14         scalefactor = 1 +  $\beta$ 

15     foreach cpoint in controlpoints() // apply new rates
16         set_rate(cpoint, lopri, scalefactor * get_rate(cpoint, lopri))

```

# Example: LatencySLO

H High Priority Workflows

L Low Priority Workflows

**Select the high priority workflow  $W$  with worst performance**

```

1  foreach candidate in H
2      miss[candidate] = latency(candidate) / guarantee[candidate]
3   $W = \text{candidate in } H \text{ with max miss[candidate]}$ 

```

**Weight low priority workflows by their interference with  $W$**

```

4  foreach rsrc in resources() // calculate importance of each resource for hipri
5      importance[rsrc] = latency( $W$ , rsrc) * log(slowdown(rsrc))

6  foreach lopri in L           // calculate low priority workflow interference
7      interference[lopri] =  $\sum_{rsrcc} \text{importance}[rsrcc] * \text{load}(lopri, rsrc) / \text{load}(rsrcc)$ 

```

**Throttle low priority workflows proportionally to their weight**

```

8  foreach lopri in L           // normalize interference
9      interference[lopri] /=  $\sum_k \text{interference}[k]$ 

10 foreach lopri in L
11     if miss[ $W$ ] > 1 // throttle
12         scalefactor =  $1 - \alpha^{(miss[W] - 1)} * \text{interference}[lopri]$ 
13     else                   // release
14         scalefactor =  $1 + \beta$ 

15 foreach cpoint in controlpoints() // apply new rates
16     set_rate(cpoint, lopri, scalefactor * get_rate(cpoint, lopri))

```

# Other types of policy...

# Other types of policy...

**Policy**

Bottleneck Fairness

Detect most overloaded resource

Fair-share resource between tenants using it

# Other types of policy...

## **Policy**

### Bottleneck Fairness

Detect most overloaded resource

Fair-share resource between tenants using it

## **Policy**

### Dominant Resource Fairness

Estimate demands and capacities from measurements

# Other types of policy...

## **Policy**

### Bottleneck Fairness

Detect most overloaded resource

Fair-share resource between tenants using it

## **Policy**

### Dominant Resource Fairness

Estimate demands and capacities from measurements

Concise

Any resources can be bottleneck (policy doesn't care)

Not system specific

# *Retro Evaluation*

# Instrumentation

# Instrumentation

Retro implementation in Java

Instrumentation Library

Central controller implementation

# Instrumentation

## Retro implementation in Java

Instrumentation Library

Central controller implementation

## To enable Retro



Propagate Workflow ID within application (like X-Trace, Dapper)



Instrument resources with wrapper classes

# Instrumentation

## Retro implementation in Java

Instrumentation Library

Central controller implementation

## To enable Retro



Propagate Workflow ID within application (like X-Trace, Dapper)



Instrument resources with wrapper classes

## Overheads

# Instrumentation

## Retro implementation in Java

Instrumentation Library

Central controller implementation

## To enable Retro

-  Propagate Workflow ID within application (like X-Trace, Dapper)
-  Instrument resources with wrapper classes

## Overheads

Resource instrumentation automatic using AspectJ

# Instrumentation

## Retro implementation in Java

Instrumentation Library

Central controller implementation

## To enable Retro

-  Propagate Workflow ID within application (like X-Trace, Dapper)
-  Instrument resources with wrapper classes

## Overheads

Resource instrumentation automatic using AspectJ

Overall 50-200 lines per system to modify RPCs

# Instrumentation

## Retro implementation in Java

Instrumentation Library

Central controller implementation

### To enable Retro

-  Propagate Workflow ID within application (like X-Trace, Dapper)
-  Instrument resources with wrapper classes

### Overheads

Resource instrumentation automatic using AspectJ

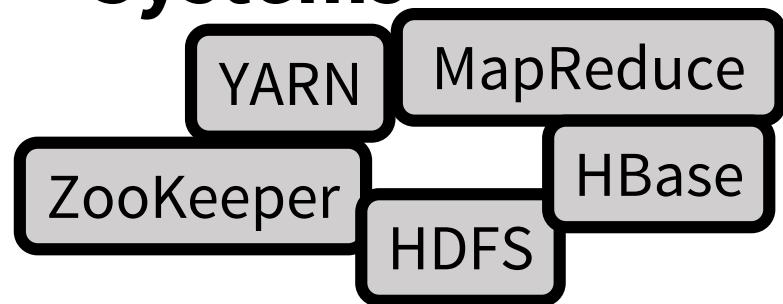
Overall 50-200 lines per system to modify RPCs

Retro overhead: max 1-2% on throughput, latency

# Experiments

# Experiments

## Systems



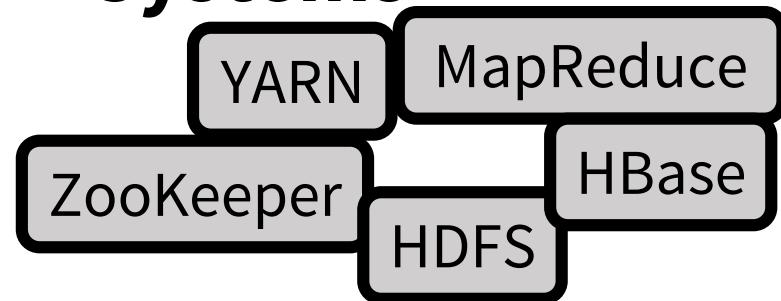
# Experiments

## Workflows



MapReduce Jobs (HiBench)  
HBase (YCSB)  
HDFS clients  
Background Data Replication  
Background Heartbeats

## Systems



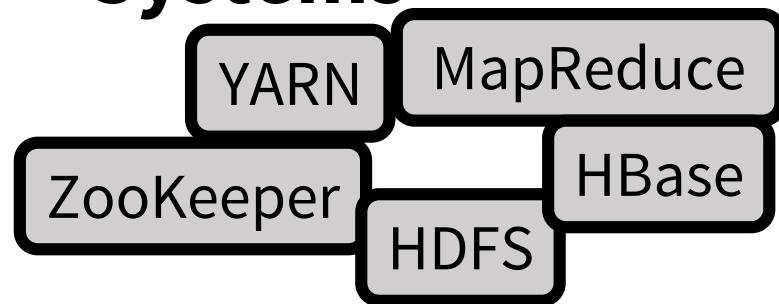
# Experiments

## Workflows



MapReduce Jobs (HiBench)  
HBase (YCSB)  
HDFS clients  
Background Data Replication  
Background Heartbeats

## Systems



## Resources



CPU, Disk, Network (All systems)  
Locks, Queues (HDFS, HBase)

# Experiments

## Workflows



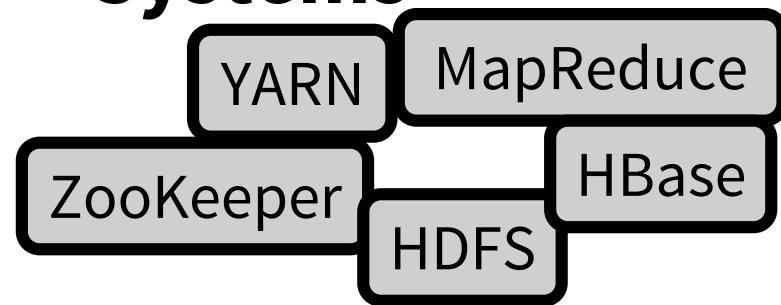
MapReduce Jobs (HiBench)  
HBase (YCSB)  
HDFS clients  
Background Data Replication  
Background Heartbeats

## Resources



CPU, Disk, Network (All systems)  
Locks, Queues (HDFS, HBase)

## Systems



## Policies

LatencySLO

Bottleneck Fairness

Dominant Resource Fairness

# Experiments

## Workflows



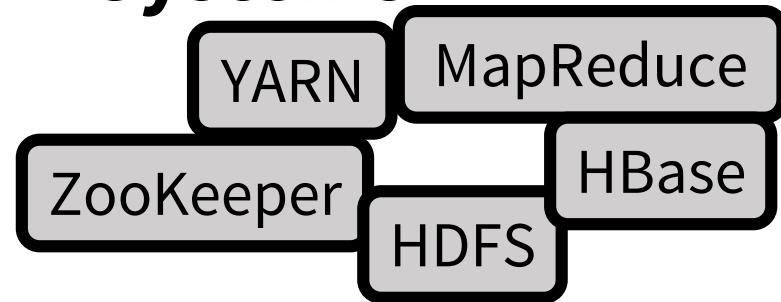
MapReduce Jobs (HiBench)  
HBase (YCSB)  
HDFS clients  
Background Data Replication  
Background Heartbeats

## Resources



CPU, Disk, Network (All systems)  
Locks, Queues (HDFS, HBase)

## Systems



## Policies

LatencySLO

Bottleneck Fairness

Dominant Resource Fairness

Policies for a mixture of systems, workflows, and resources

Results on clusters up to 200 nodes

See paper for full experiment results

# Experiments

## Workflows



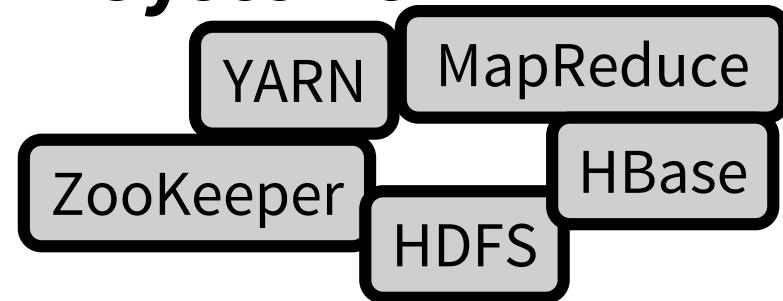
MapReduce Jobs (HiBench)  
HBase (YCSB)  
HDFS clients  
Background Data Replication  
Background Heartbeats

## Resources



CPU, Disk, Network (All systems)  
Locks, Queues (HDFS, HBase)

## Systems



## Policies

LatencySLO

Bottleneck Fairness

Dominant Resource Fairness

Policies for a mixture of systems, workflows, and resources

Results on clusters up to 200 nodes

See paper for full experiment results

**This talk** LatencySLO policy results

# Experiments

## Workflows



MapReduce Jobs (HiBench)

HBase (YCSB)

HDFS clients

Background Data Replication

Background Heartbeats

## Resources



CPU, Disk, Network (All systems)

Locks, Queues (HDFS, HBase)

## Systems



## Policies

**LatencySLO**

Bottleneck Fairness

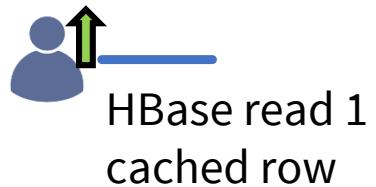
Dominant Resource Fairness

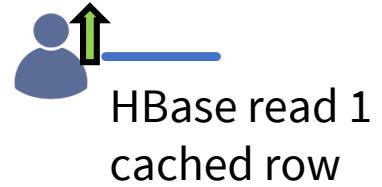
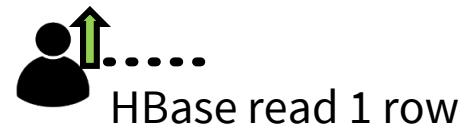
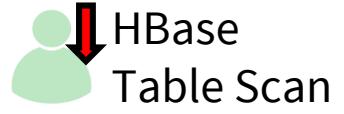
Policies for a mixture of systems, workflows, and resources

Results on clusters up to 200 nodes

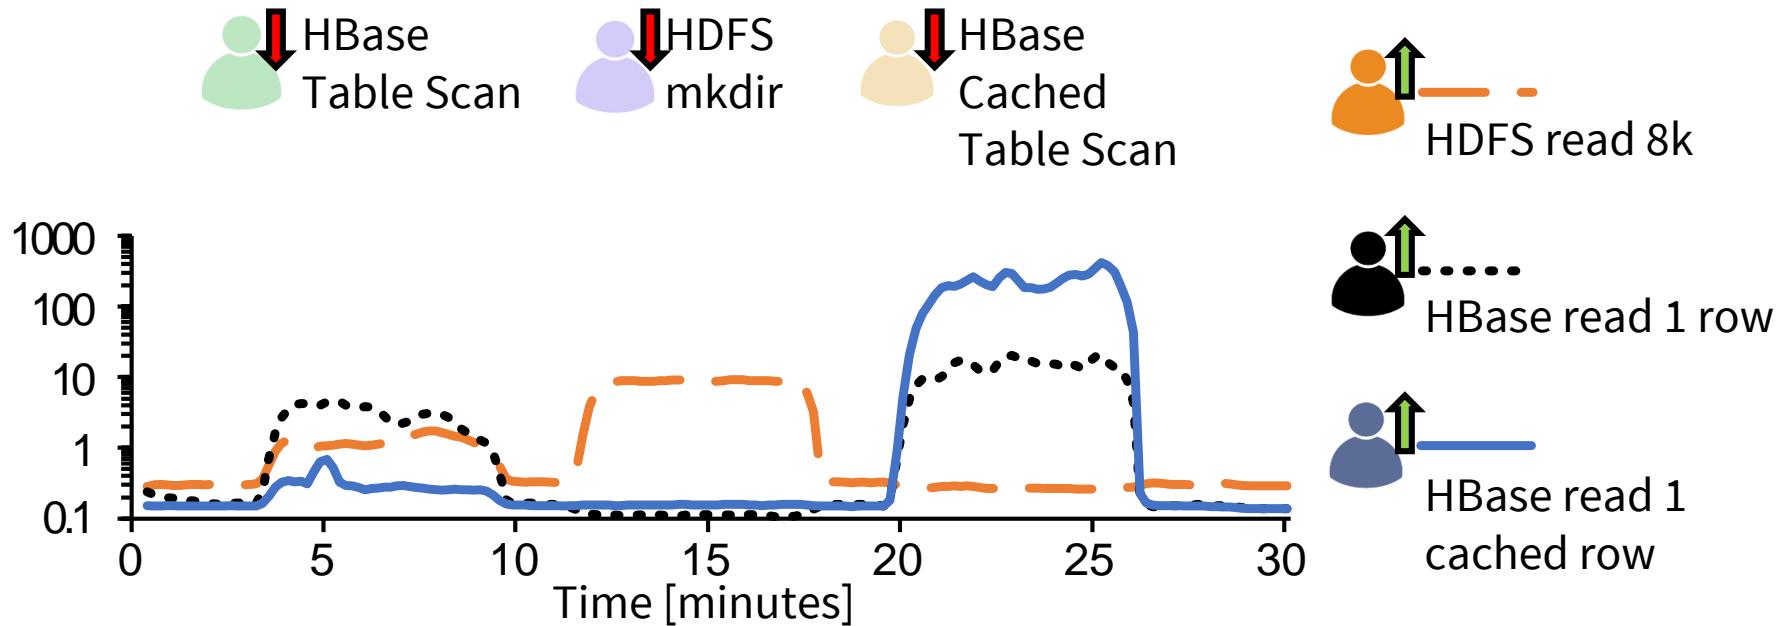
See paper for full experiment results

**This talk** LatencySLO policy results

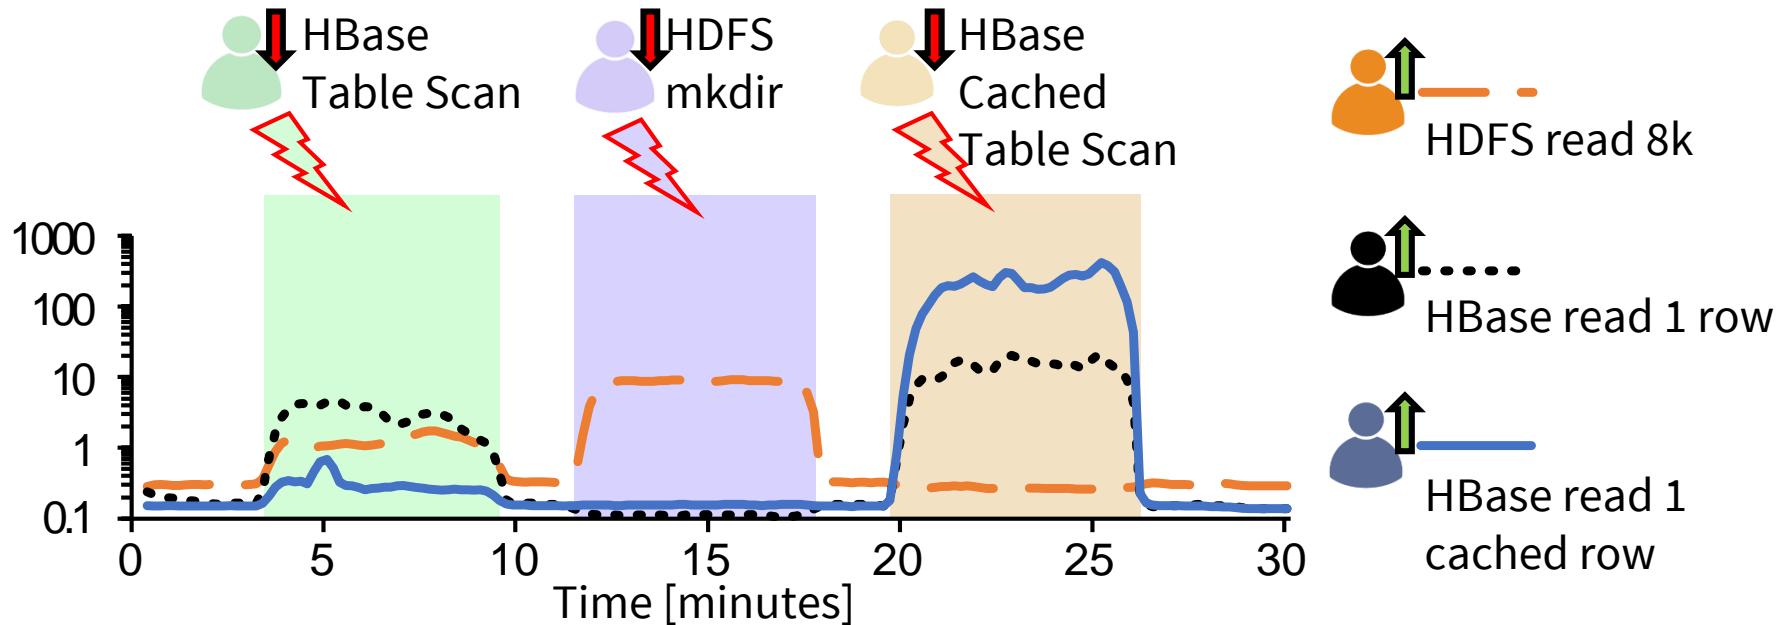




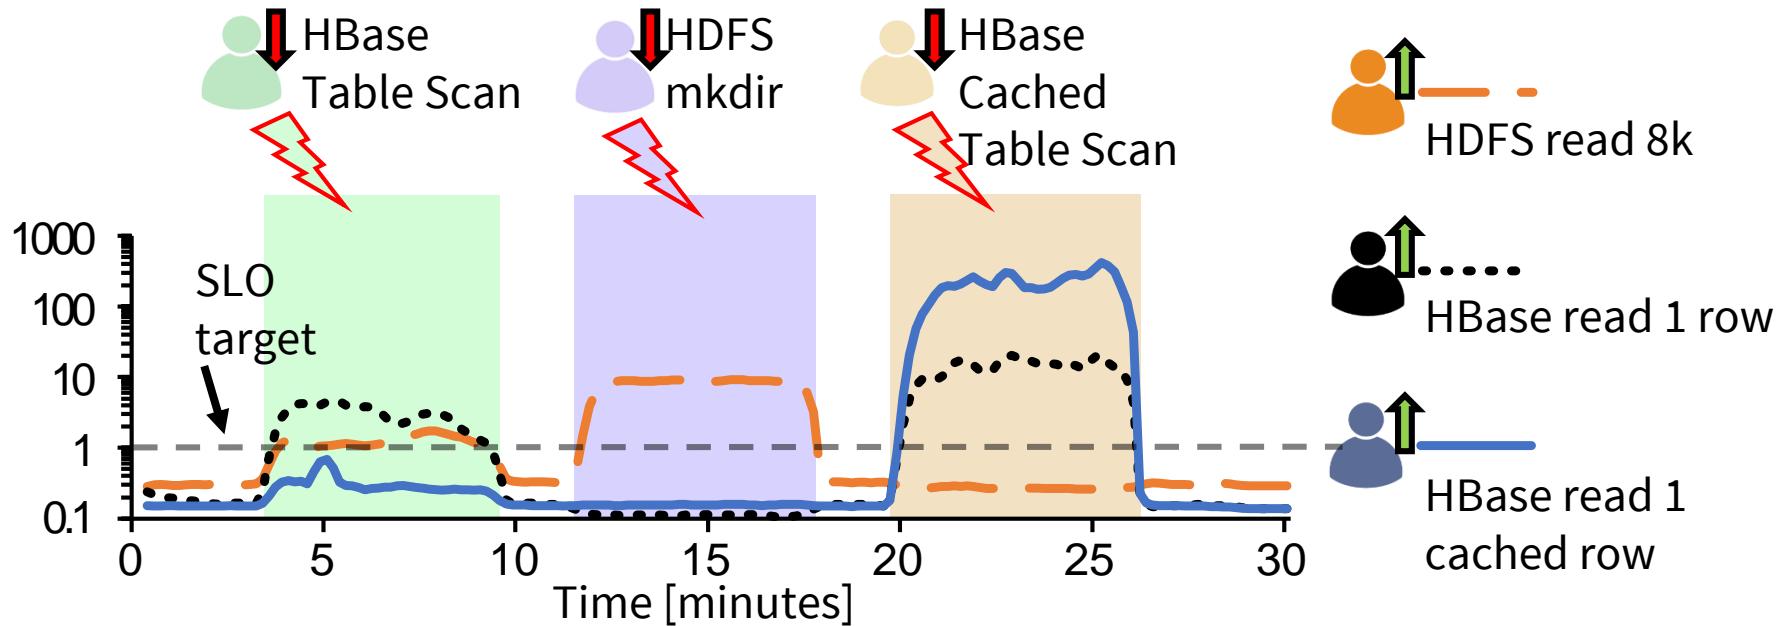
SLO-Normalized Latency



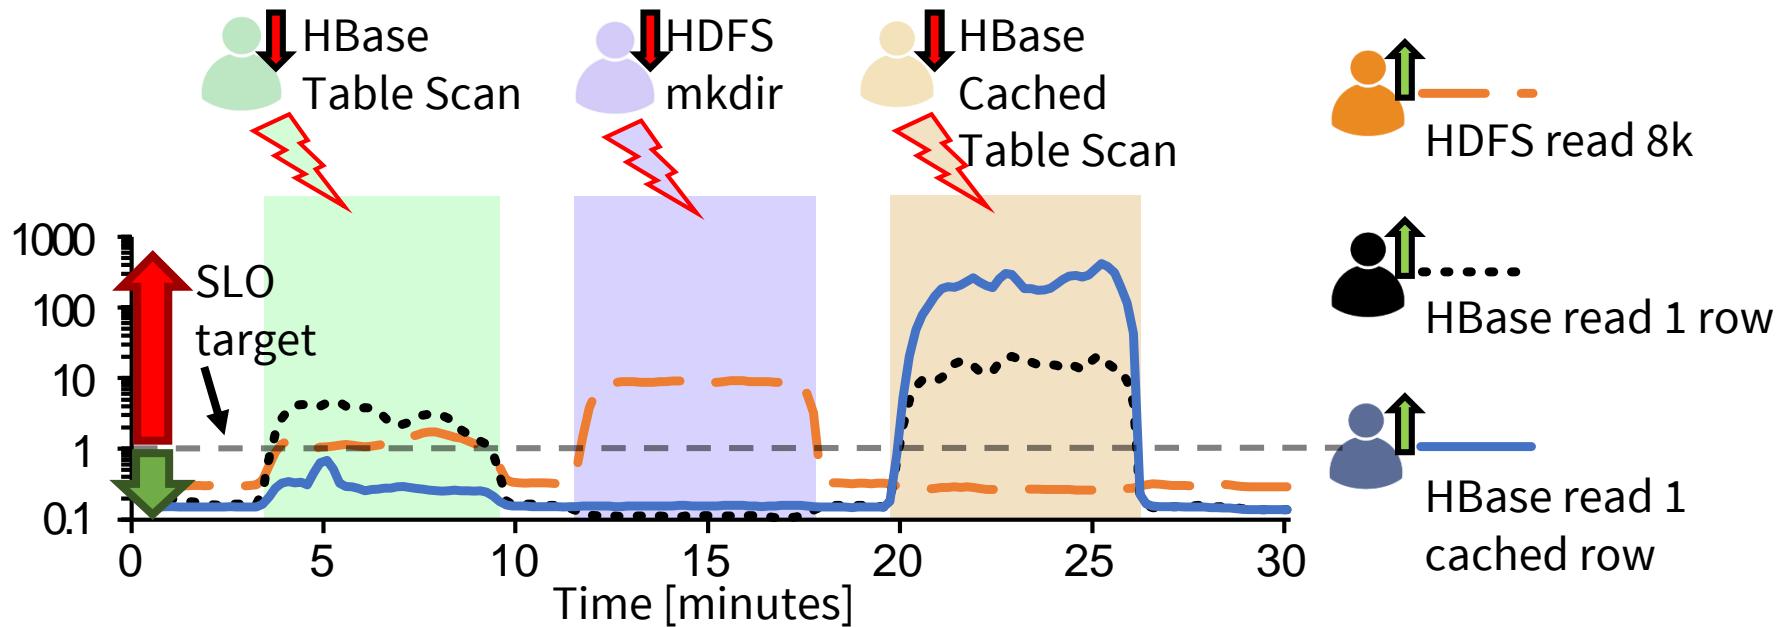
SLO-Normalized Latency



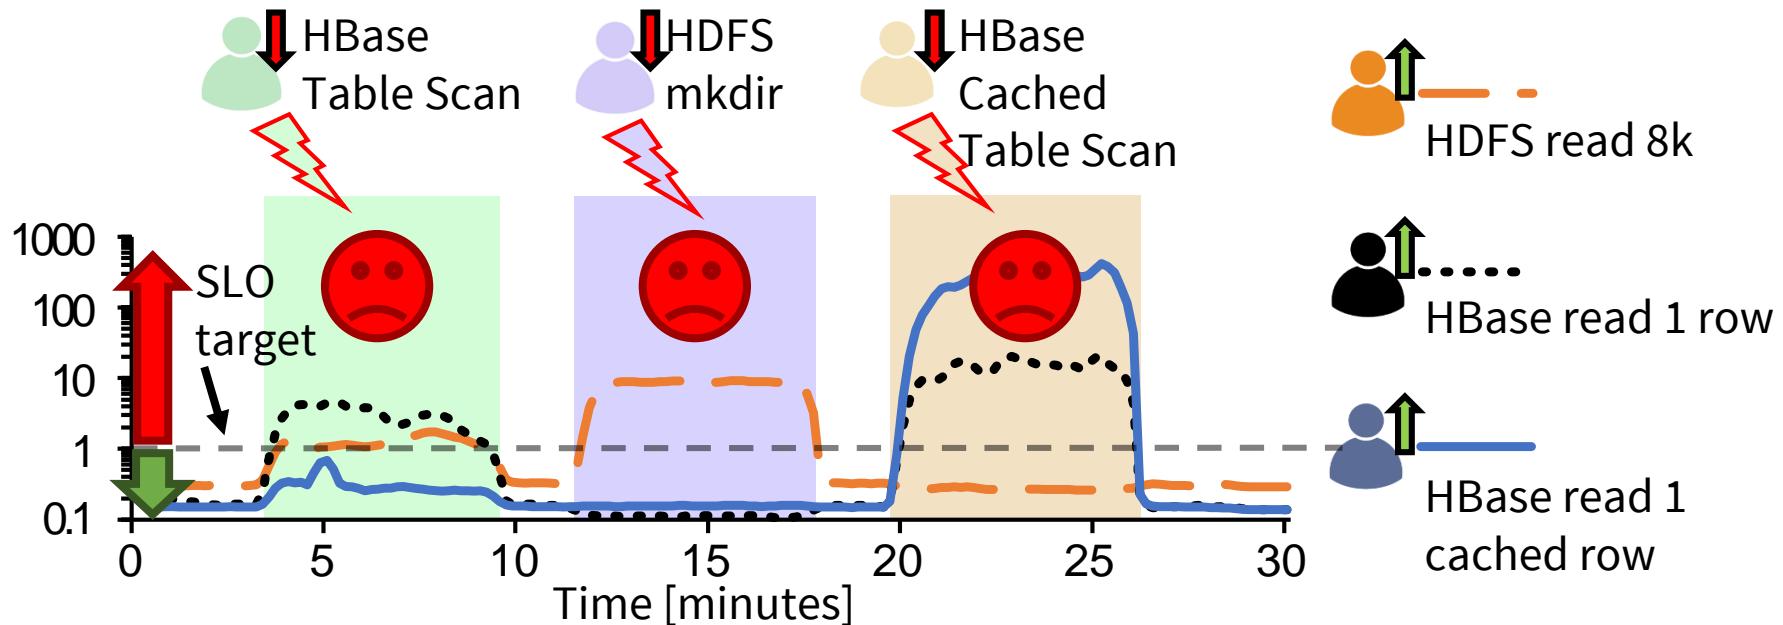
SLO-Normalized Latency

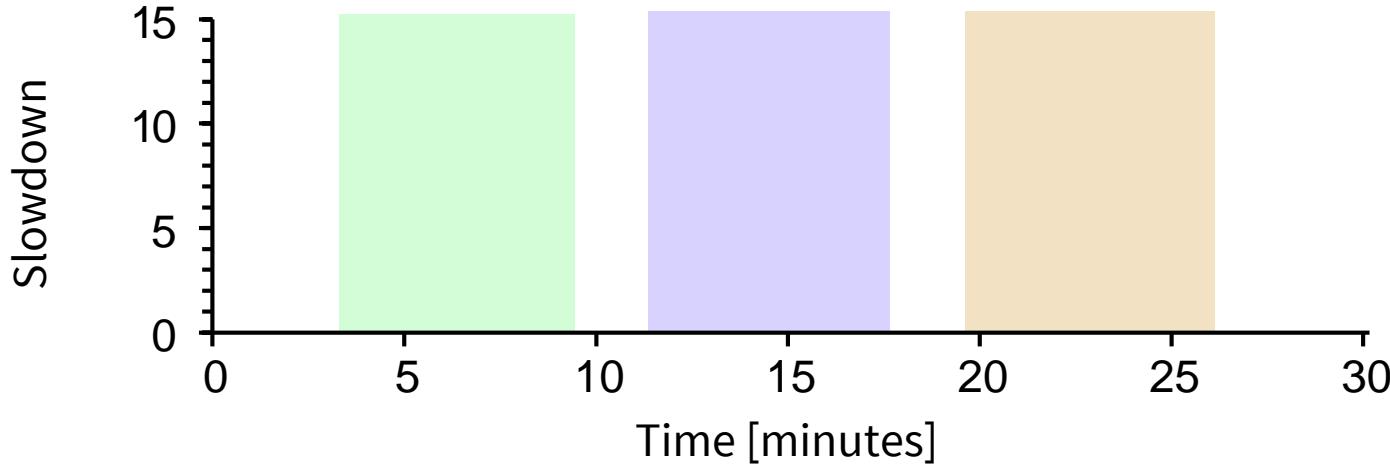
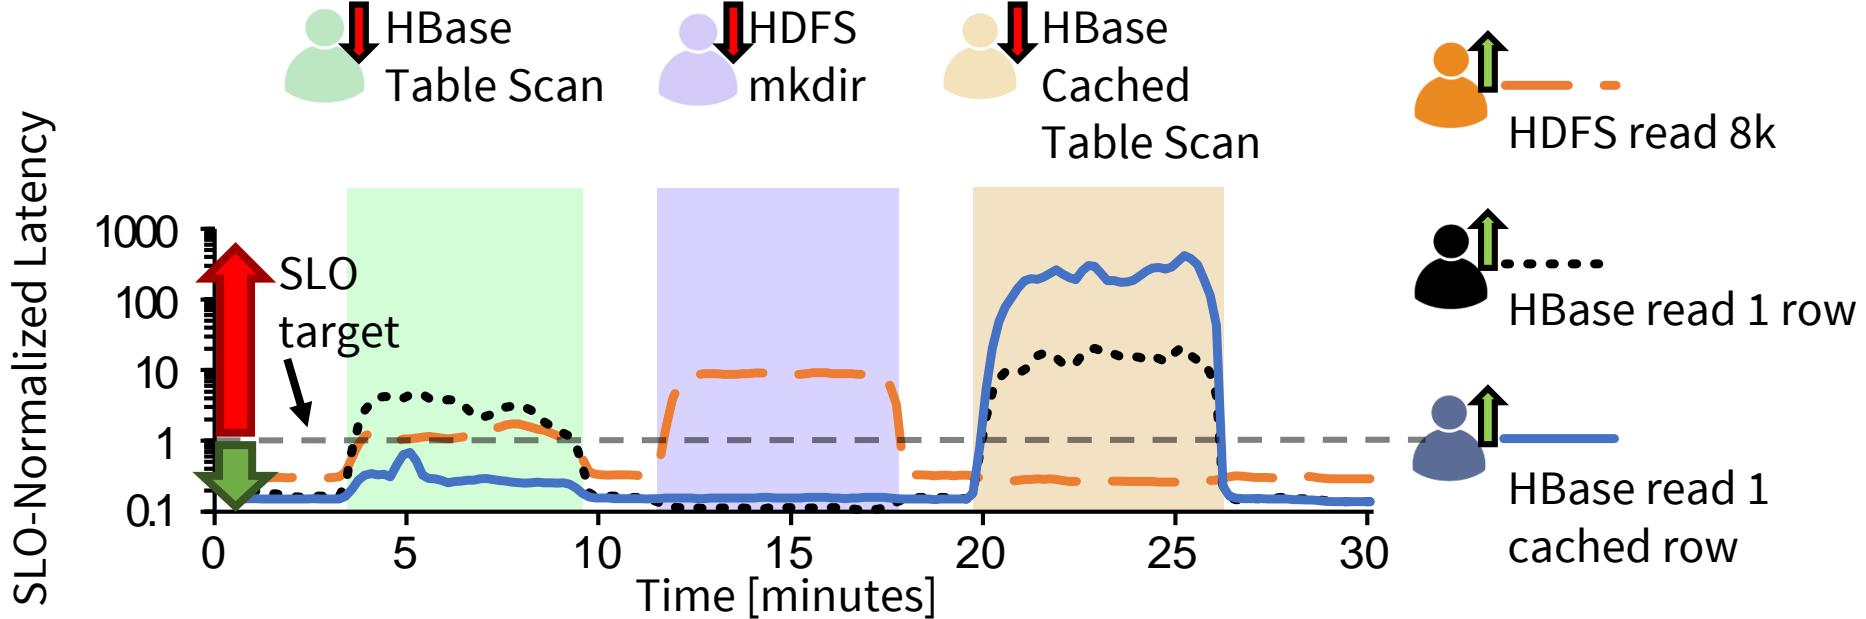


SLO-Normalized Latency

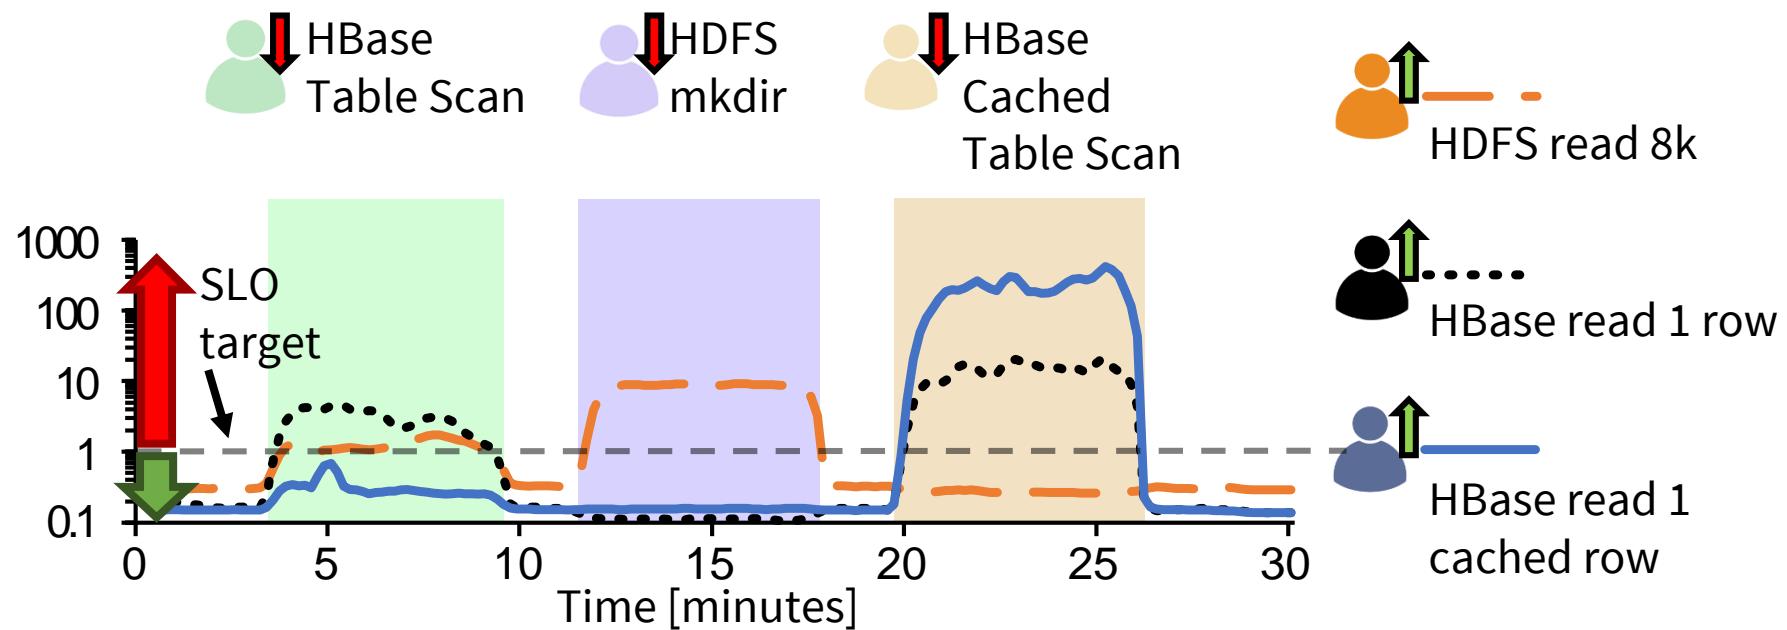


SLO-Normalized Latency

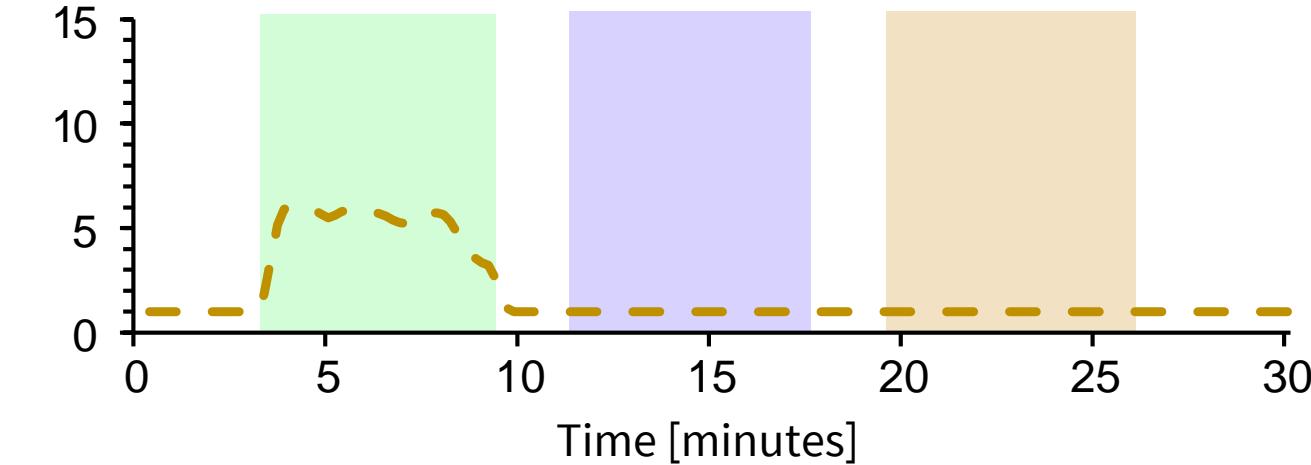




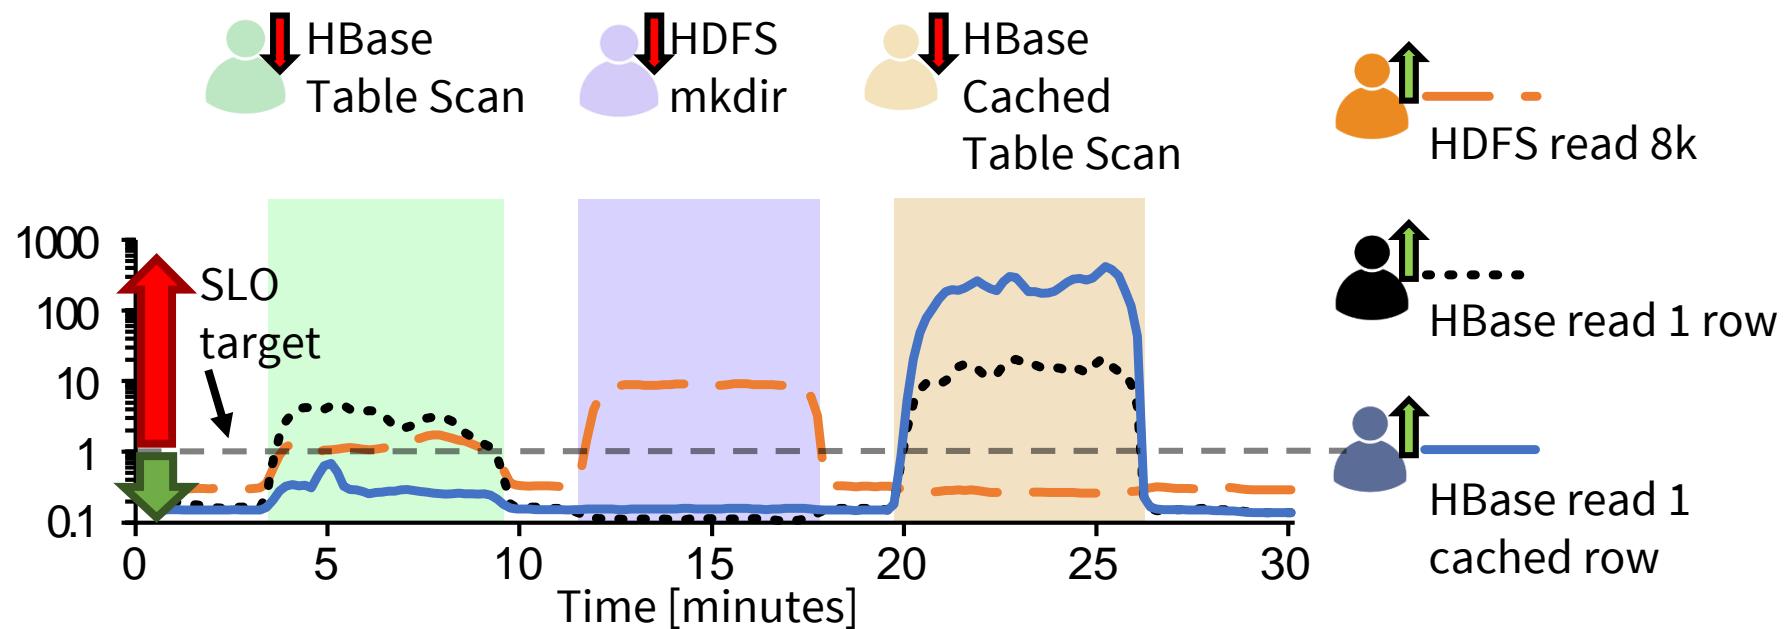
SLO-Normalized Latency



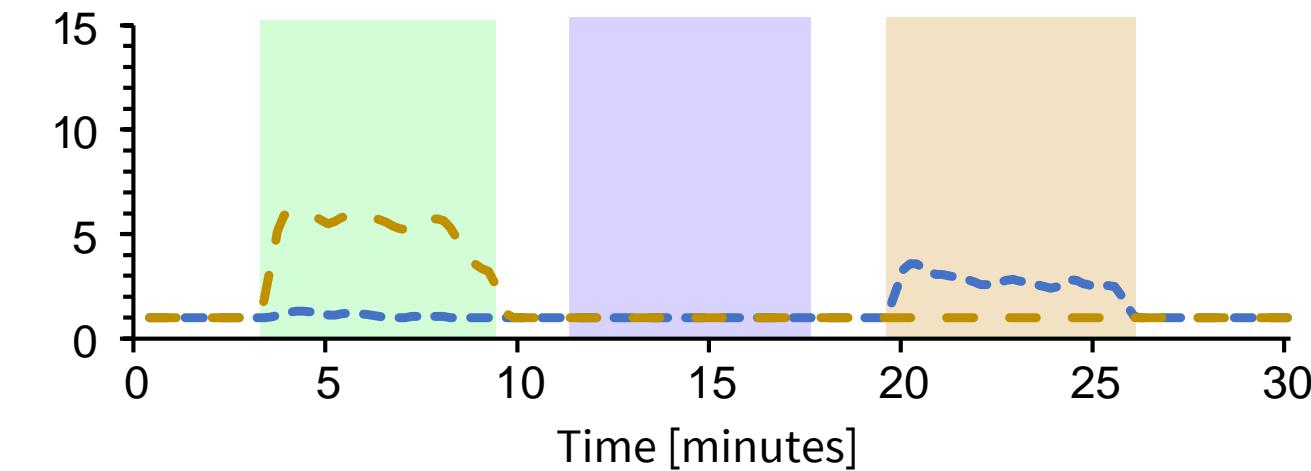
Slowdown



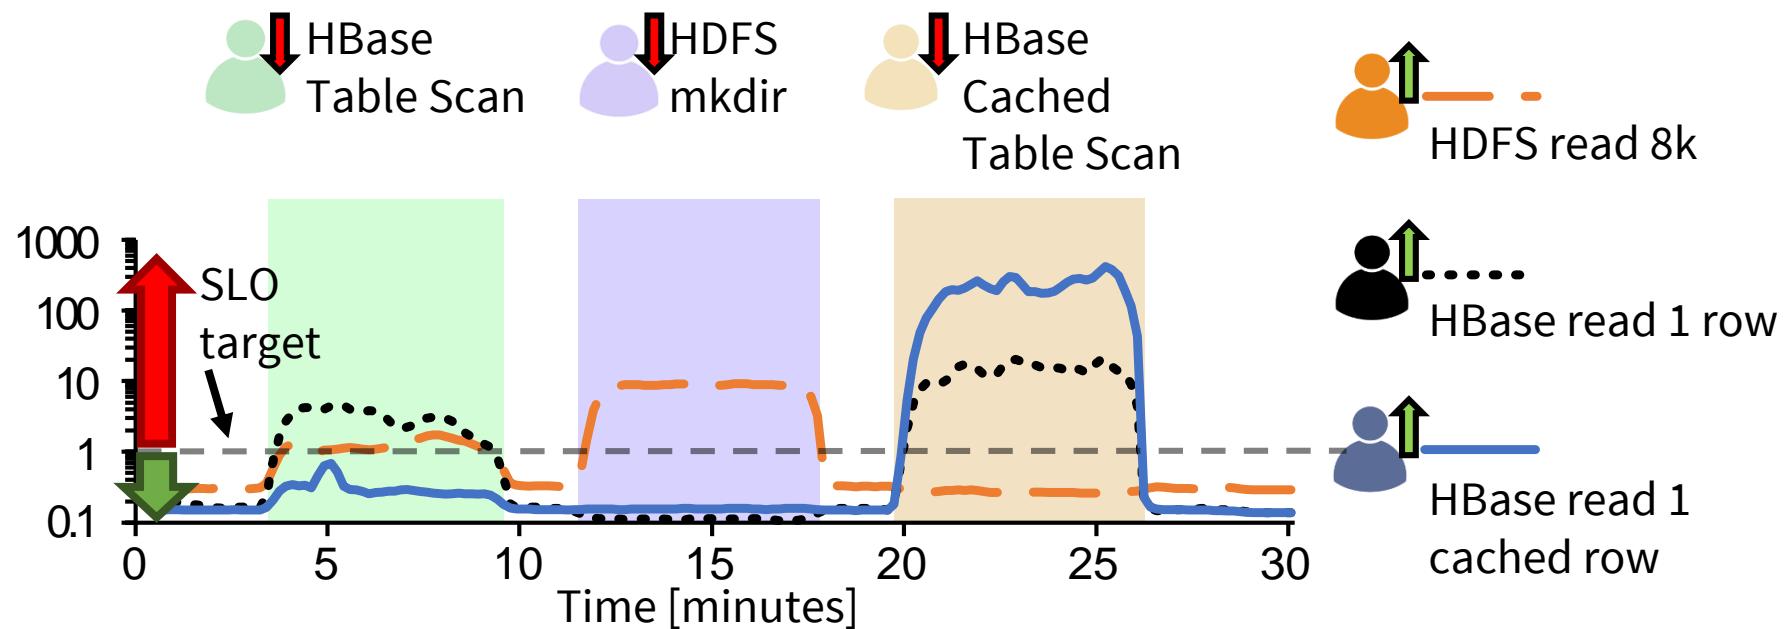
SLO-Normalized Latency



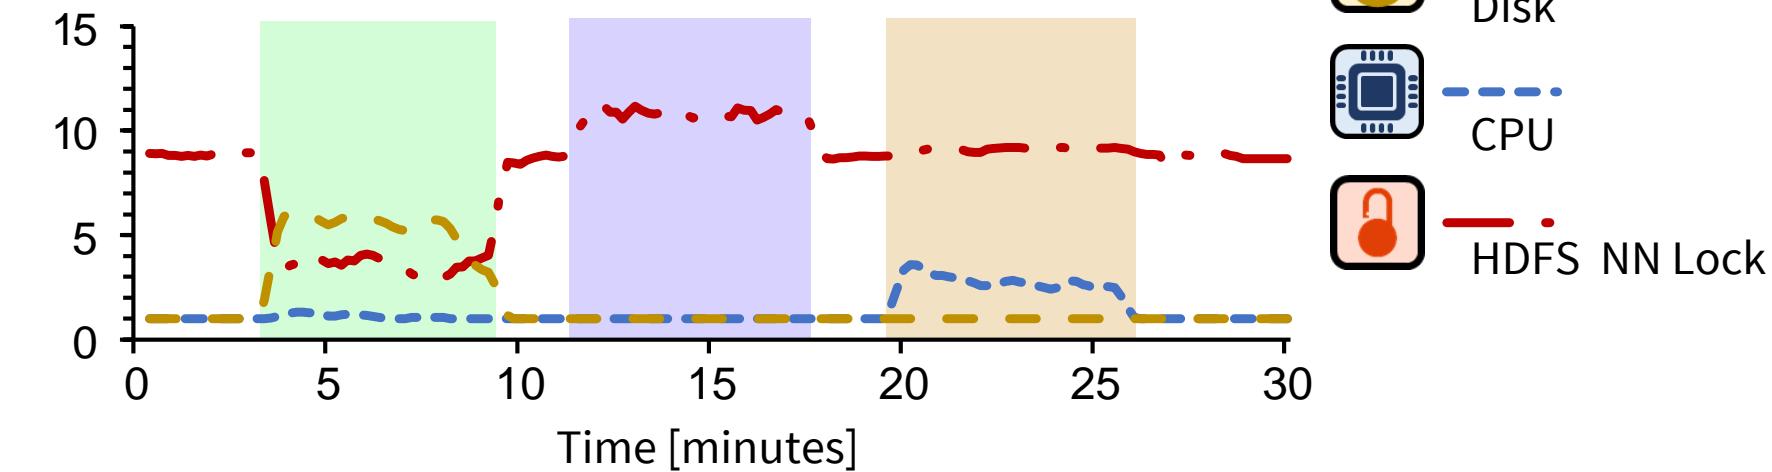
Slowdown



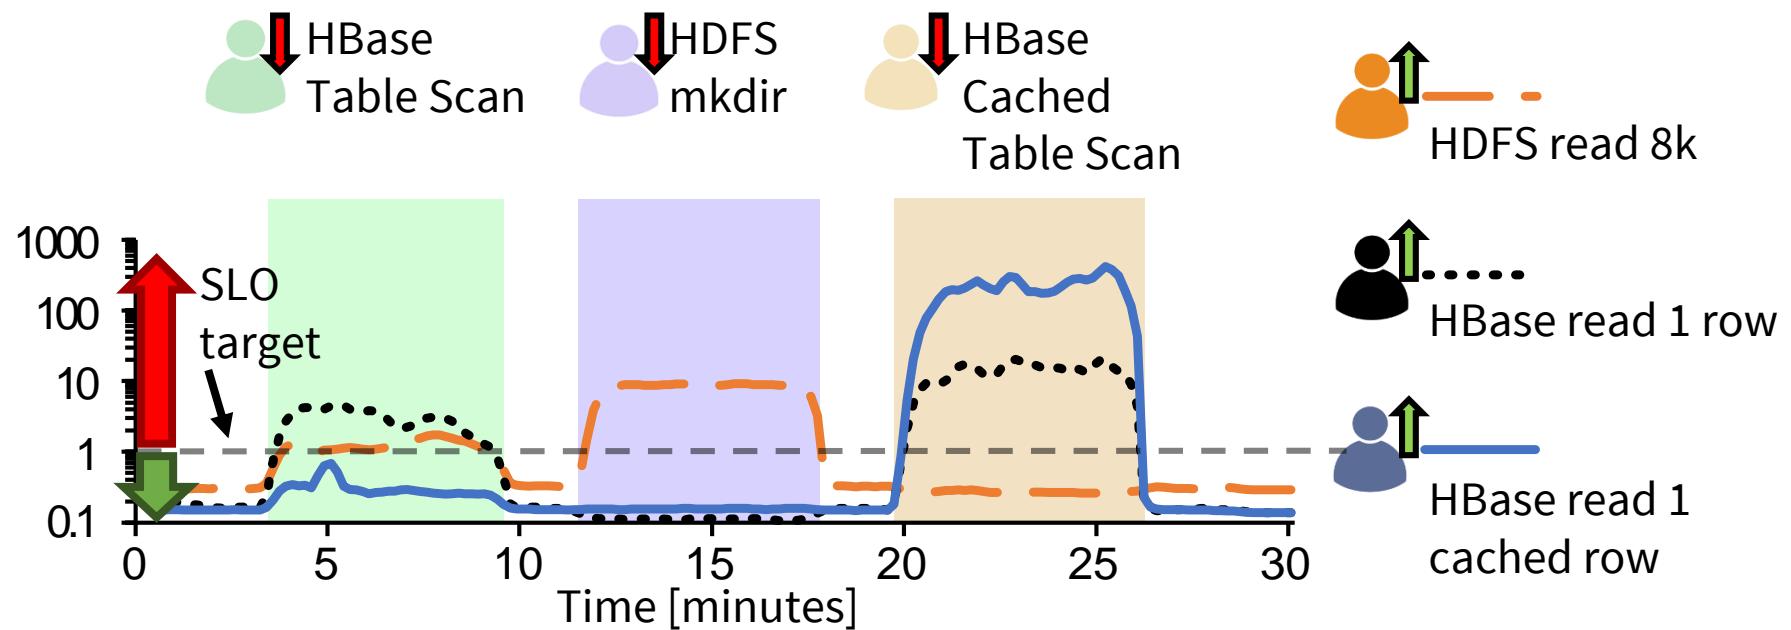
SLO-Normalized Latency



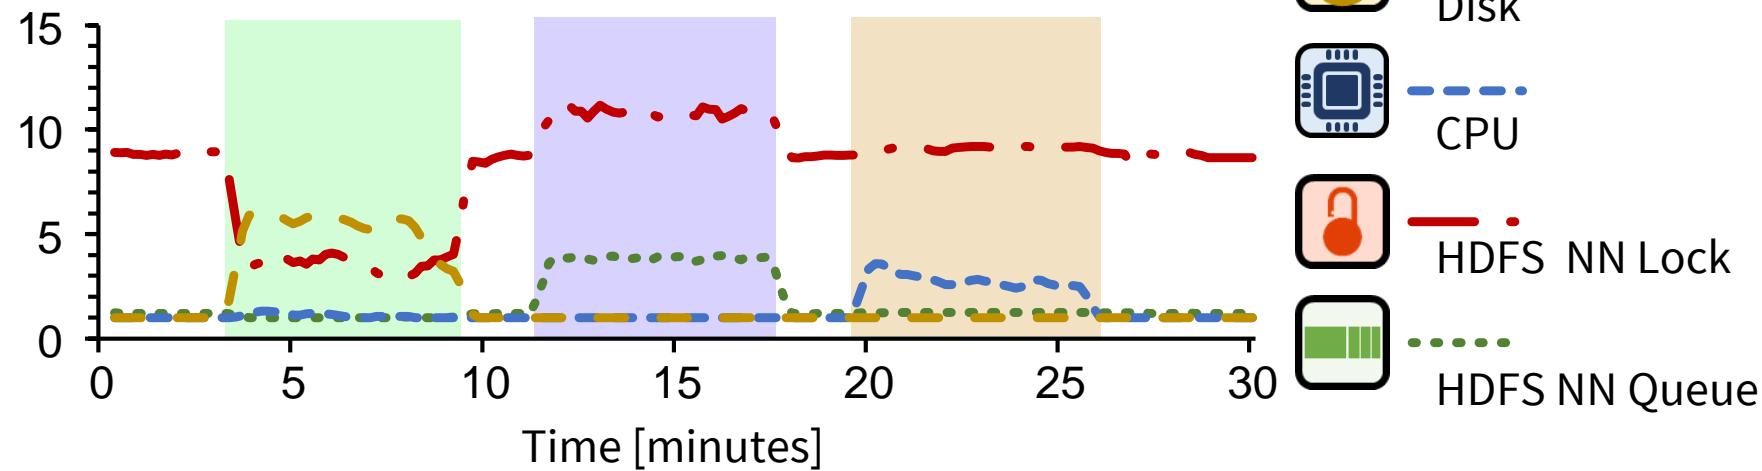
Slowdown



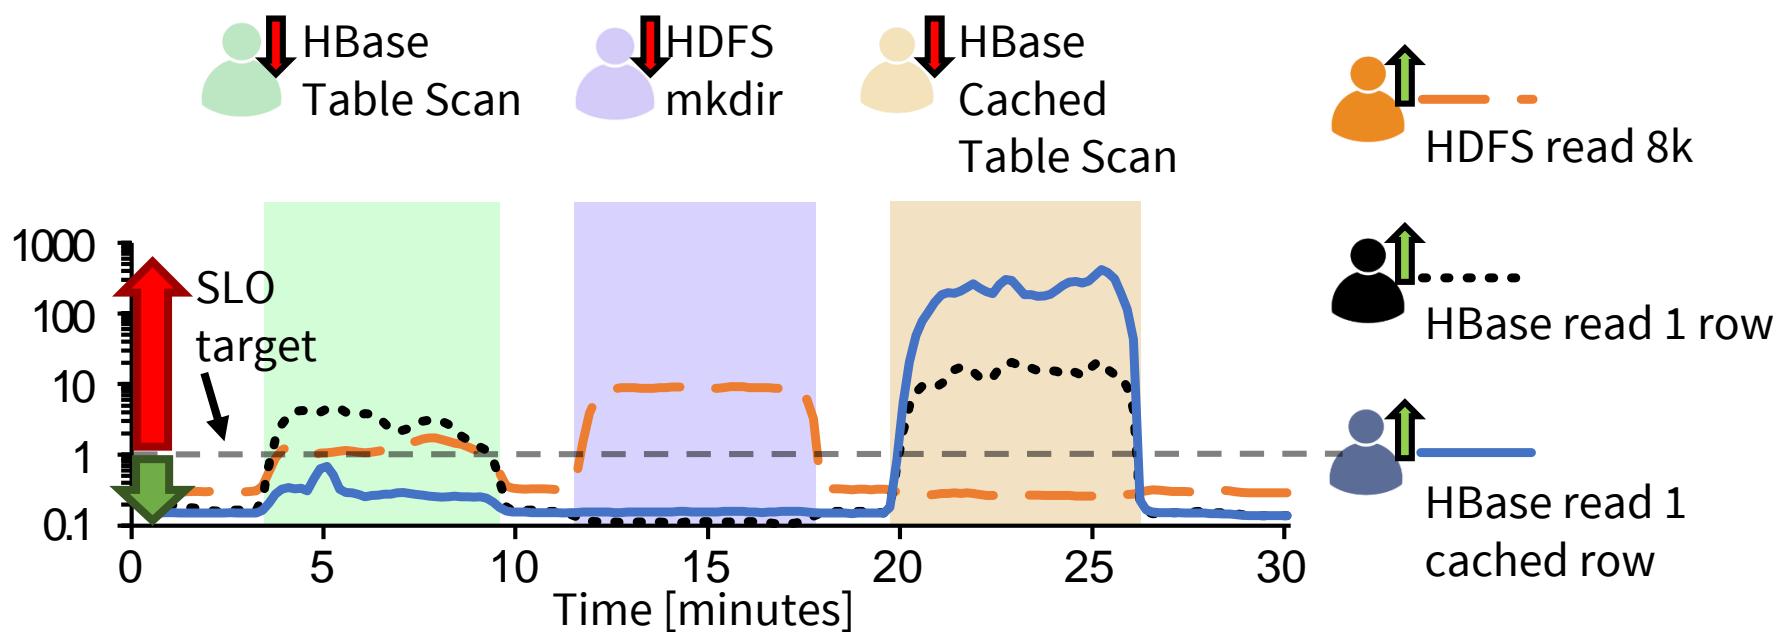
SLO-Normalized Latency



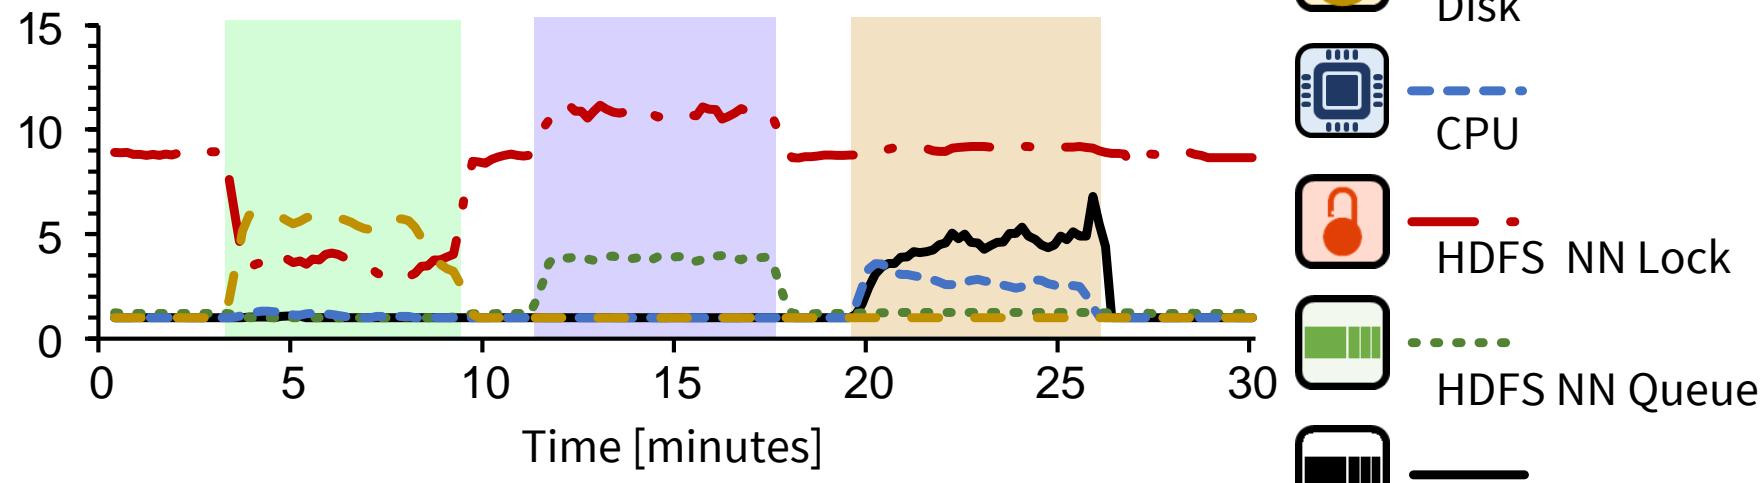
Slowdown



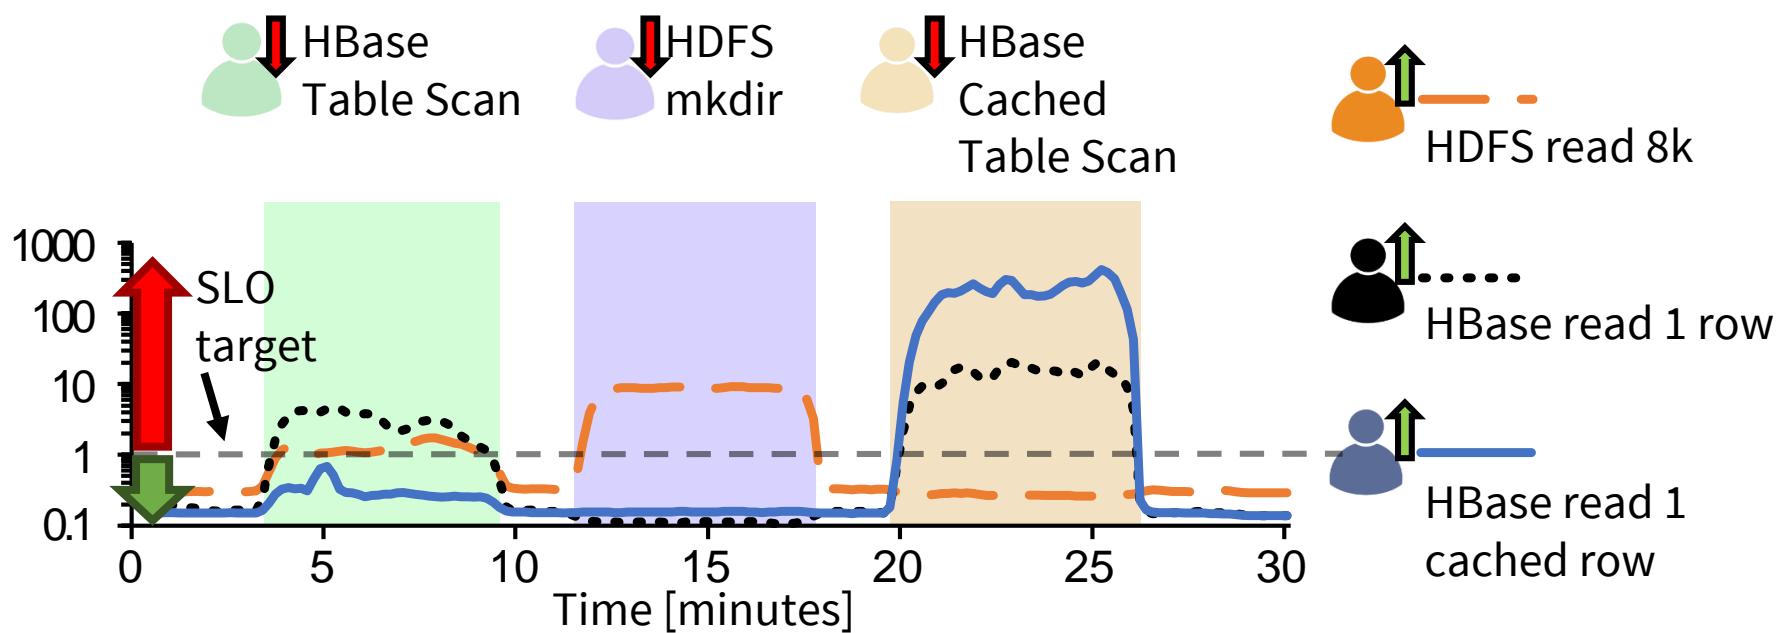
SLO-Normalized Latency



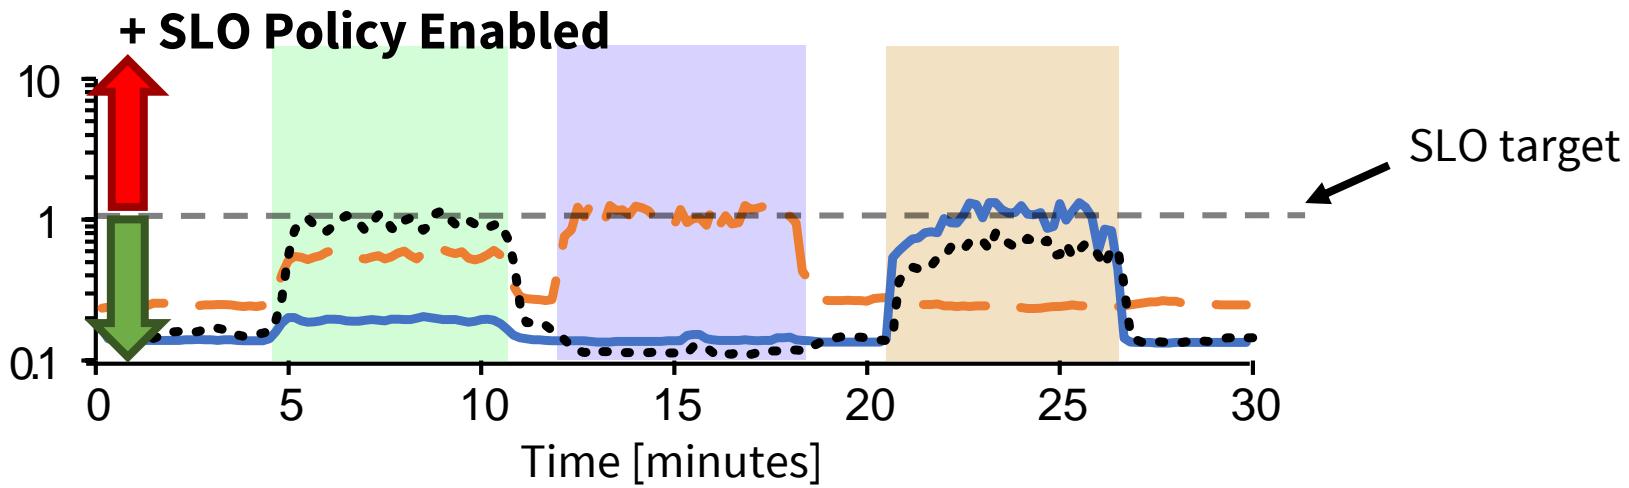
Slowdown



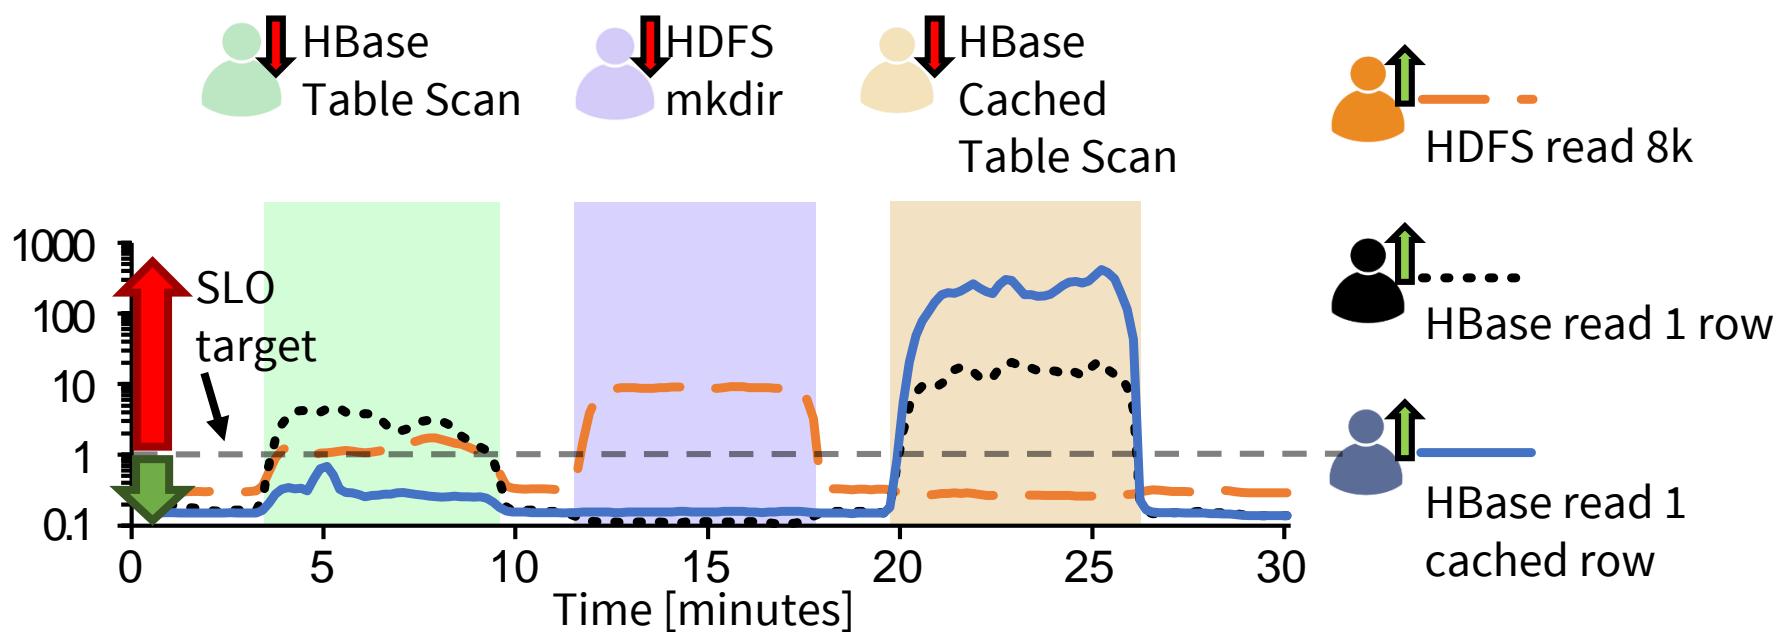
SLO-Normalized Latency



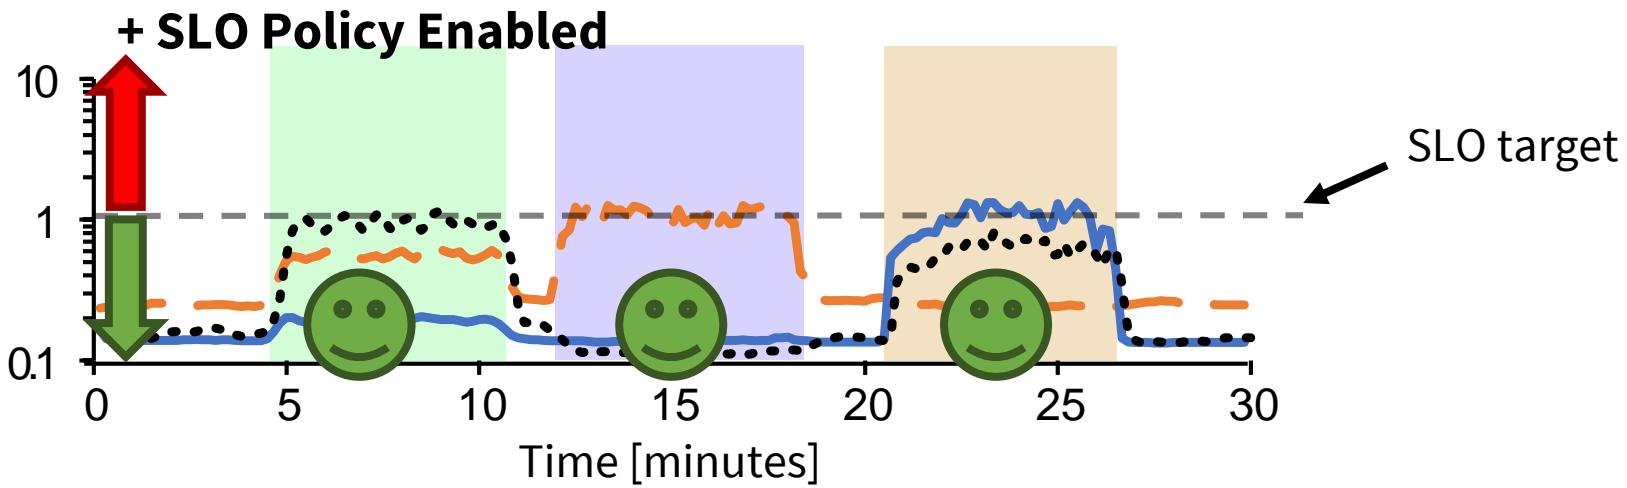
SLO-Normalized Latency



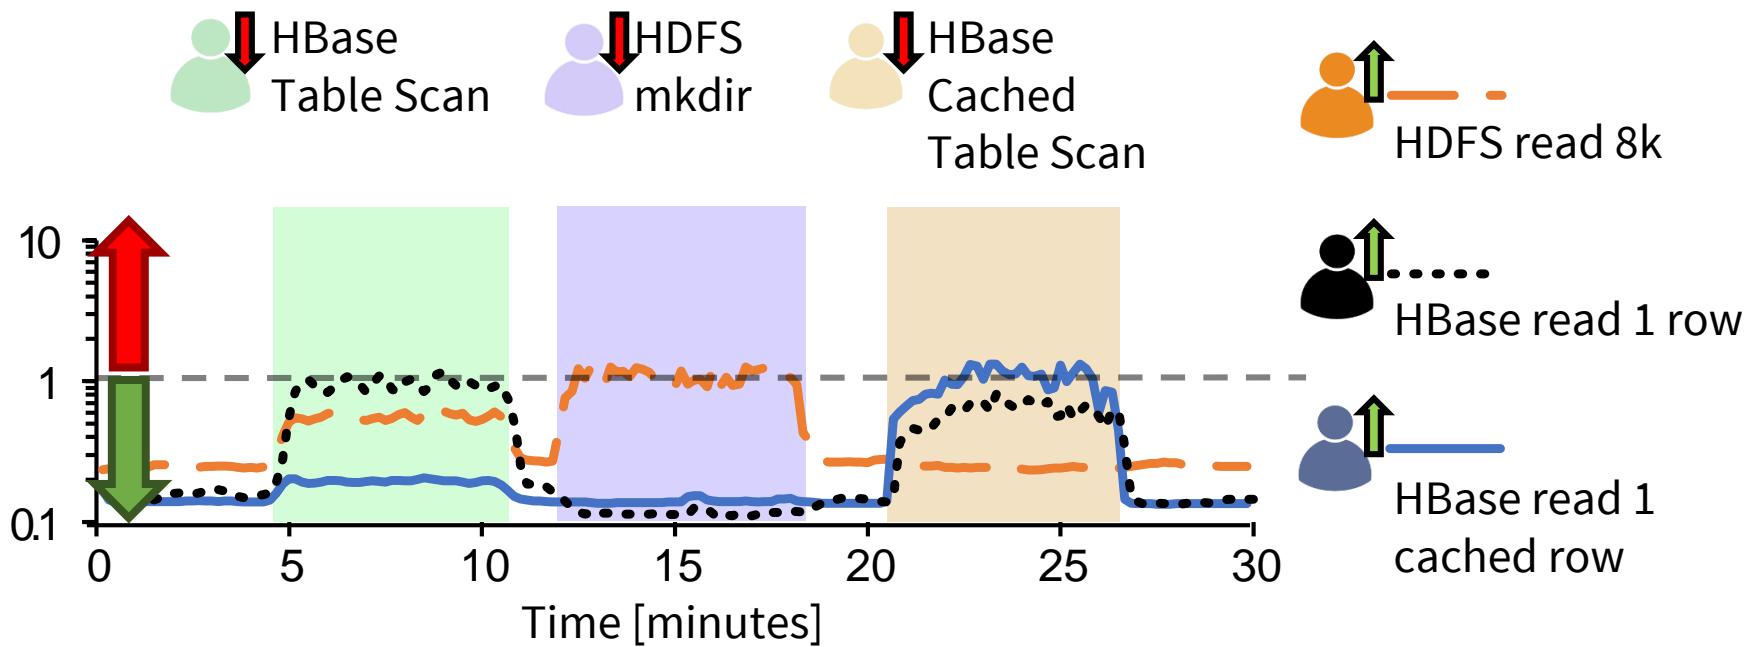
SLO-Normalized Latency



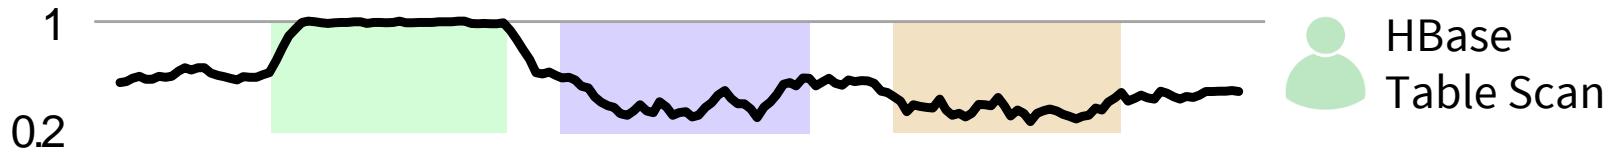
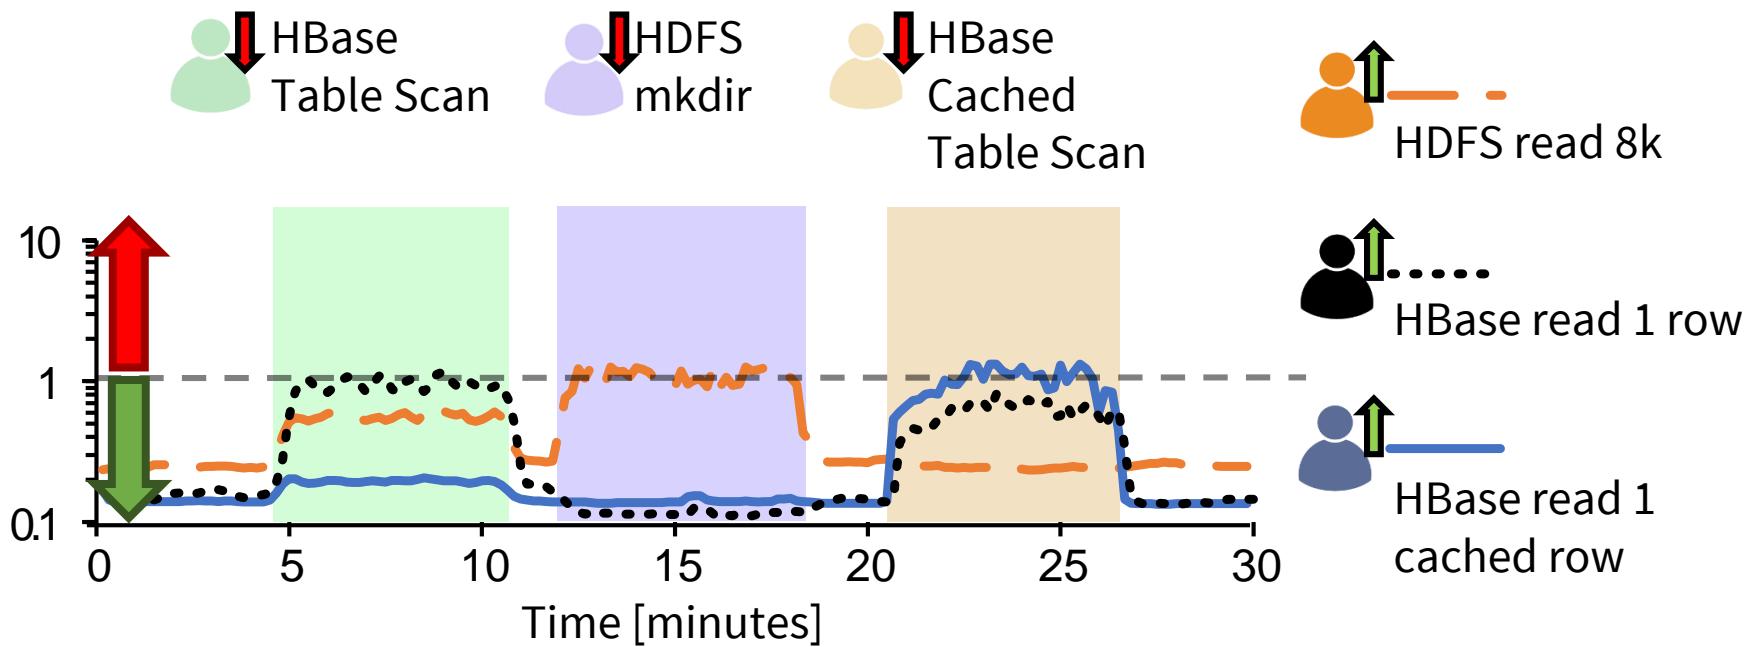
SLO-Normalized Latency



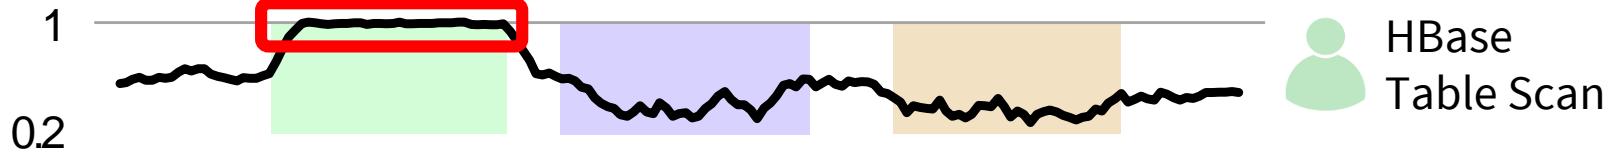
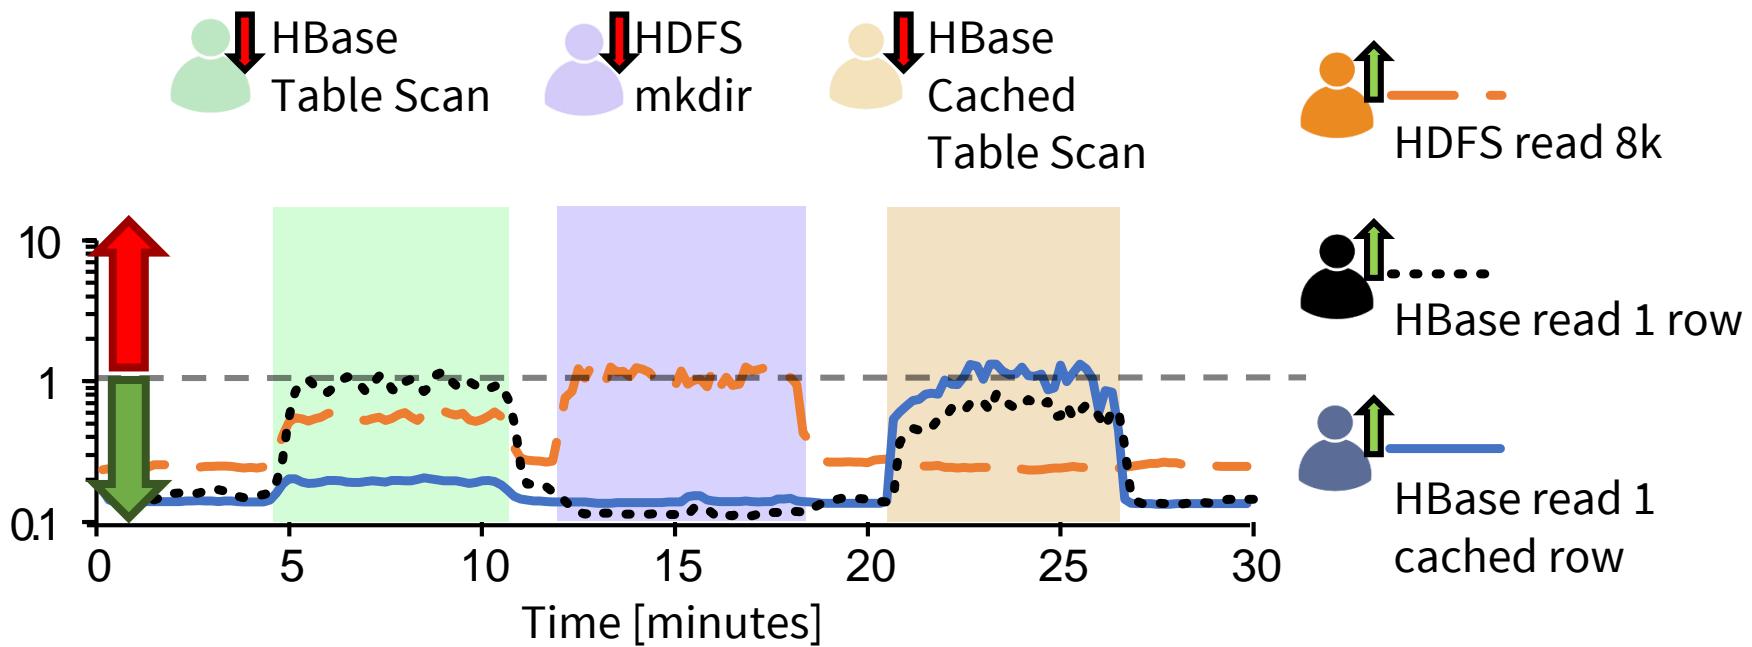
SLO-Normalized Latency



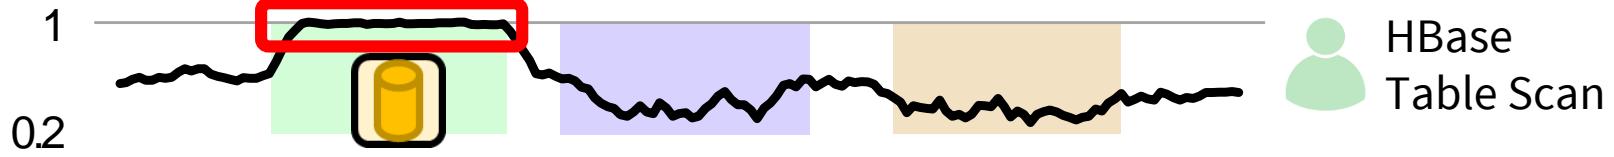
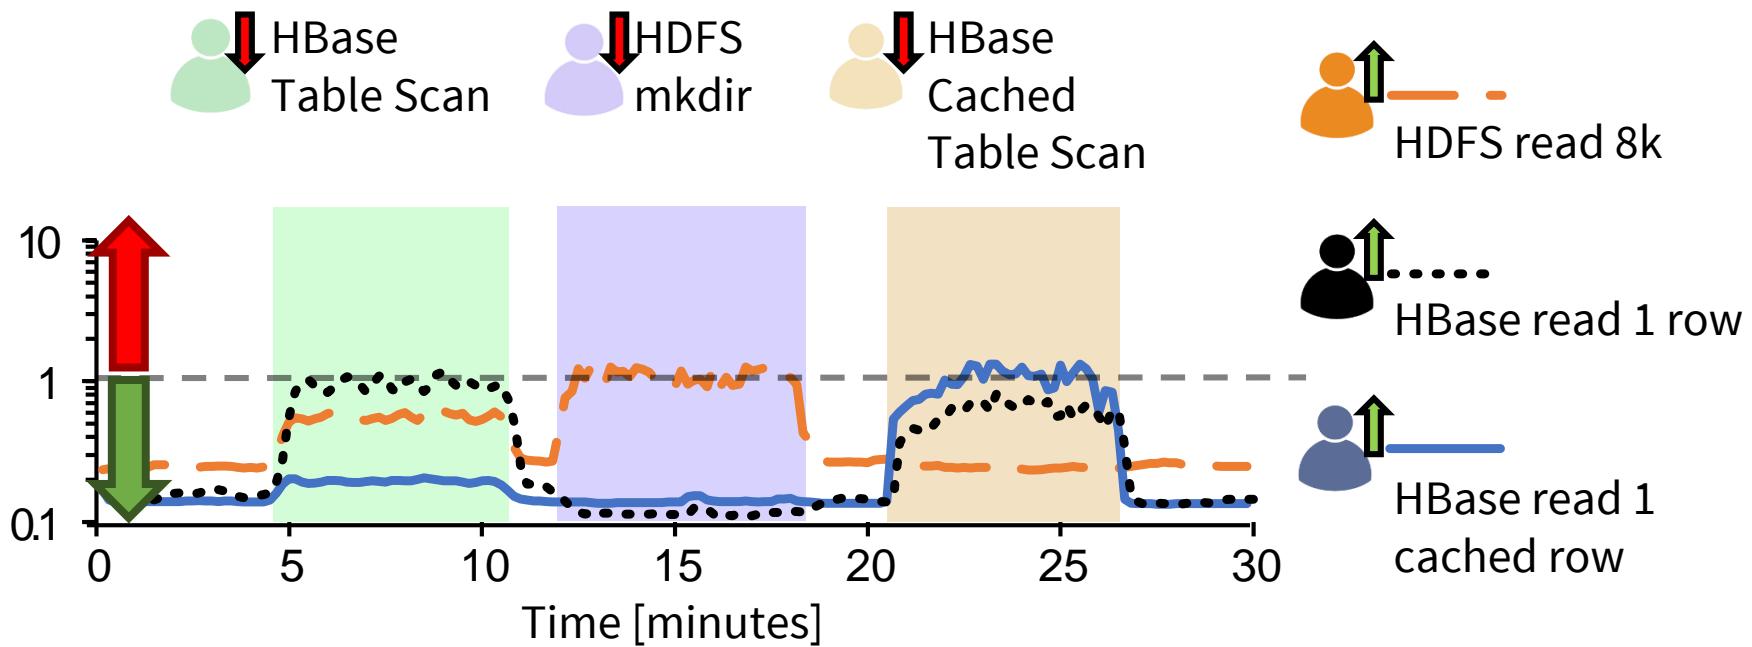
SLO-Normalized Latency



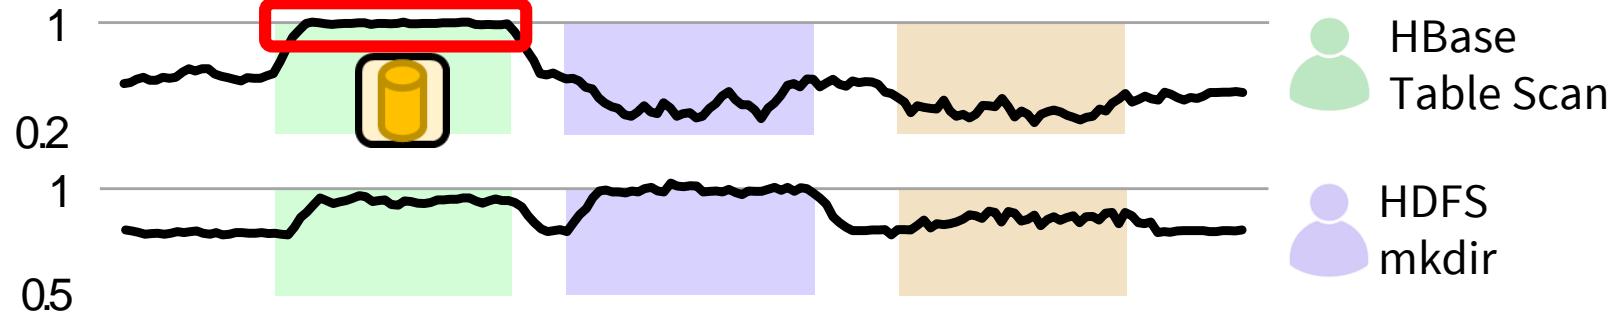
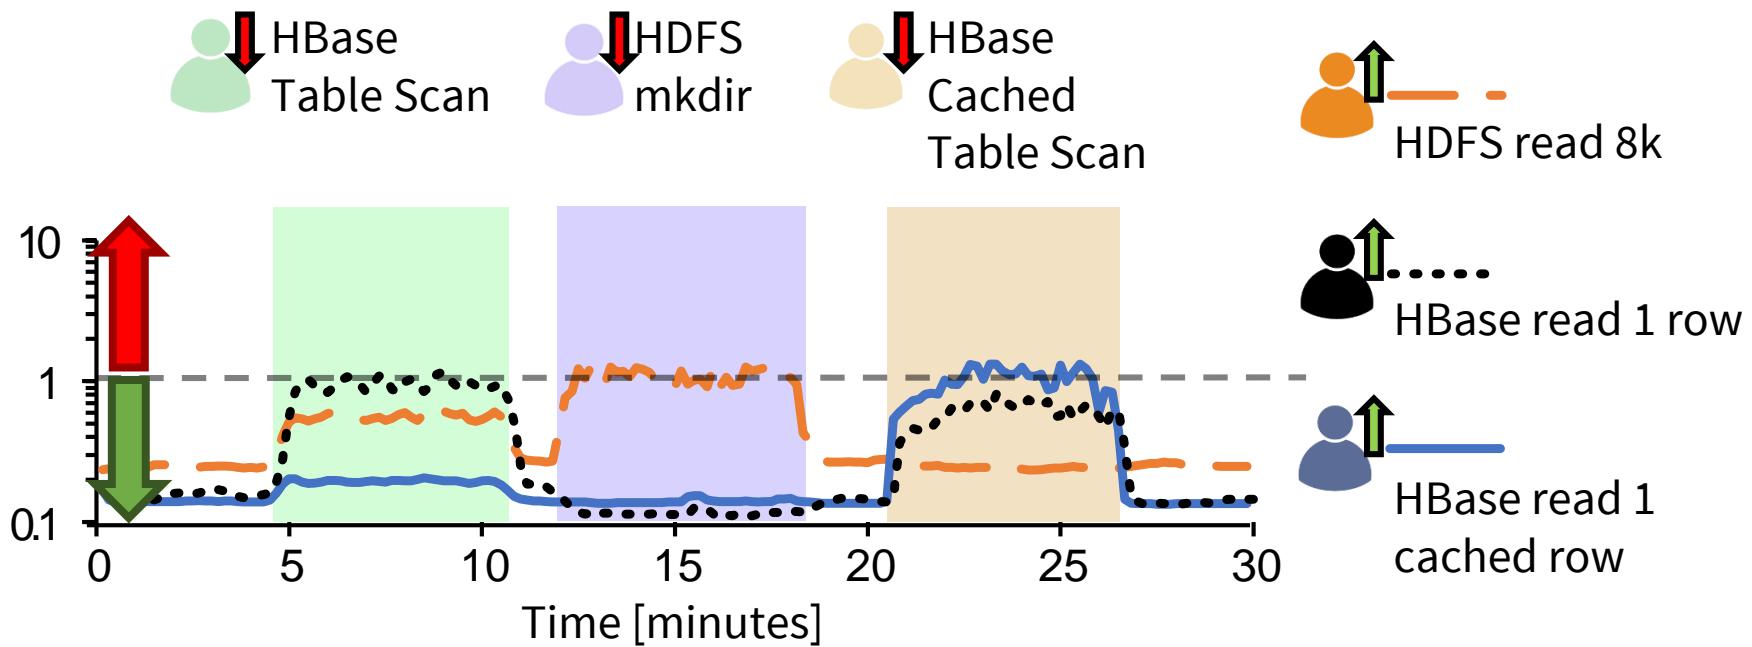
SLO-Normalized Latency



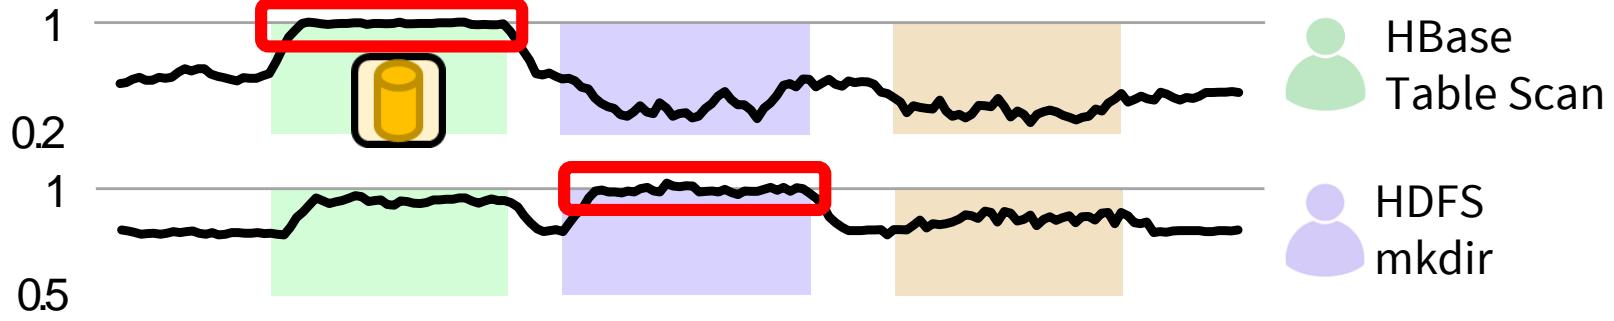
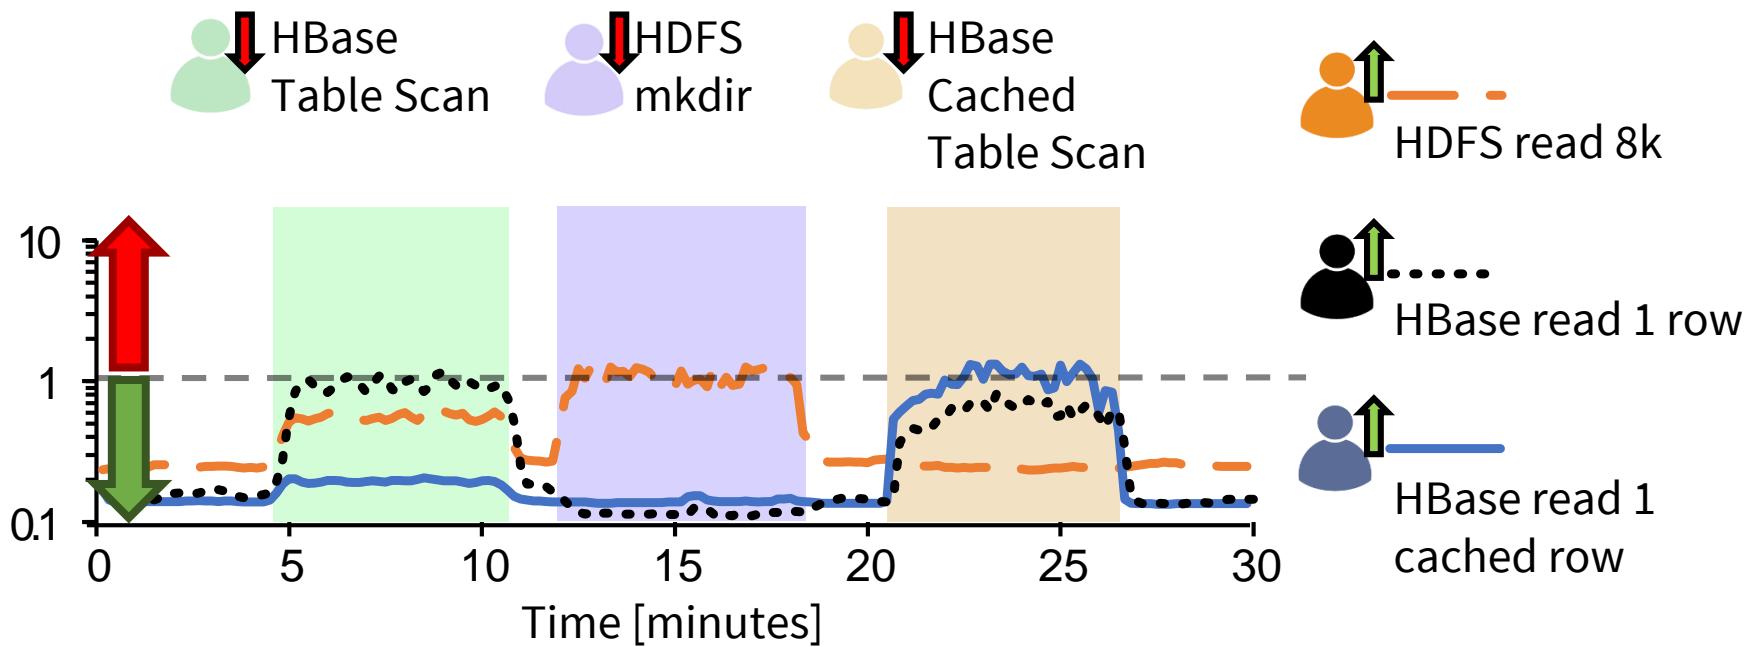
SLO-Normalized Latency



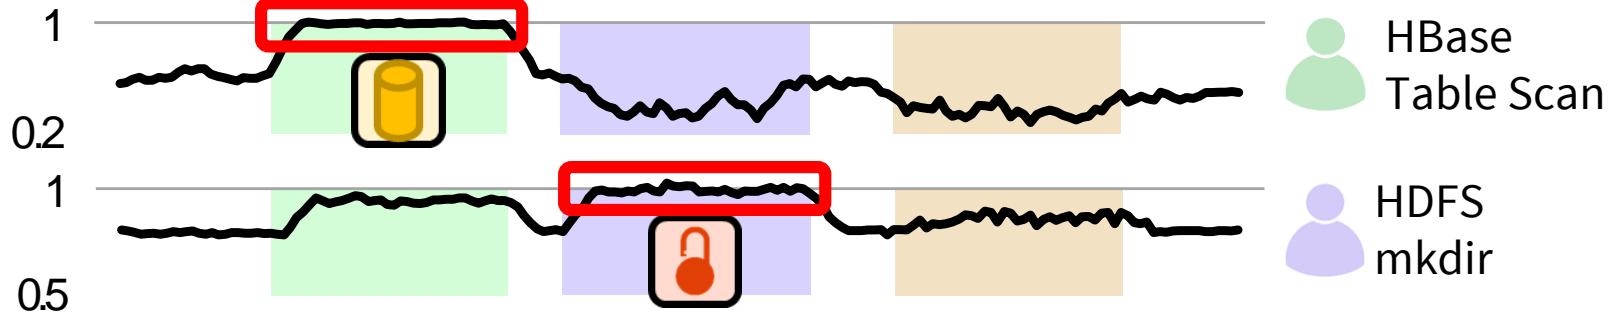
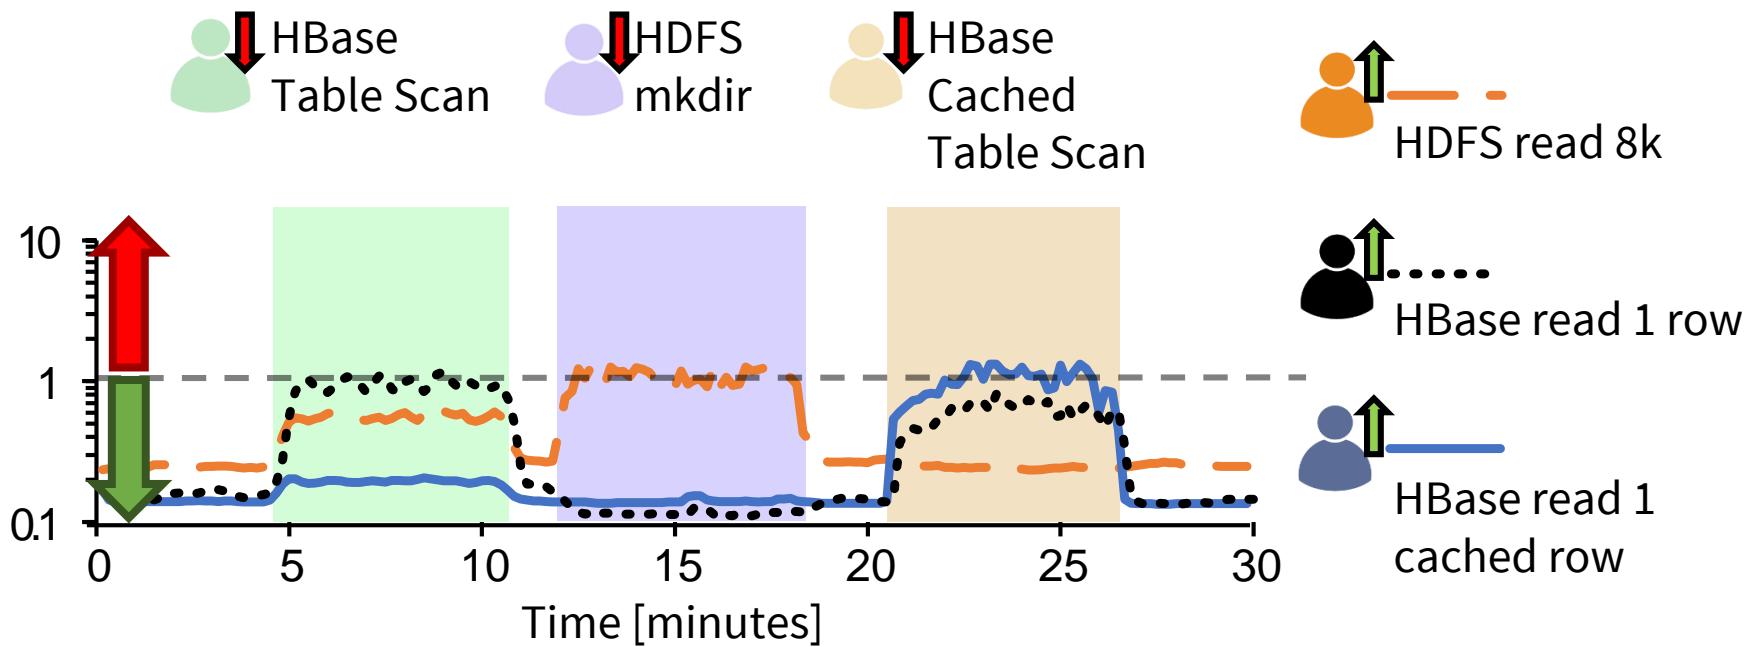
SLO-Normalized Latency



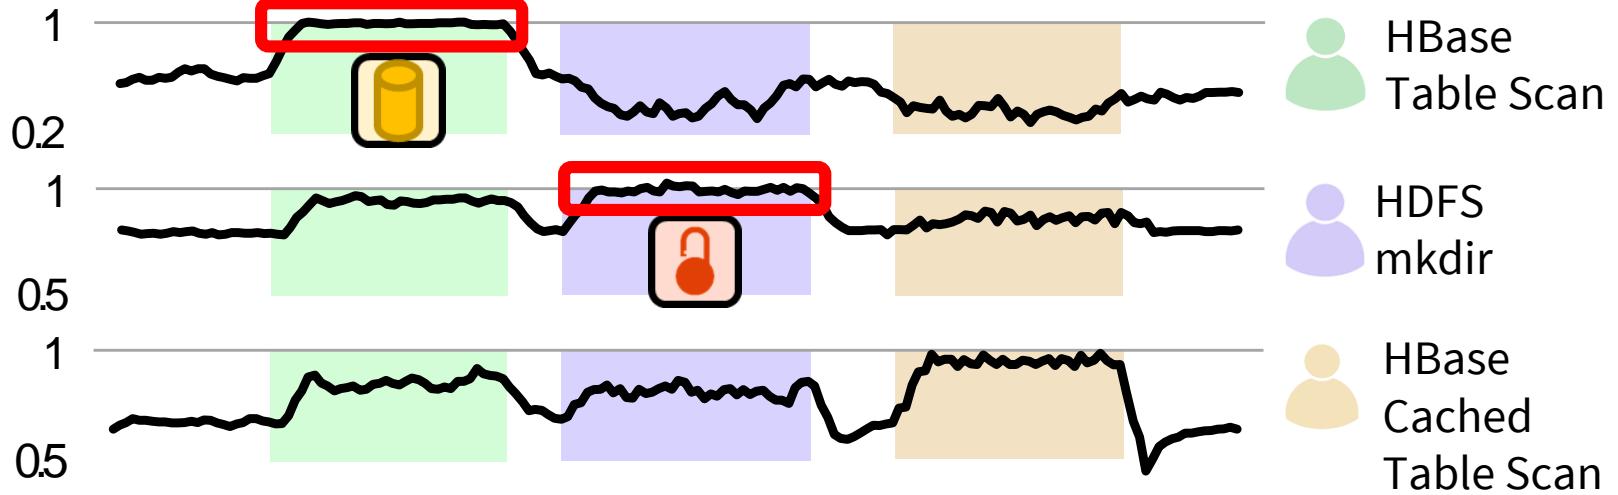
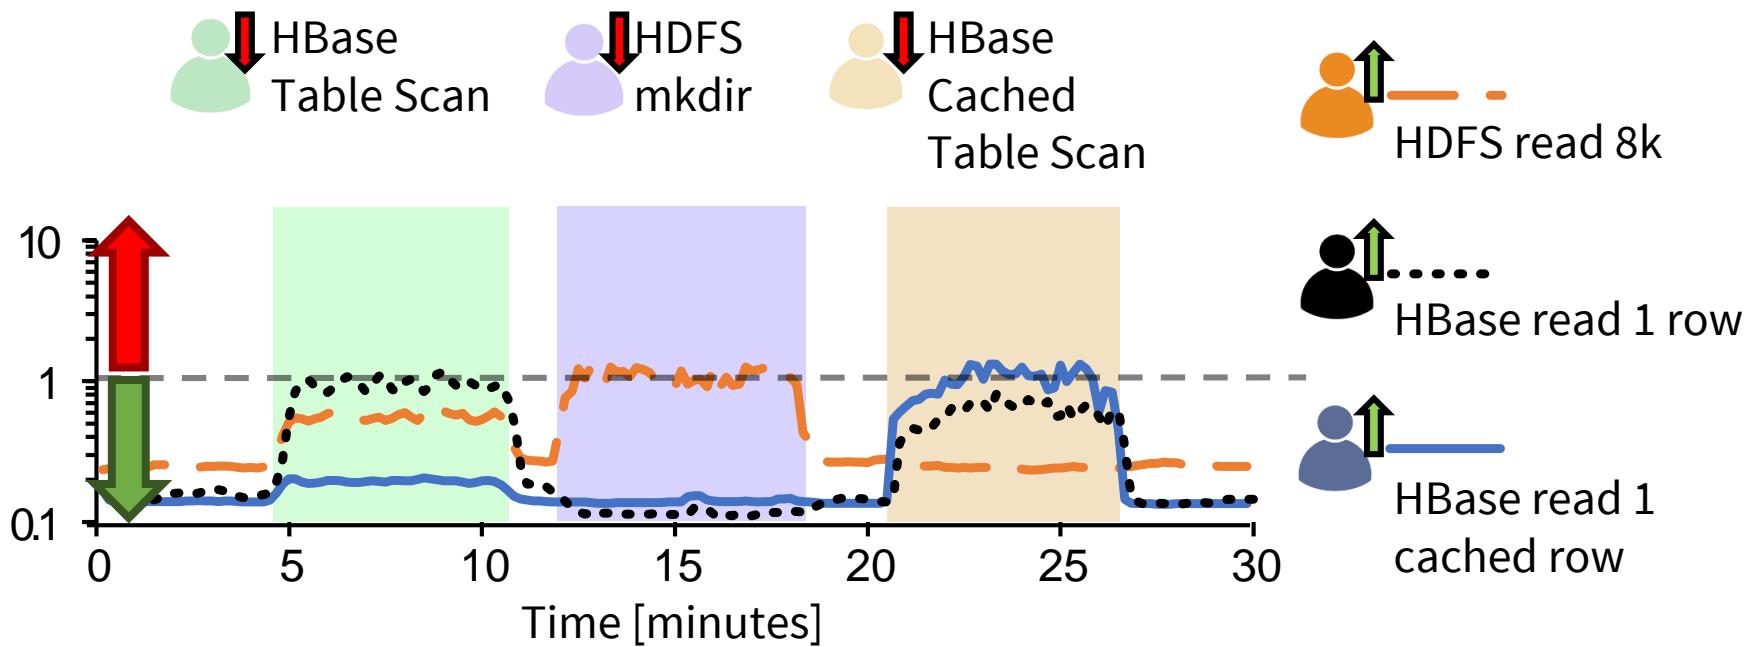
SLO-Normalized Latency



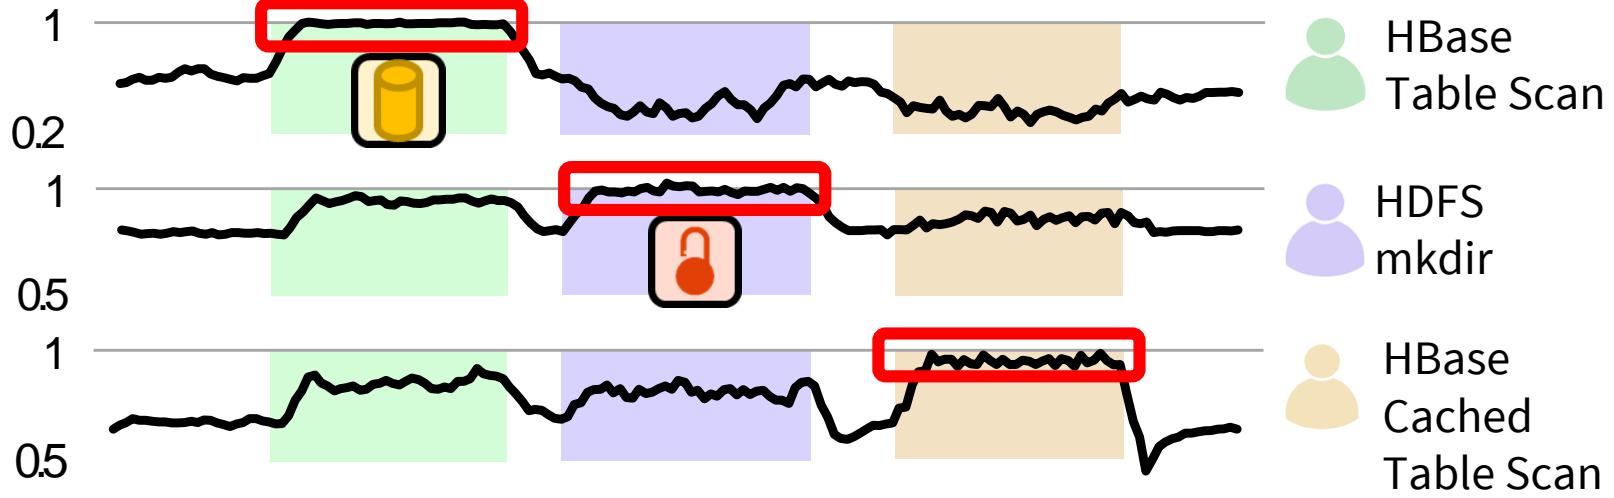
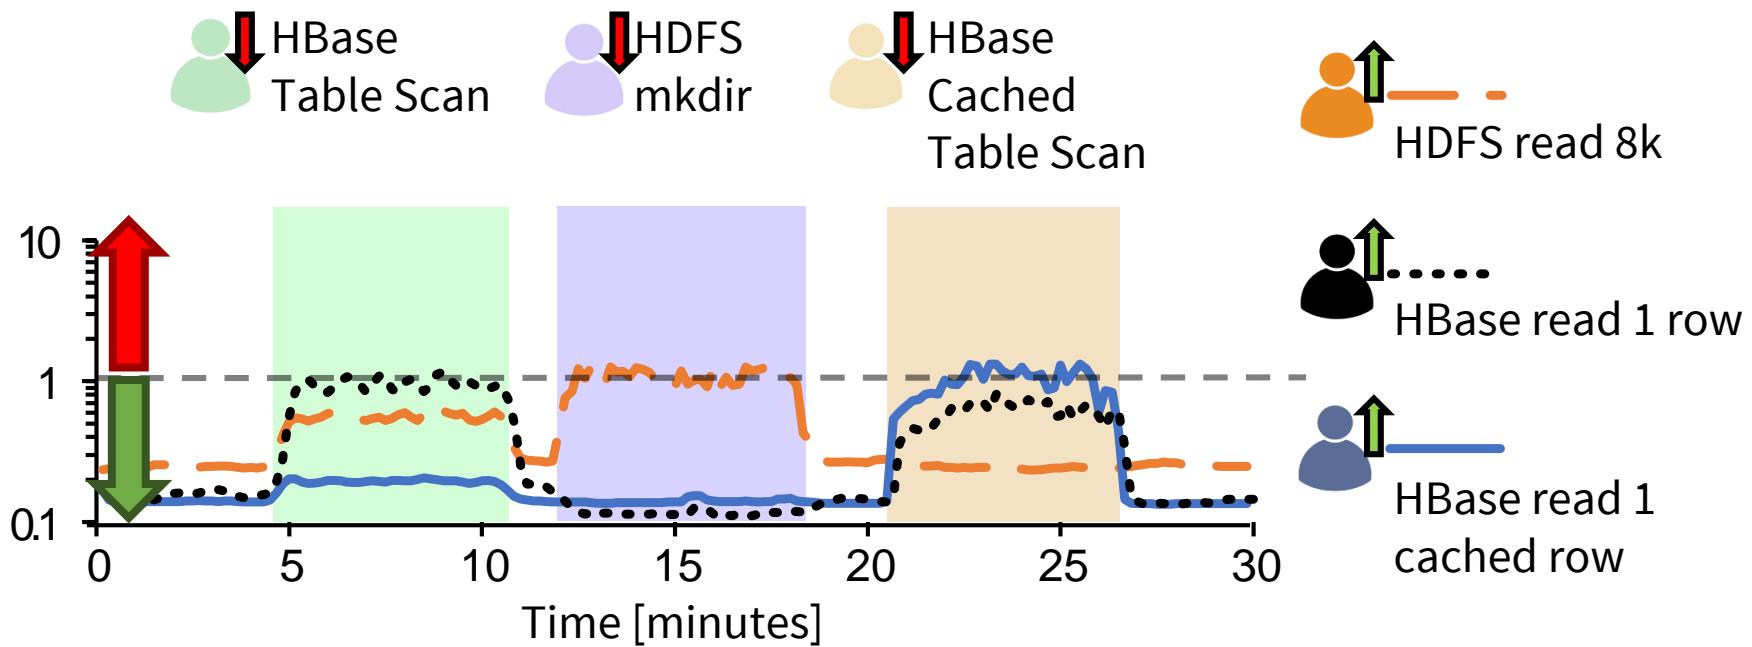
SLO-Normalized Latency



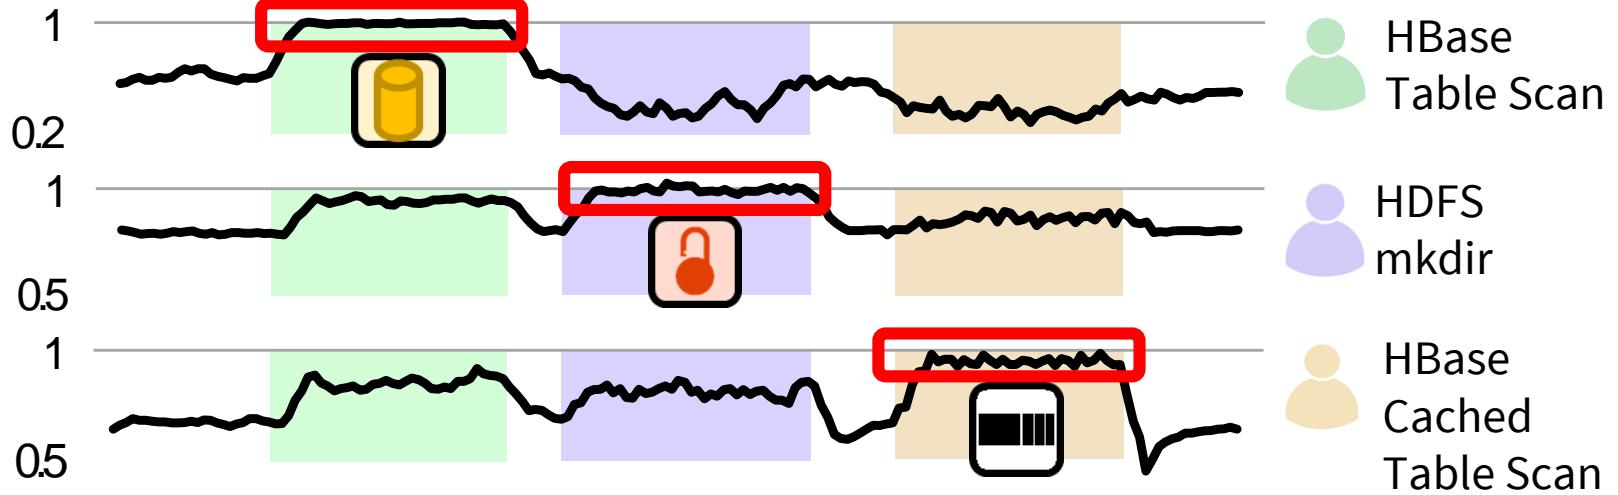
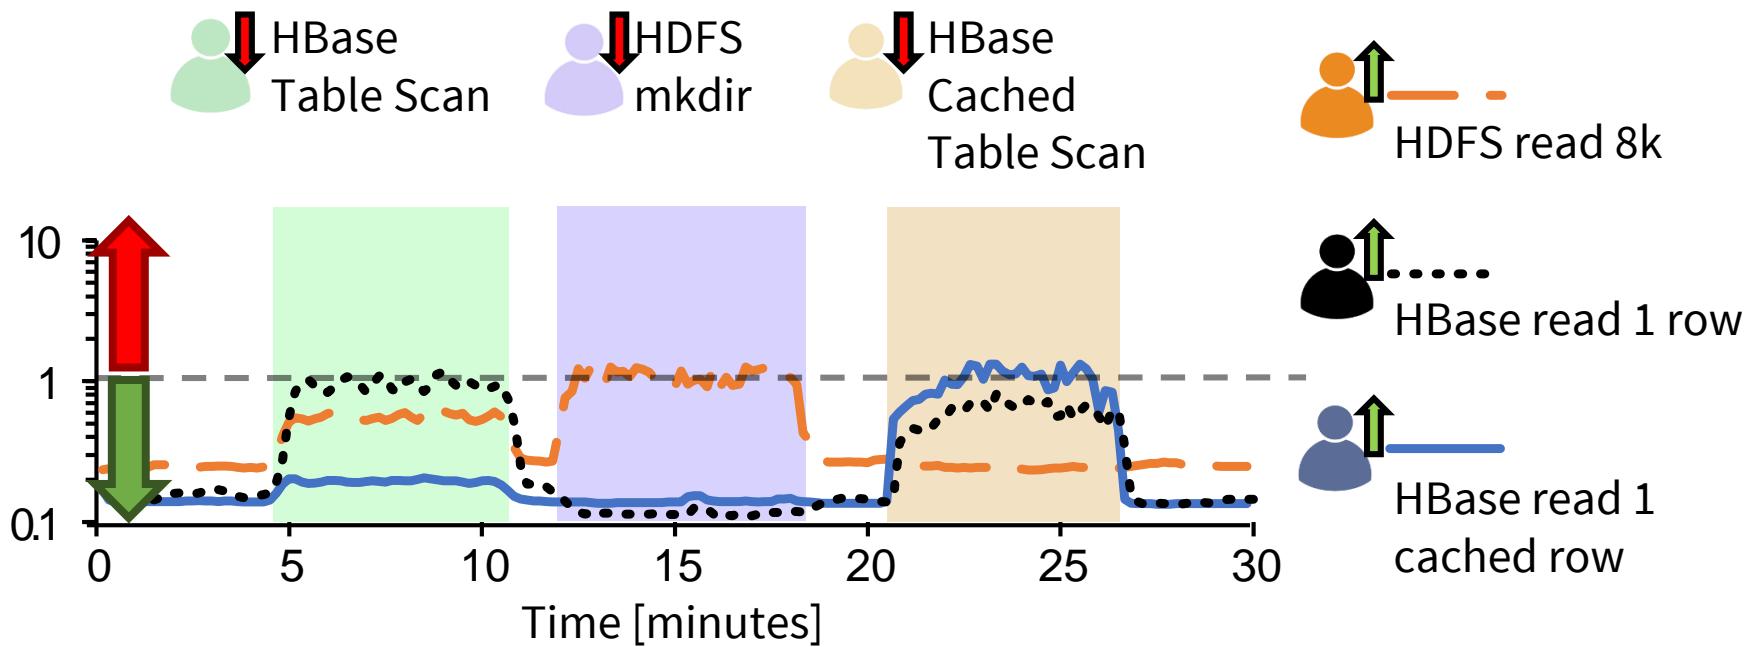
SLO-Normalized Latency



SLO-Normalized Latency

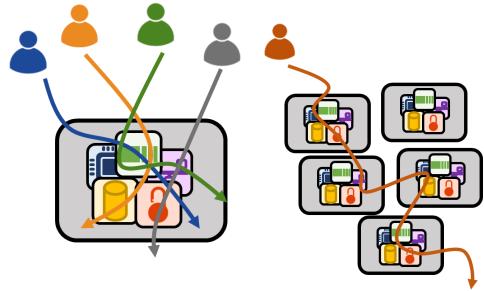


SLO-Normalized Latency

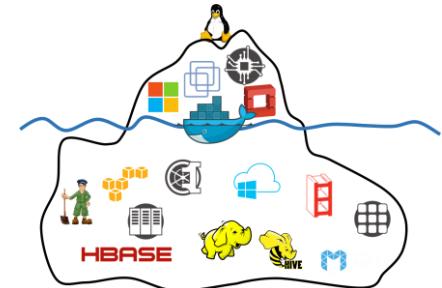


# *Retro Conclusion*

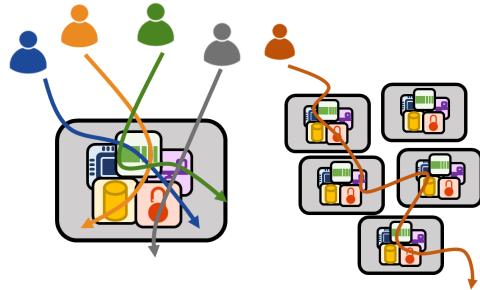
# Retro Conclusion



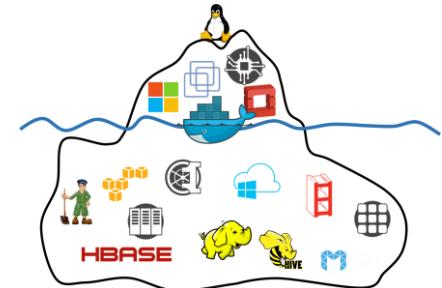
Resource management for *shared* distributed systems



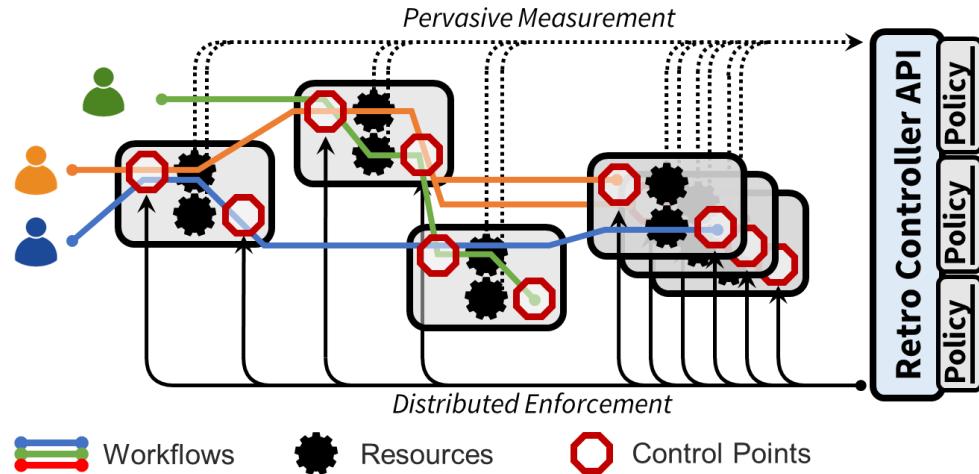
# Retro Conclusion



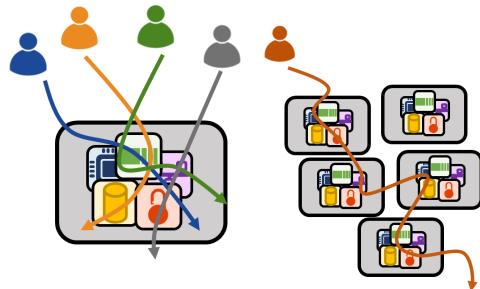
Resource management for *shared* distributed systems



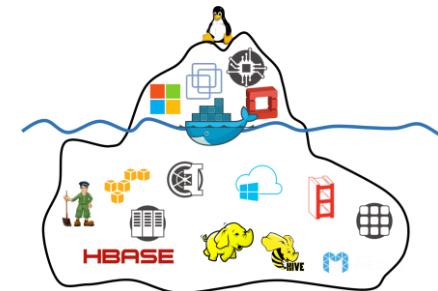
*Centralized* resource management



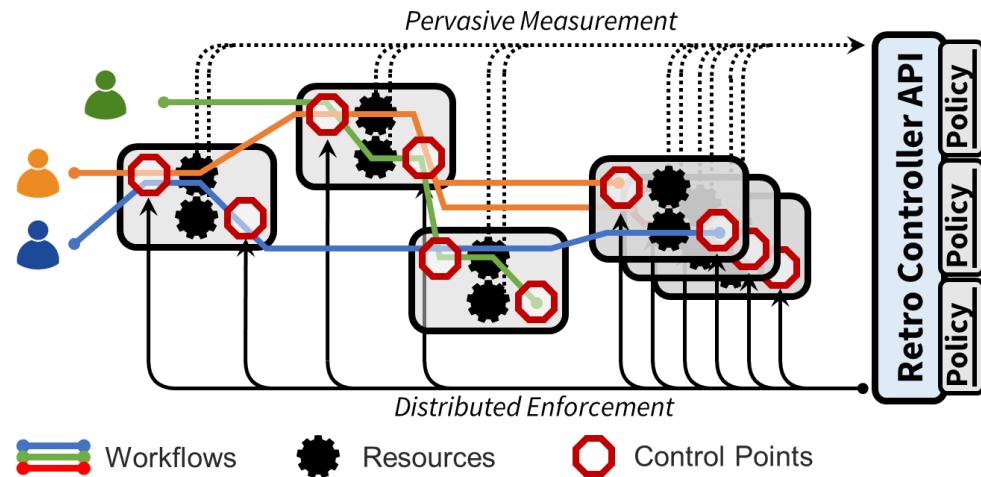
# Retro Conclusion



Resource management for *shared* distributed systems

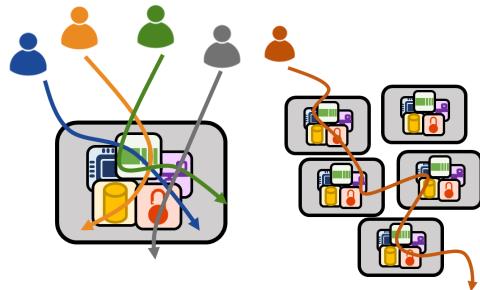


Centralized resource management

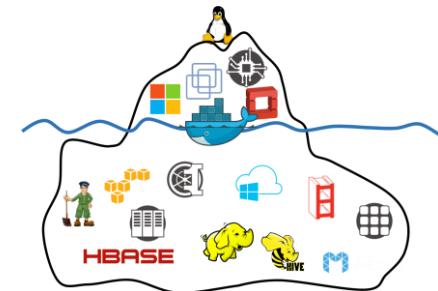


Comprehensive: resources, processes, tenants, background tasks

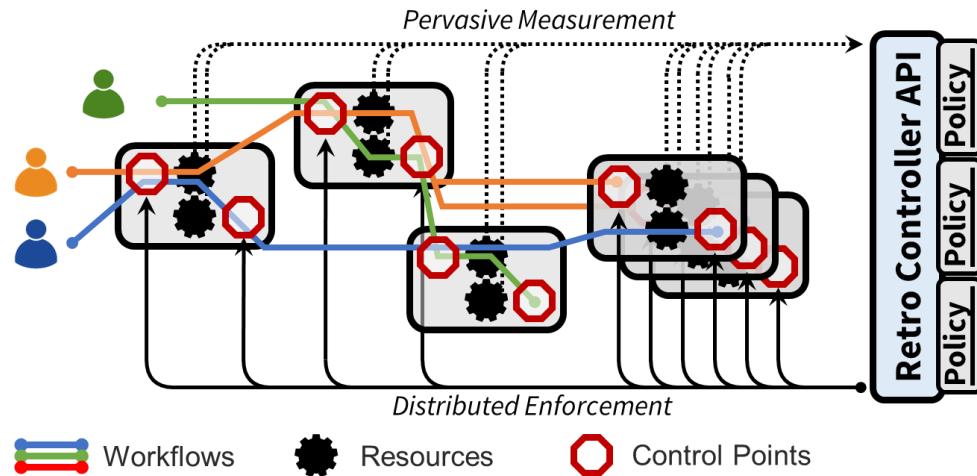
# Retro Conclusion



Resource management for *shared* distributed systems



*Centralized* resource management



Comprehensive: resources, processes, tenants, background tasks

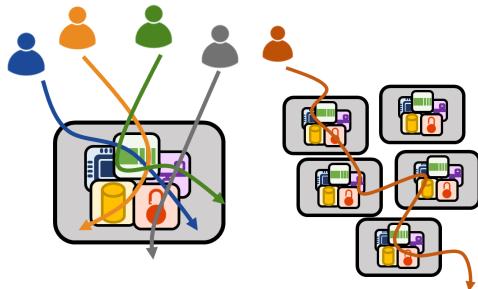
Abstractions for writing *concise, general-purpose* policies:

Workflows

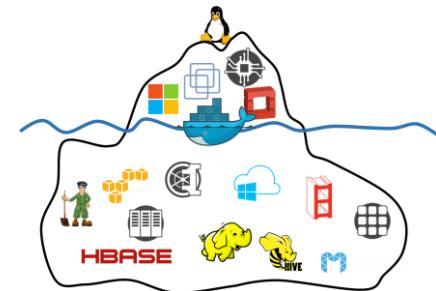
Resources (slowdown, load)

Control points

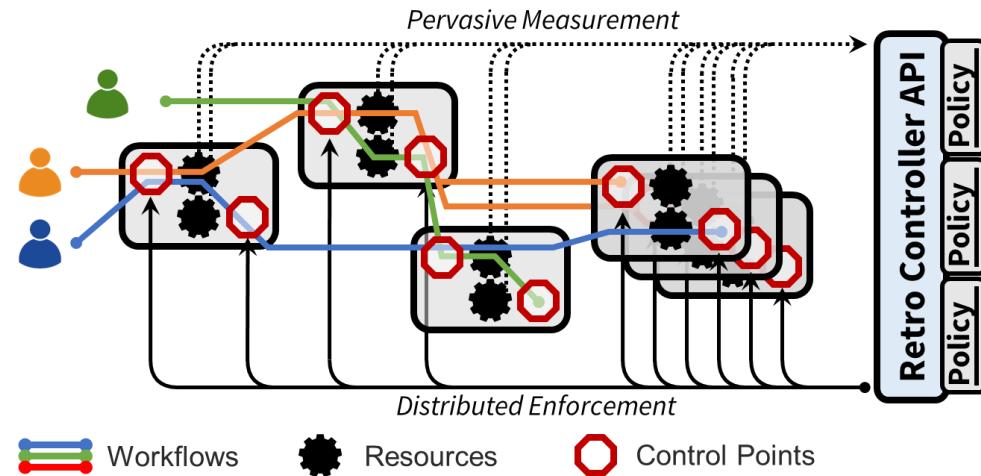
# Retro Conclusion



Resource management for *shared* distributed systems



*Centralized* resource management



*Comprehensive*: resources, processes, tenants, background tasks

Abstractions for writing *concise, general-purpose* policies:

- Workflows
- Resources (slowdown, load)
- Control points

